

# Neural Network Introduction

Ling 282/482: Deep Learning for Computational Linguistics

C.M. Downey

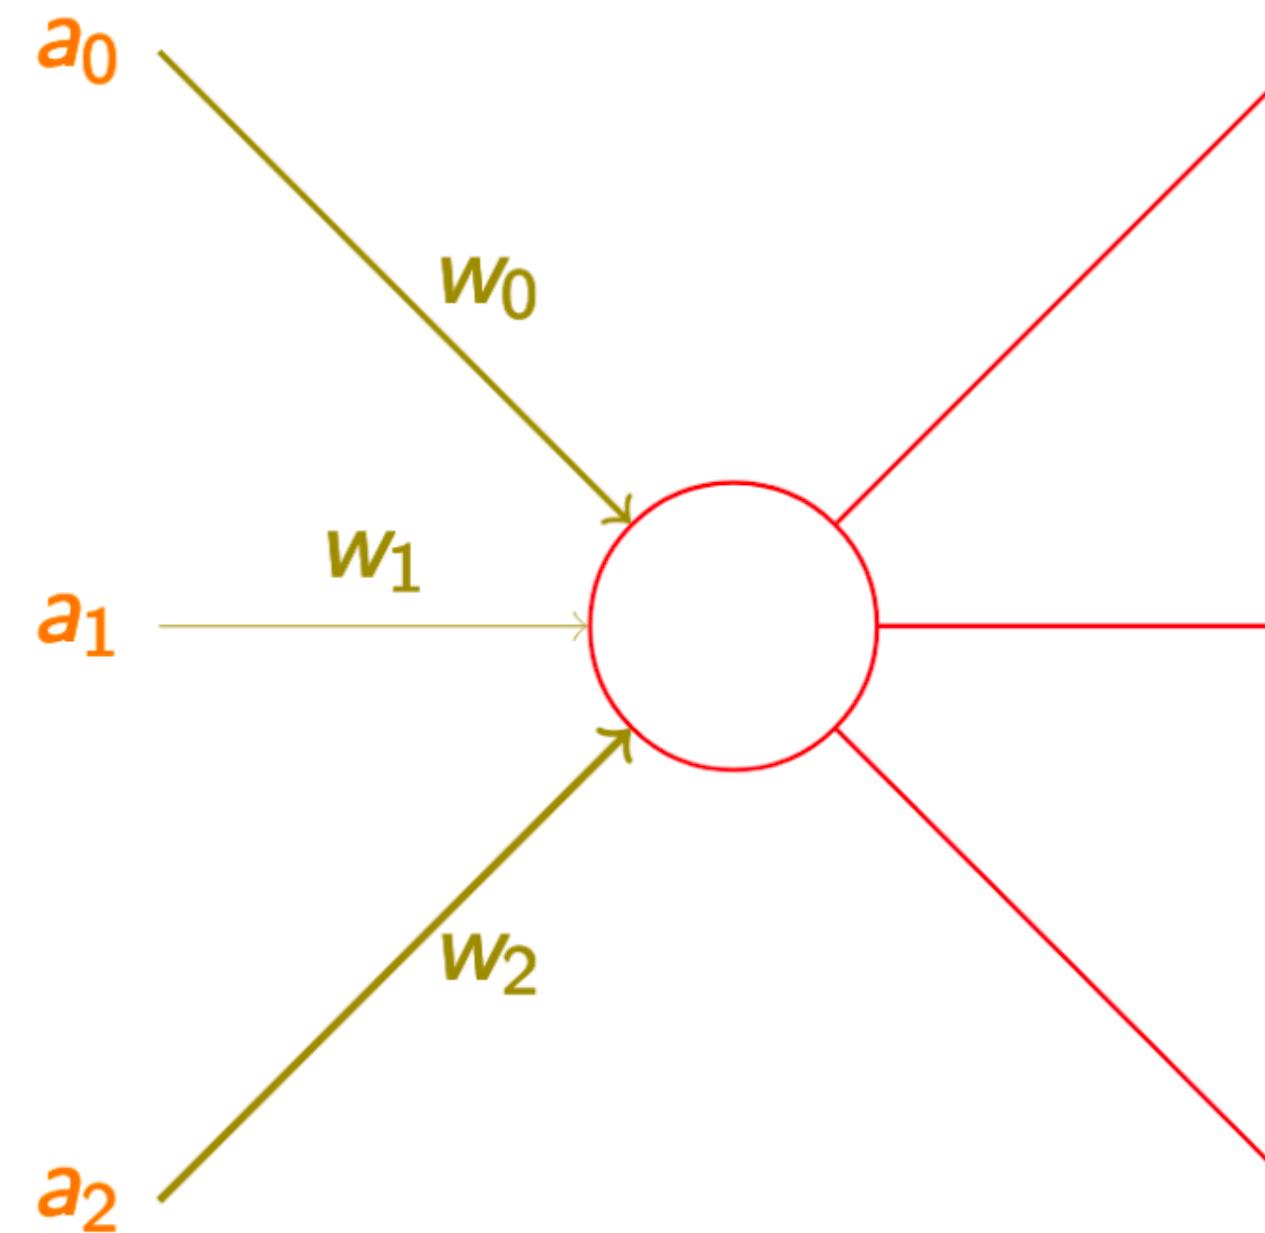
Fall 2024

# Plan for Today

- Last time:
  - Computational graph abstraction
  - Backpropagation algorithm
- Today: intro to **feed-forward neural networks**
  - Basic computation + expressive power
  - Multilayer perceptrons
  - Mini-batches
  - Hyper-parameters and regularization

# Computation: Basic Example

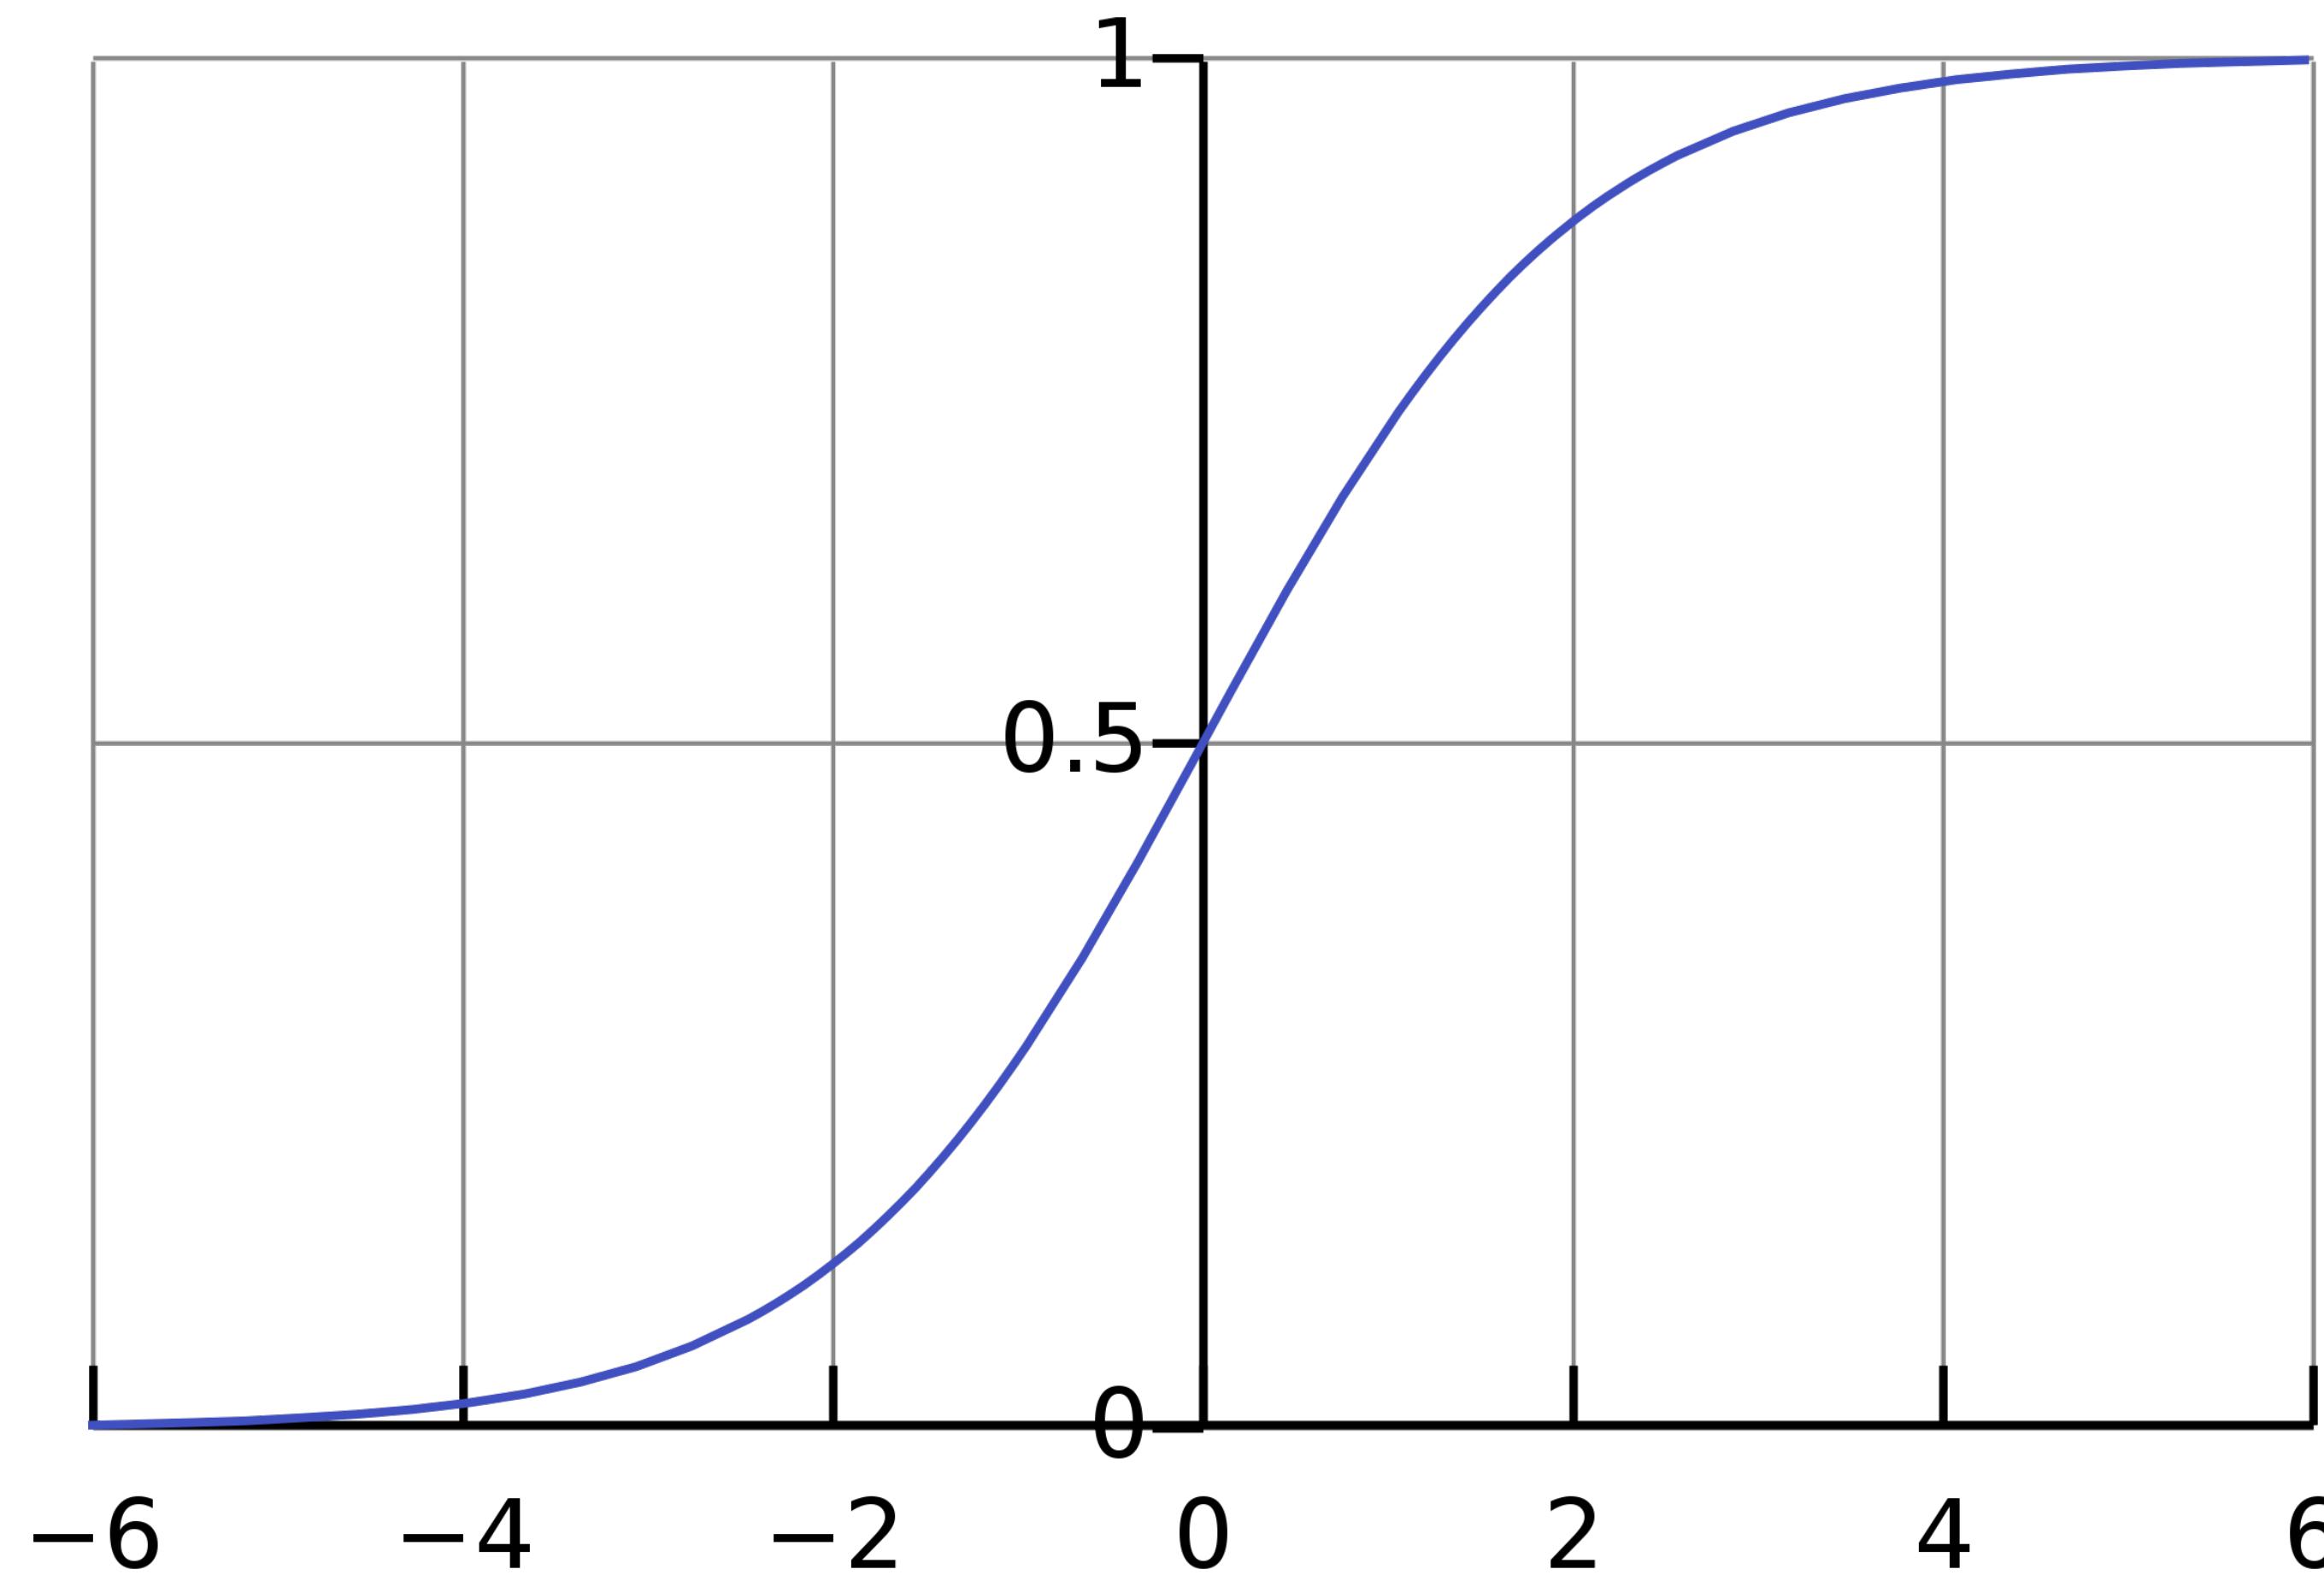
# Artificial Neuron



$$a = f(a_0 \cdot w_0 + a_1 \cdot w_1 + a_2 \cdot w_2)$$

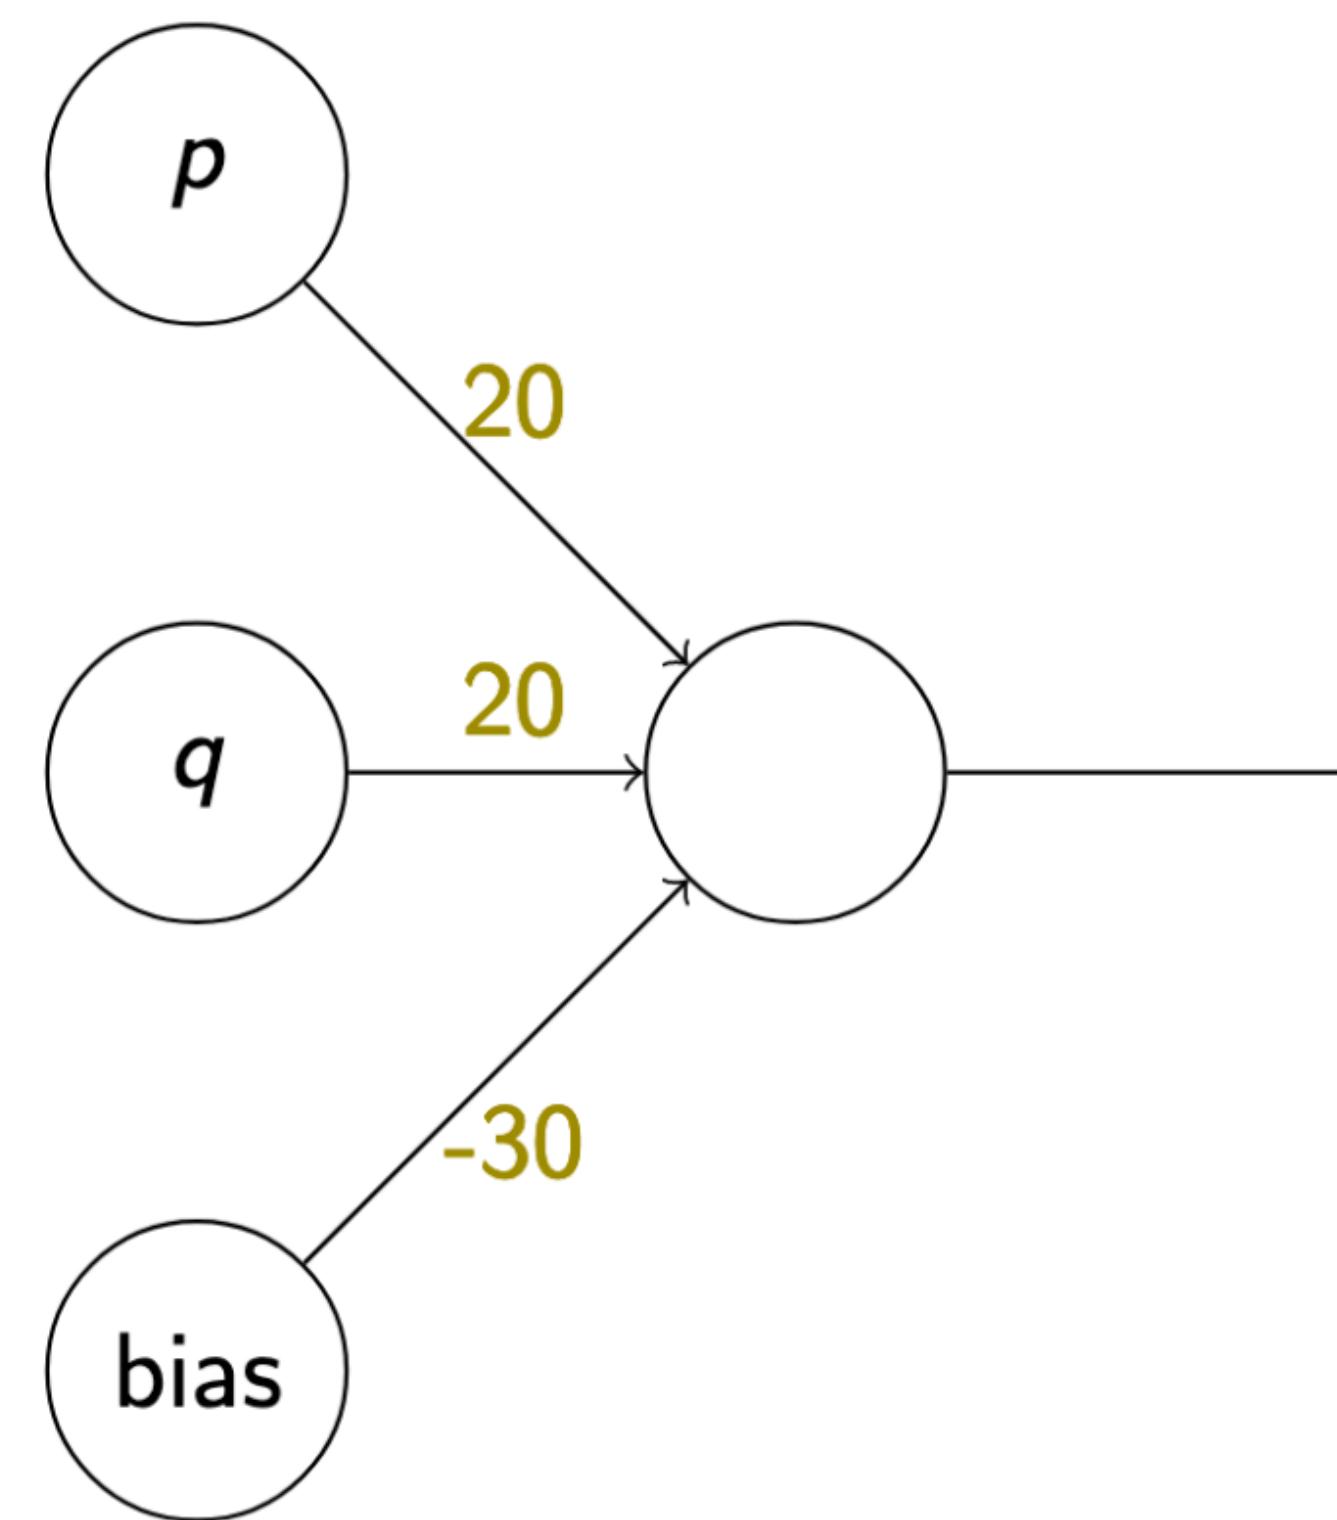
<https://github.com/shanest/nn-tutorial>

# Activation Function: Sigmoid



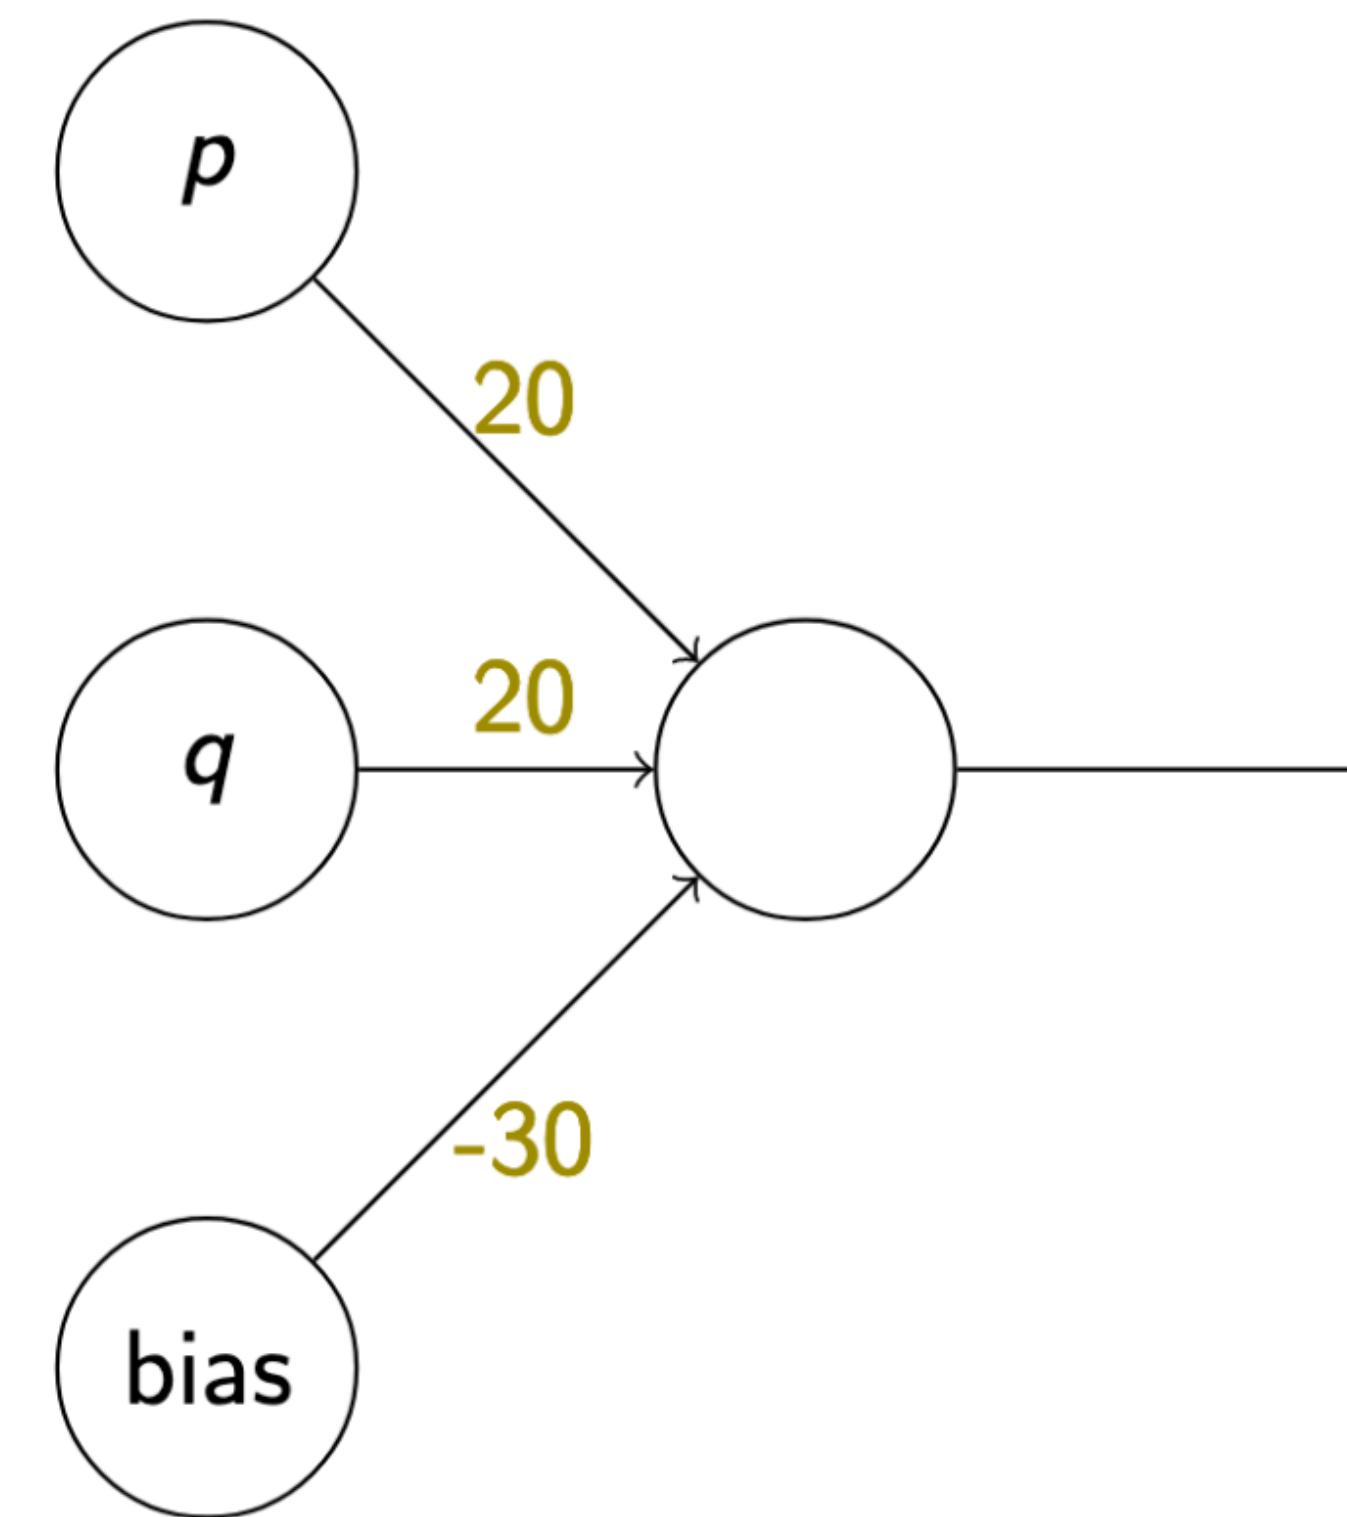
$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

# Computing a Boolean function



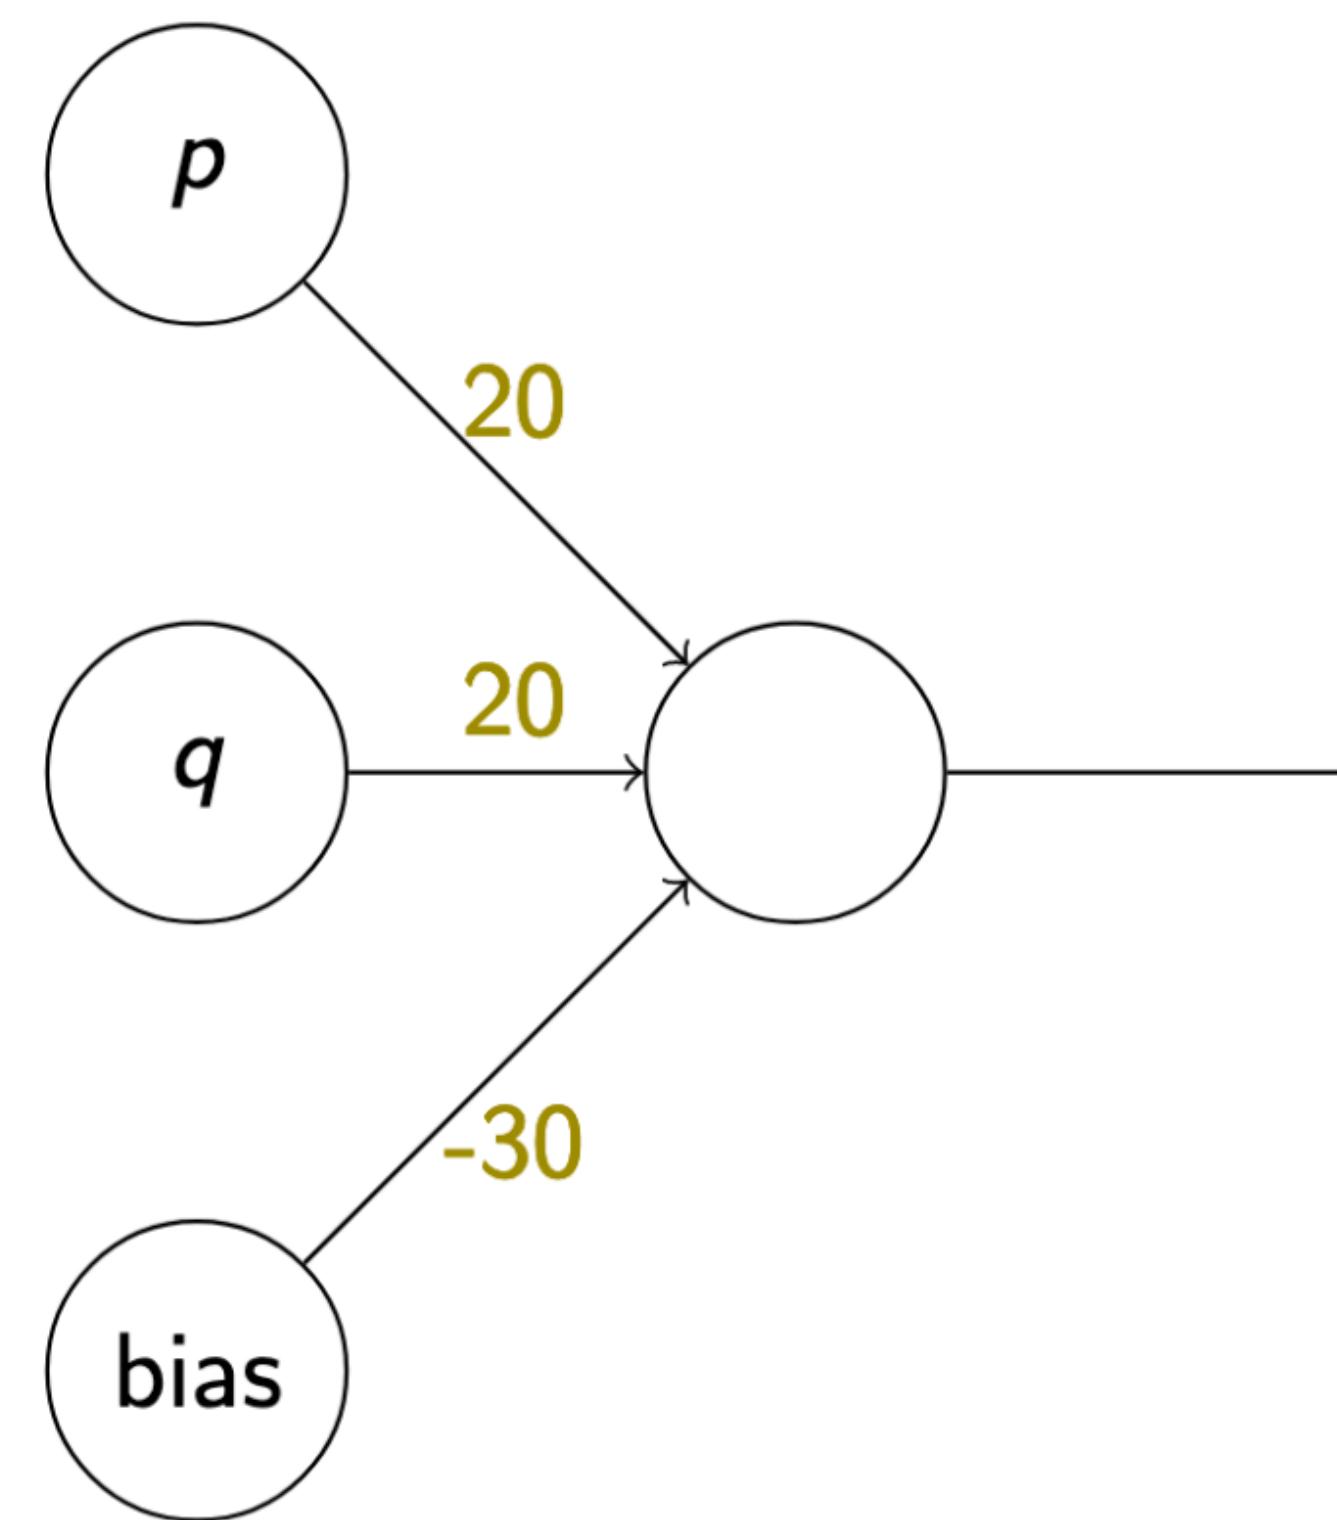
# Computing a Boolean function

p	q	a
---	---	---



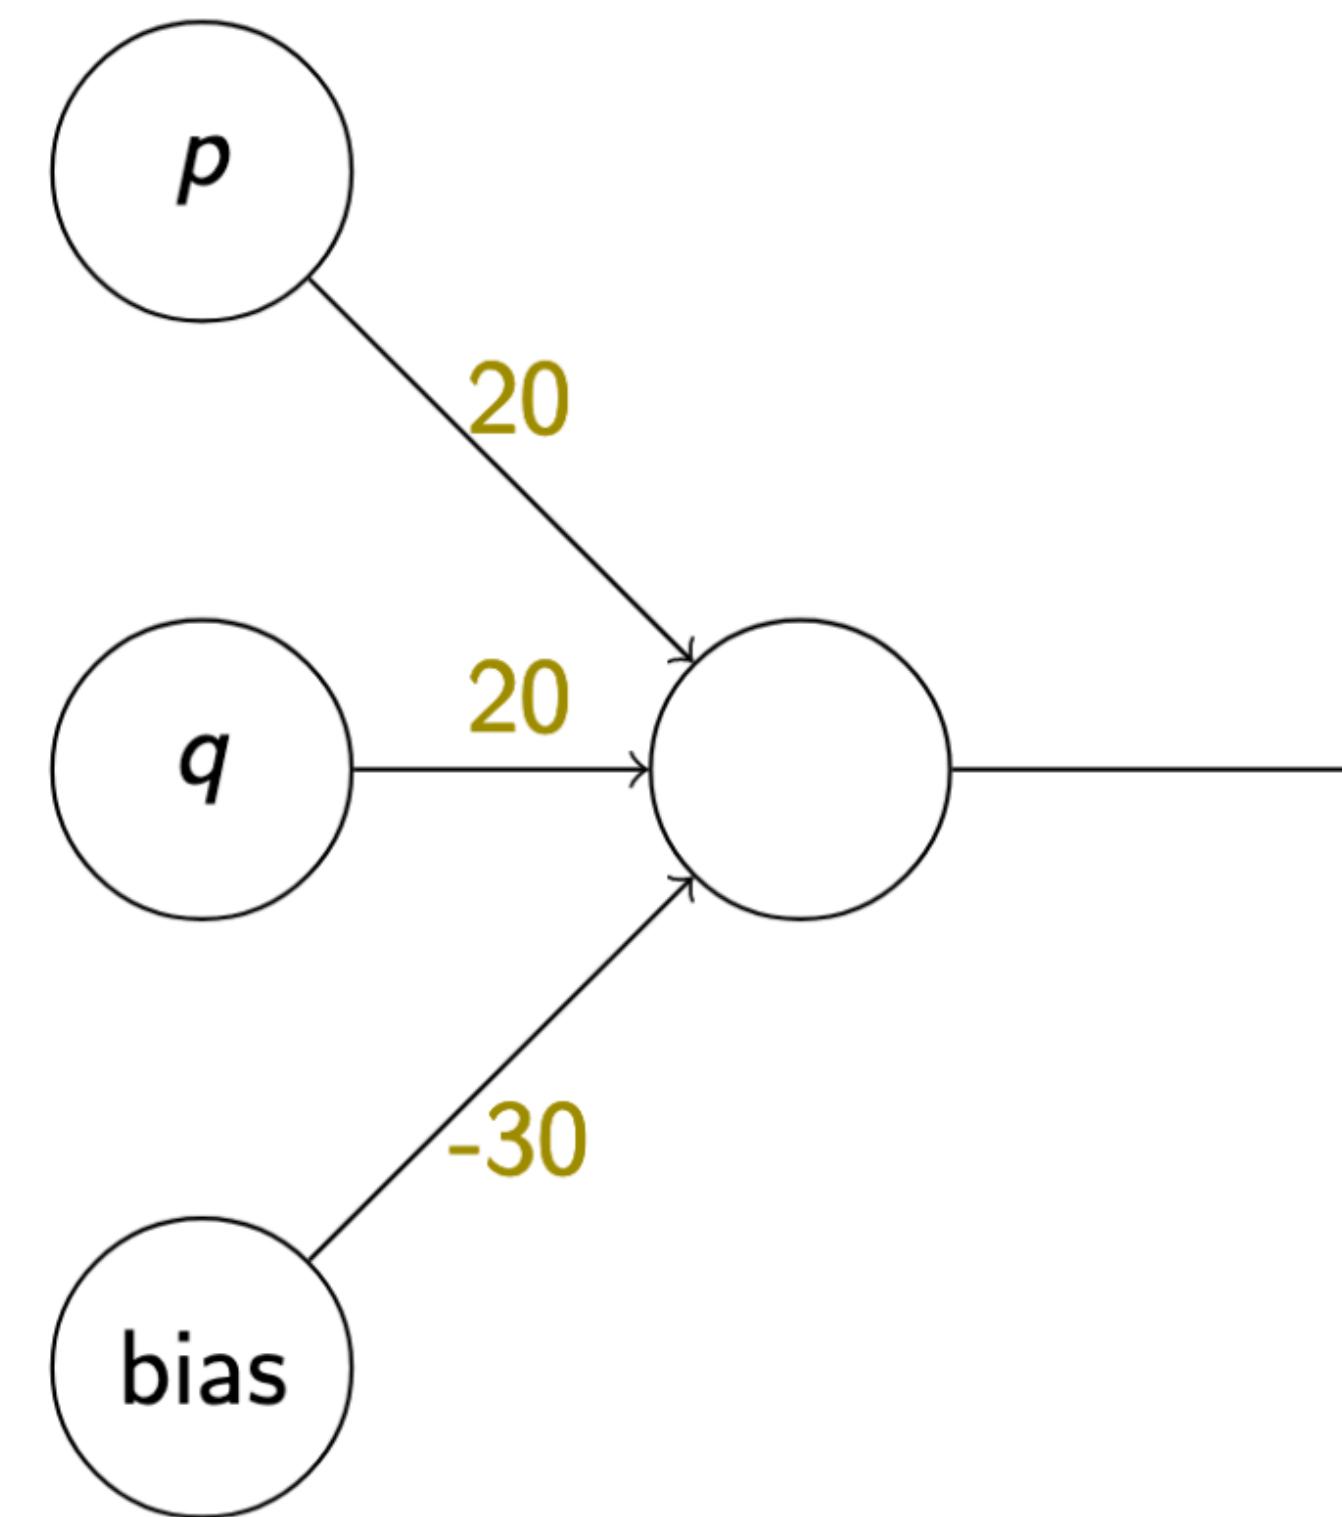
# Computing a Boolean function

p	q	a
1	1	1



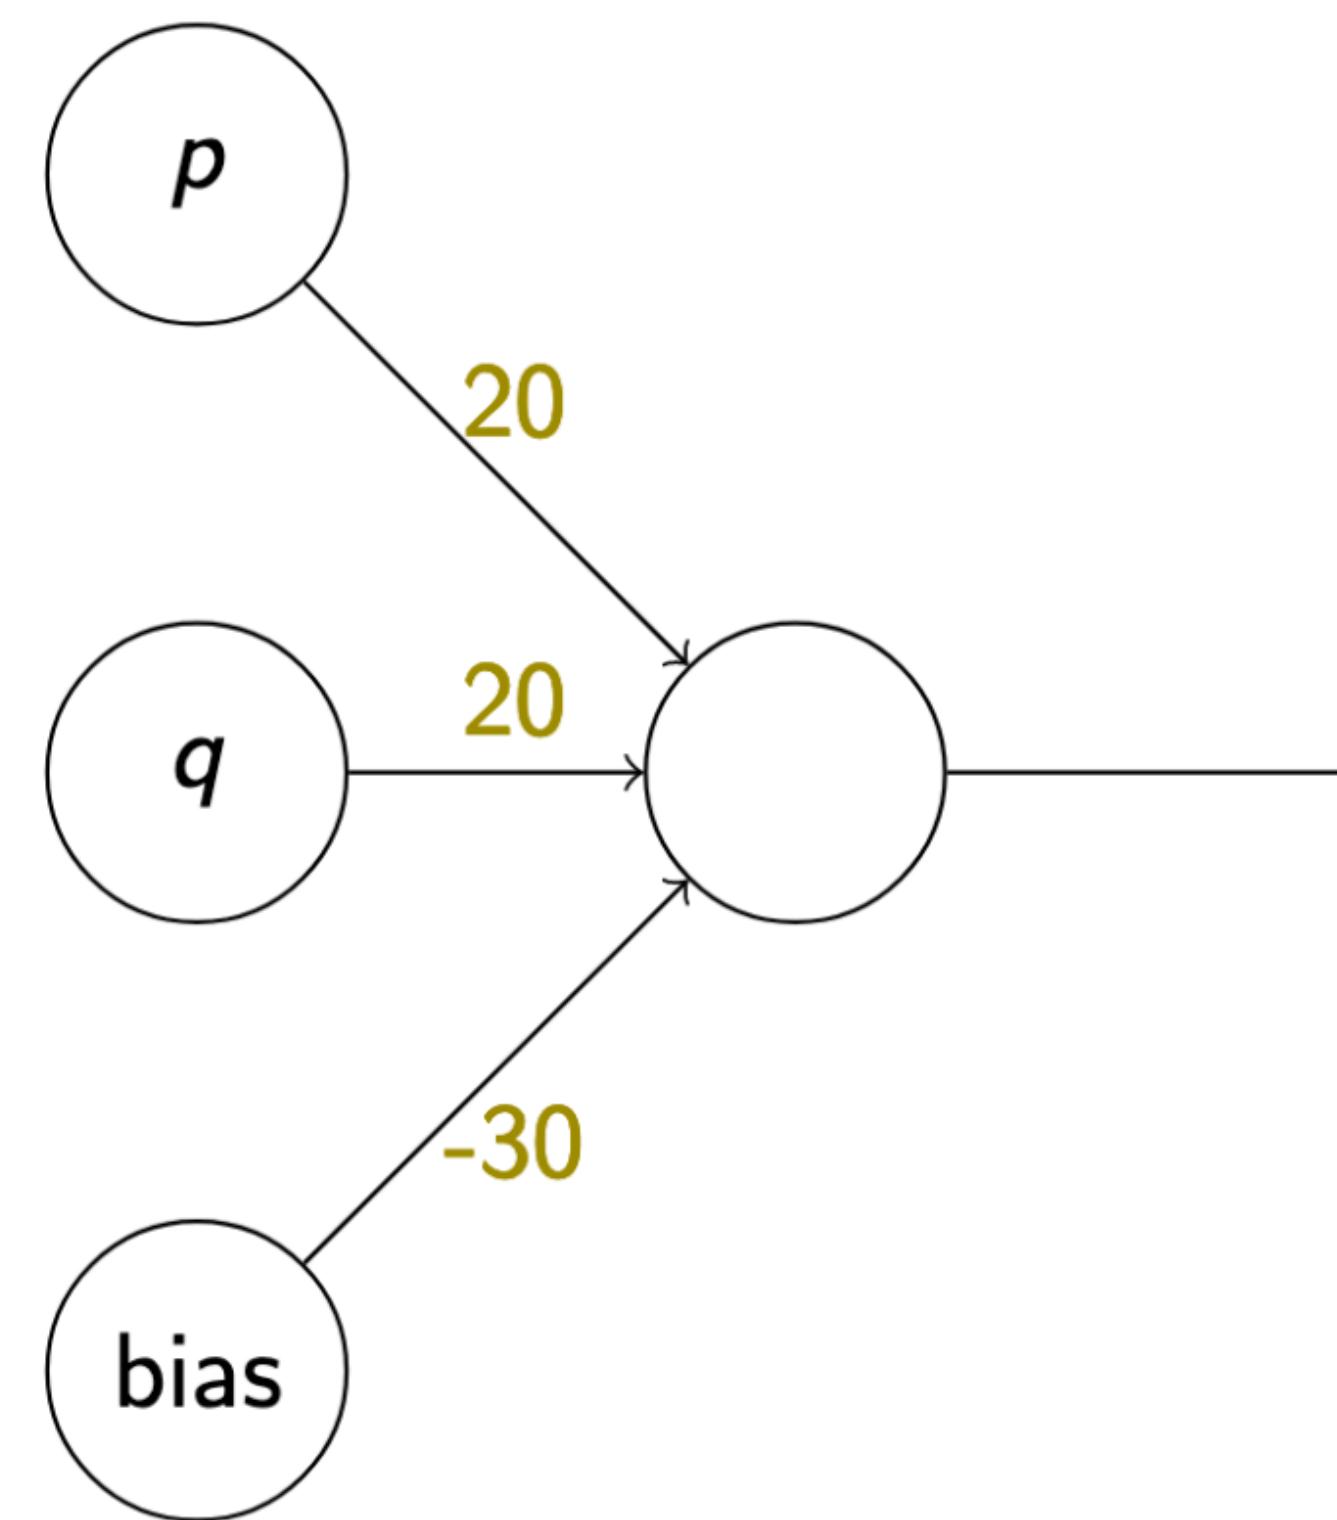
# Computing a Boolean function

p	q	a
1	1	1
1	0	0



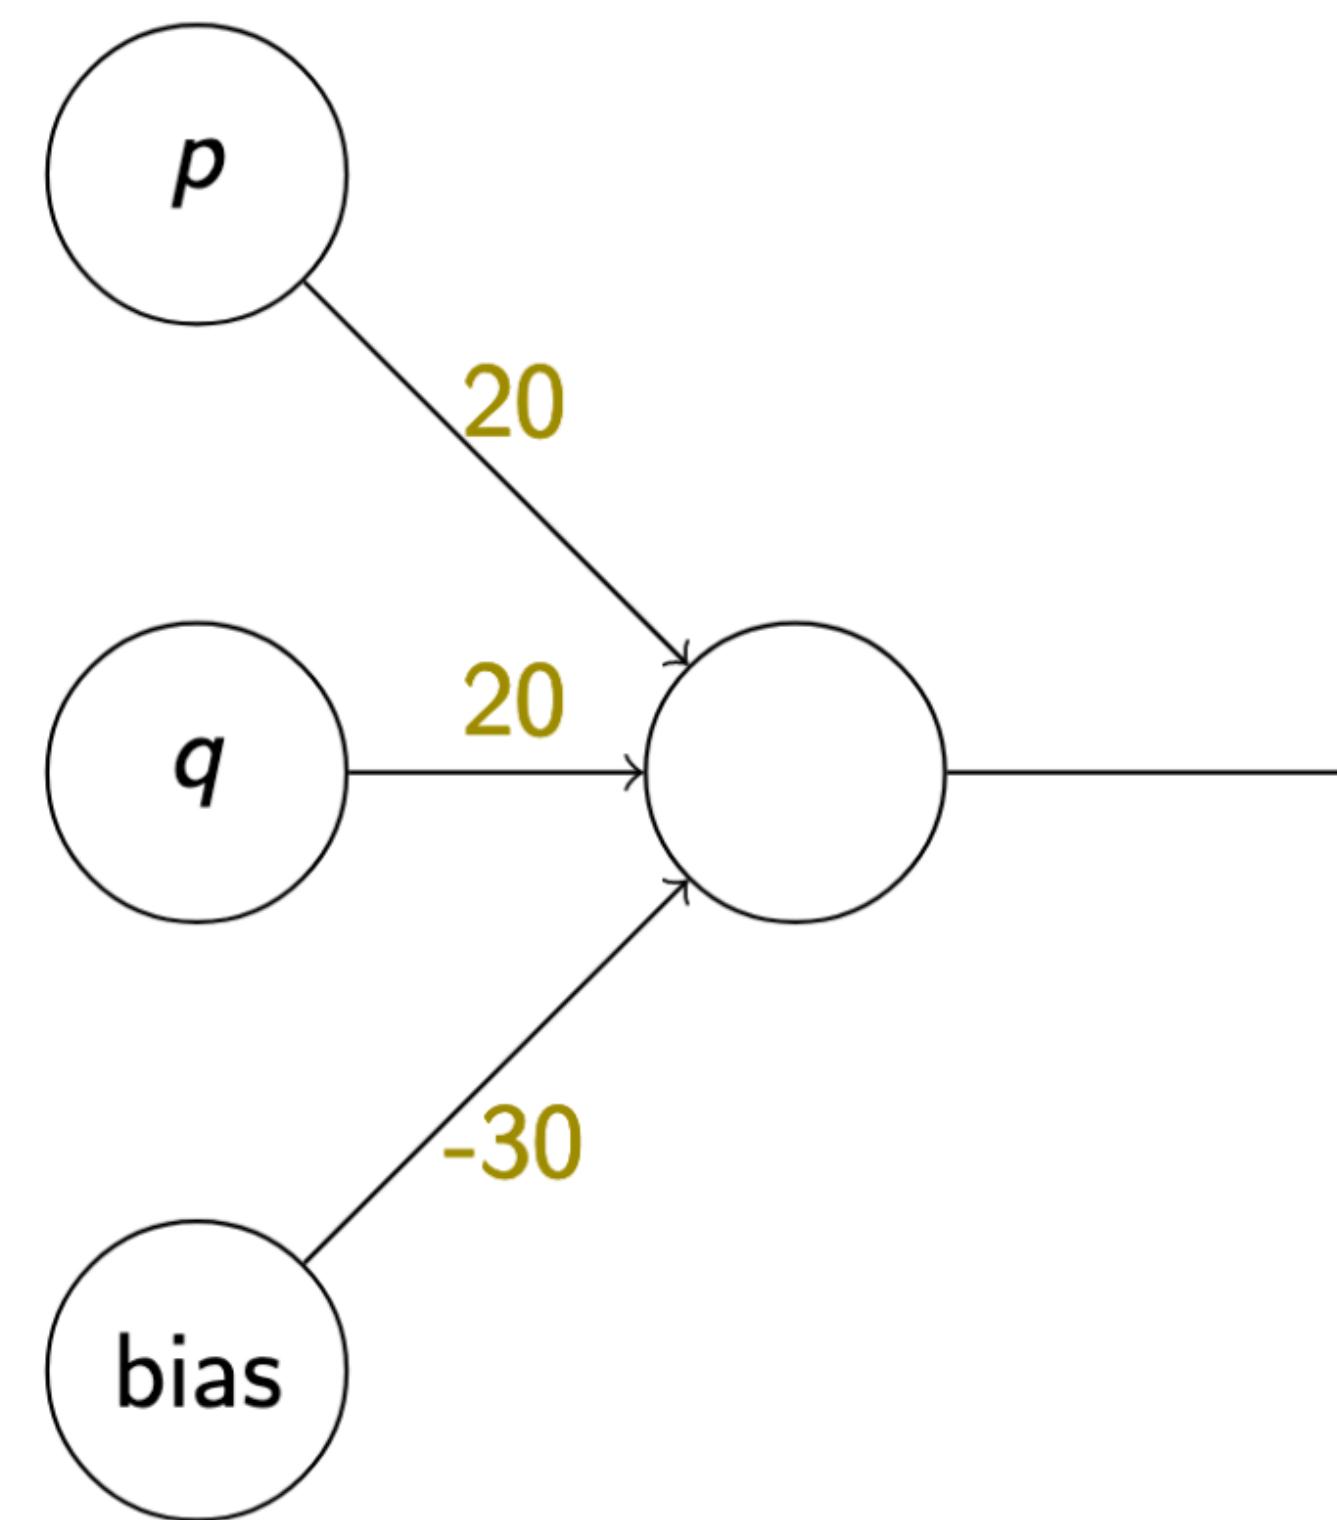
# Computing a Boolean function

p	q	a
1	1	1
1	0	0
0	1	0

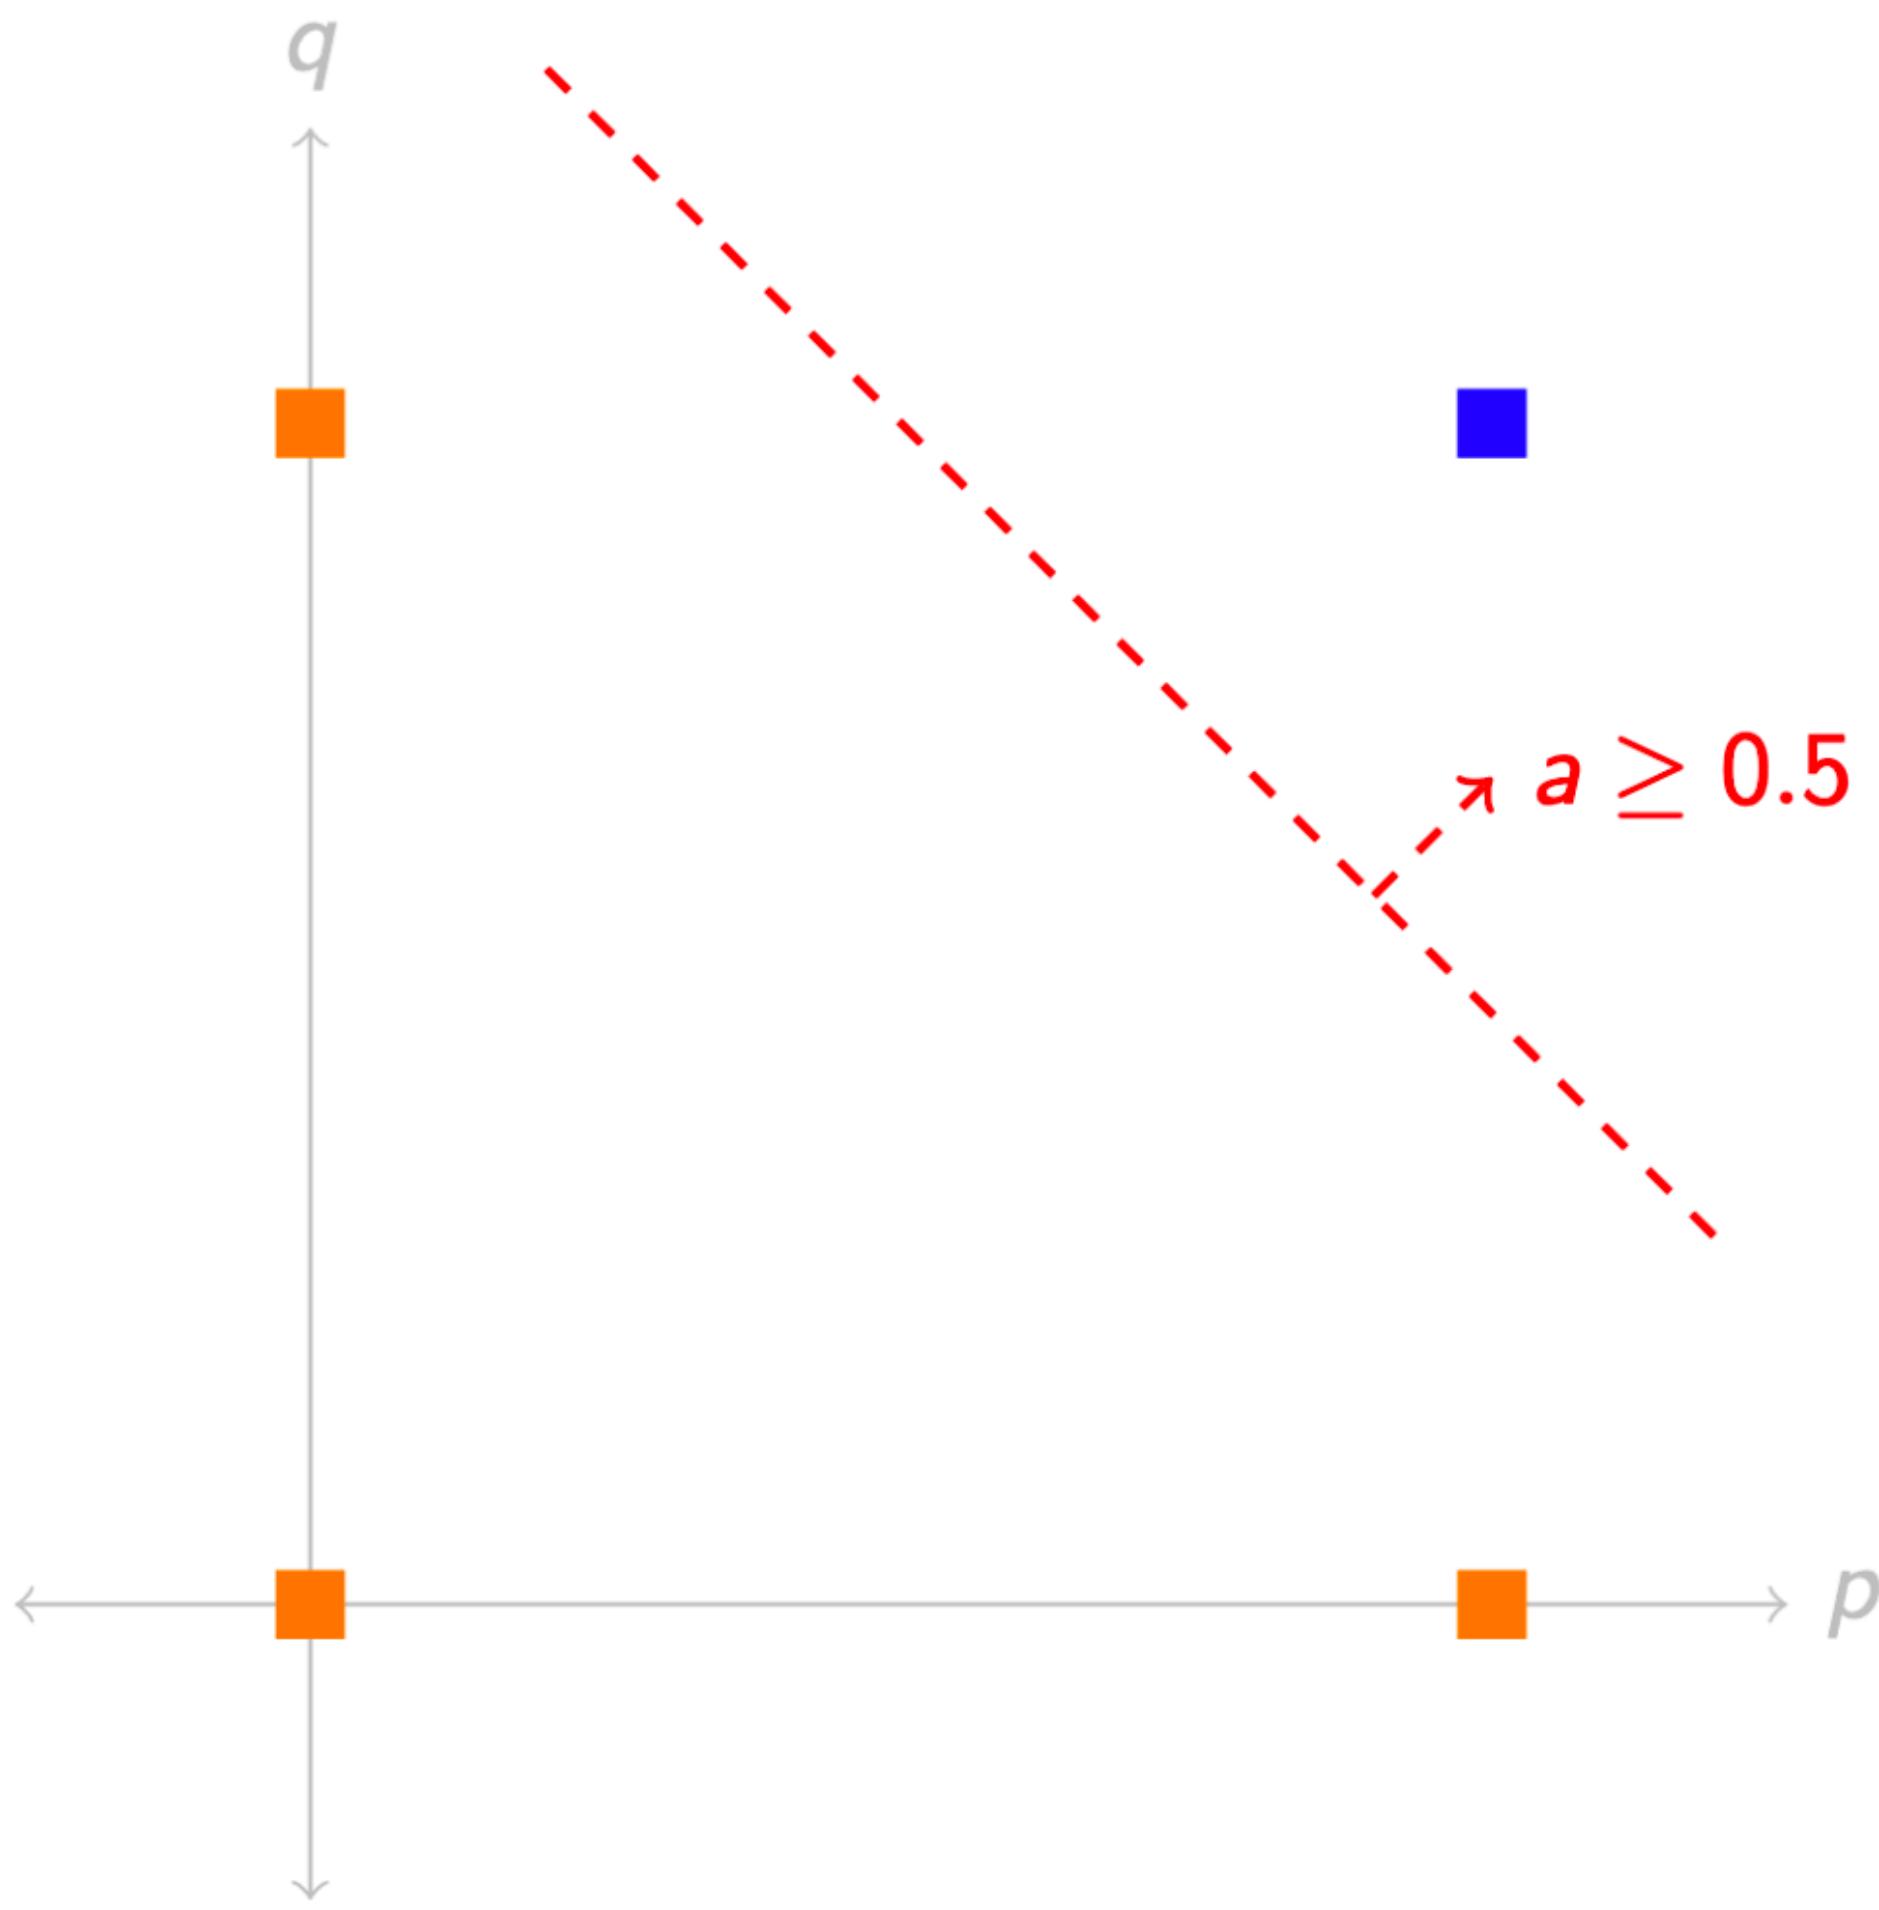


# Computing a Boolean function

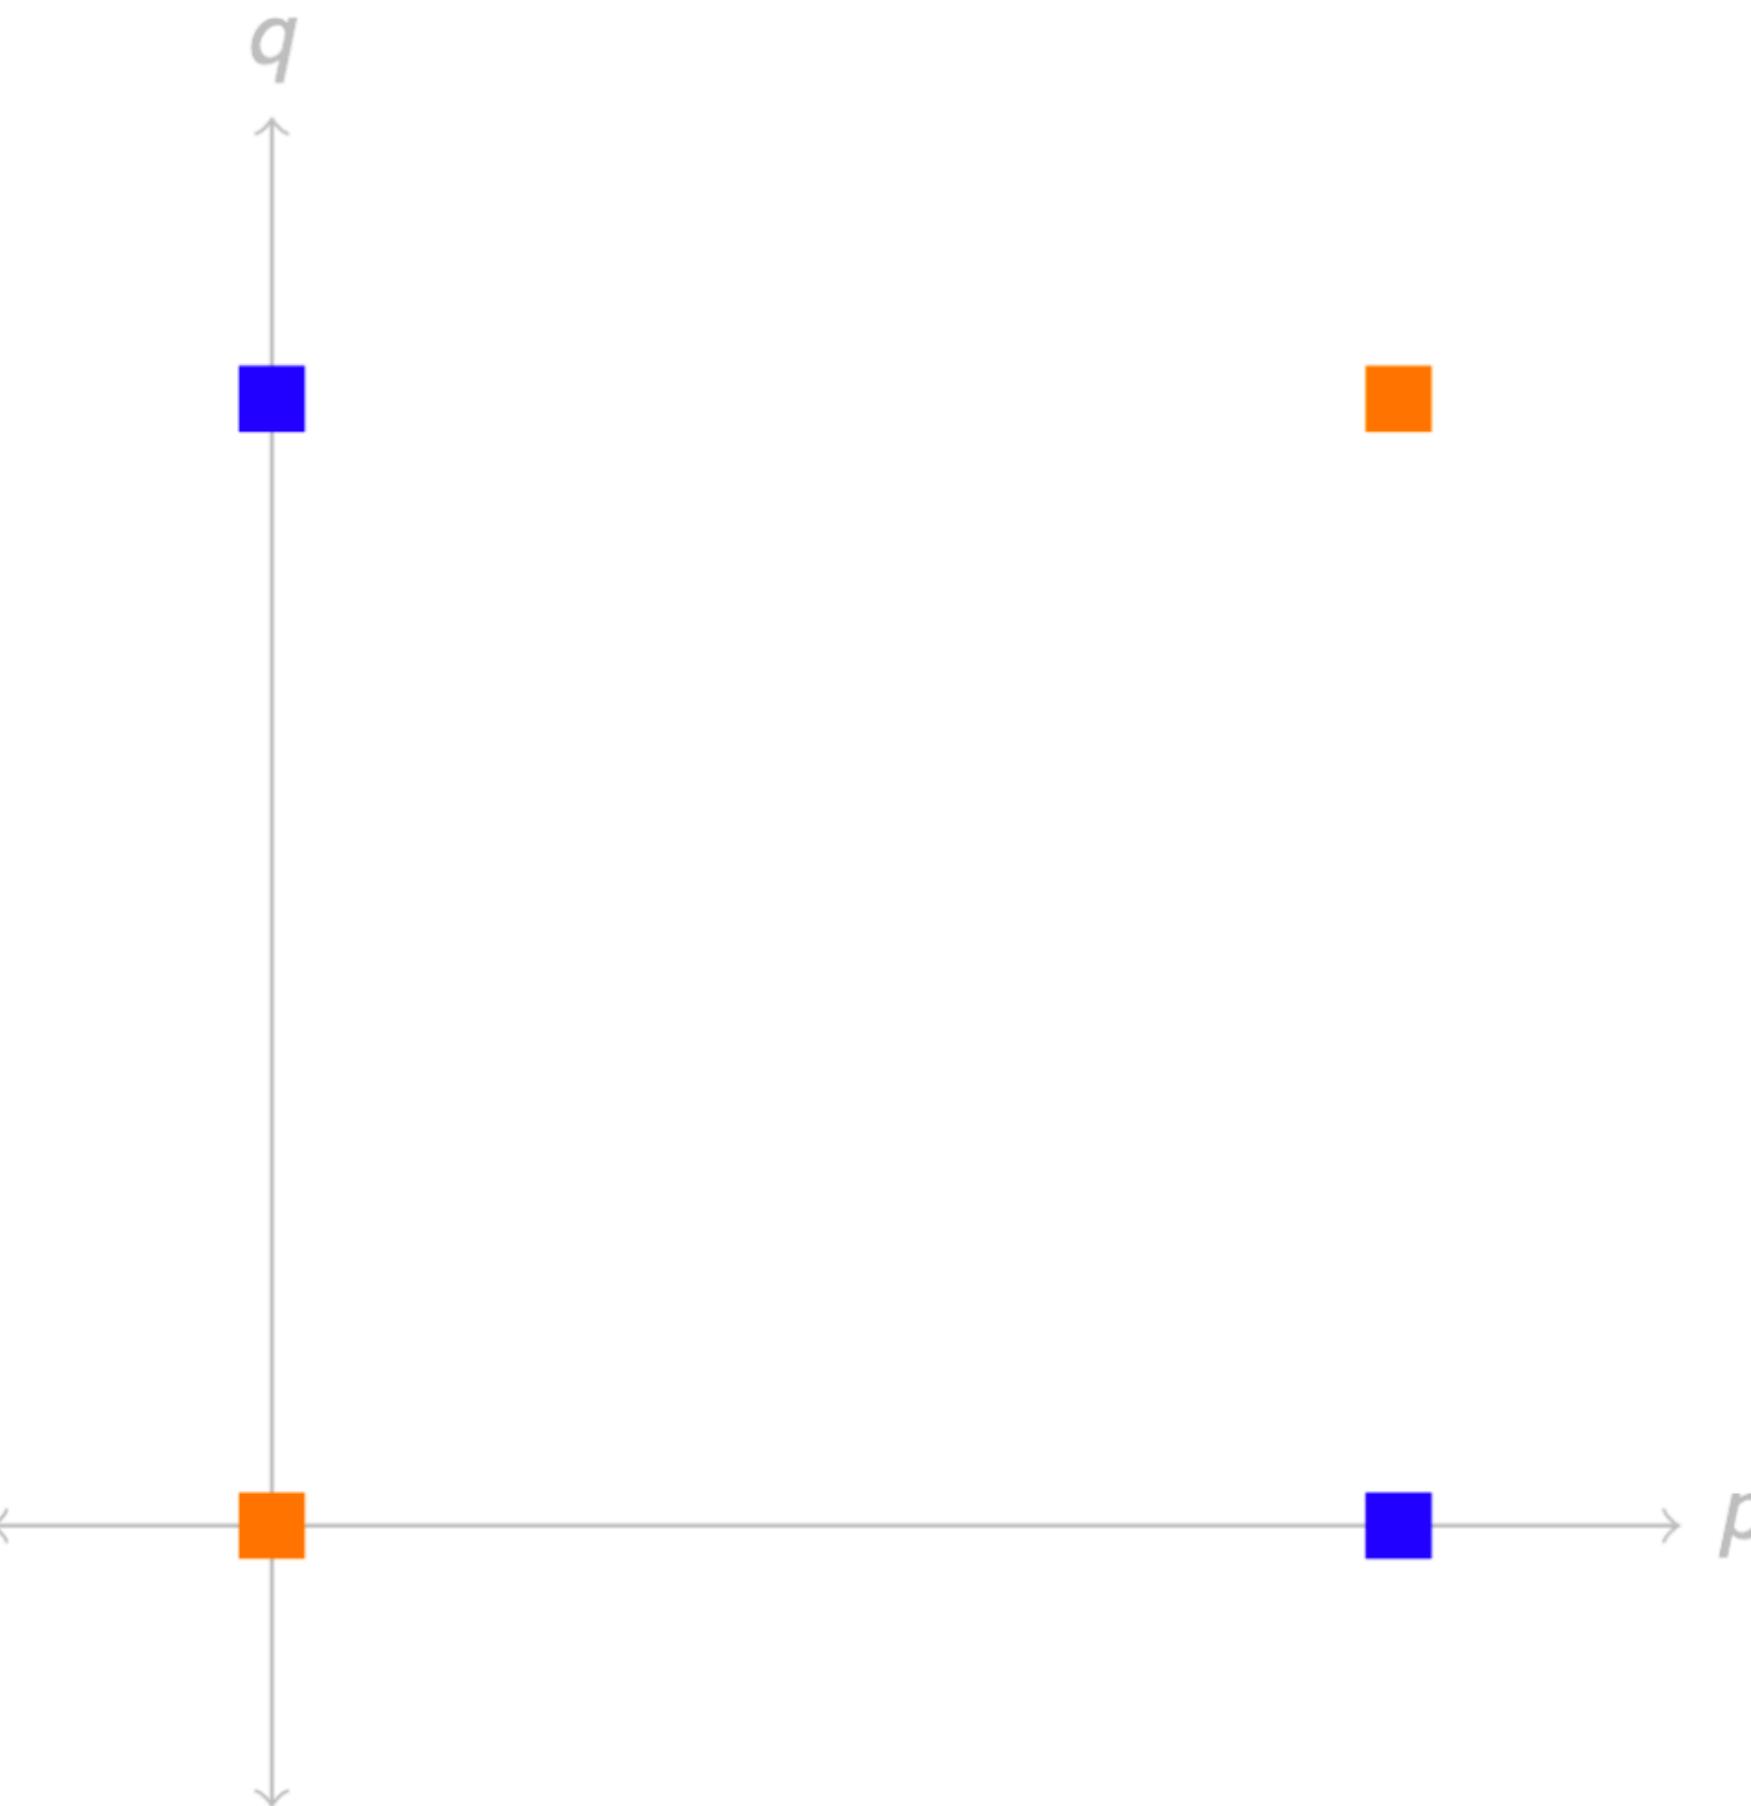
p	q	a
1	1	1
1	0	0
0	1	0
0	0	0



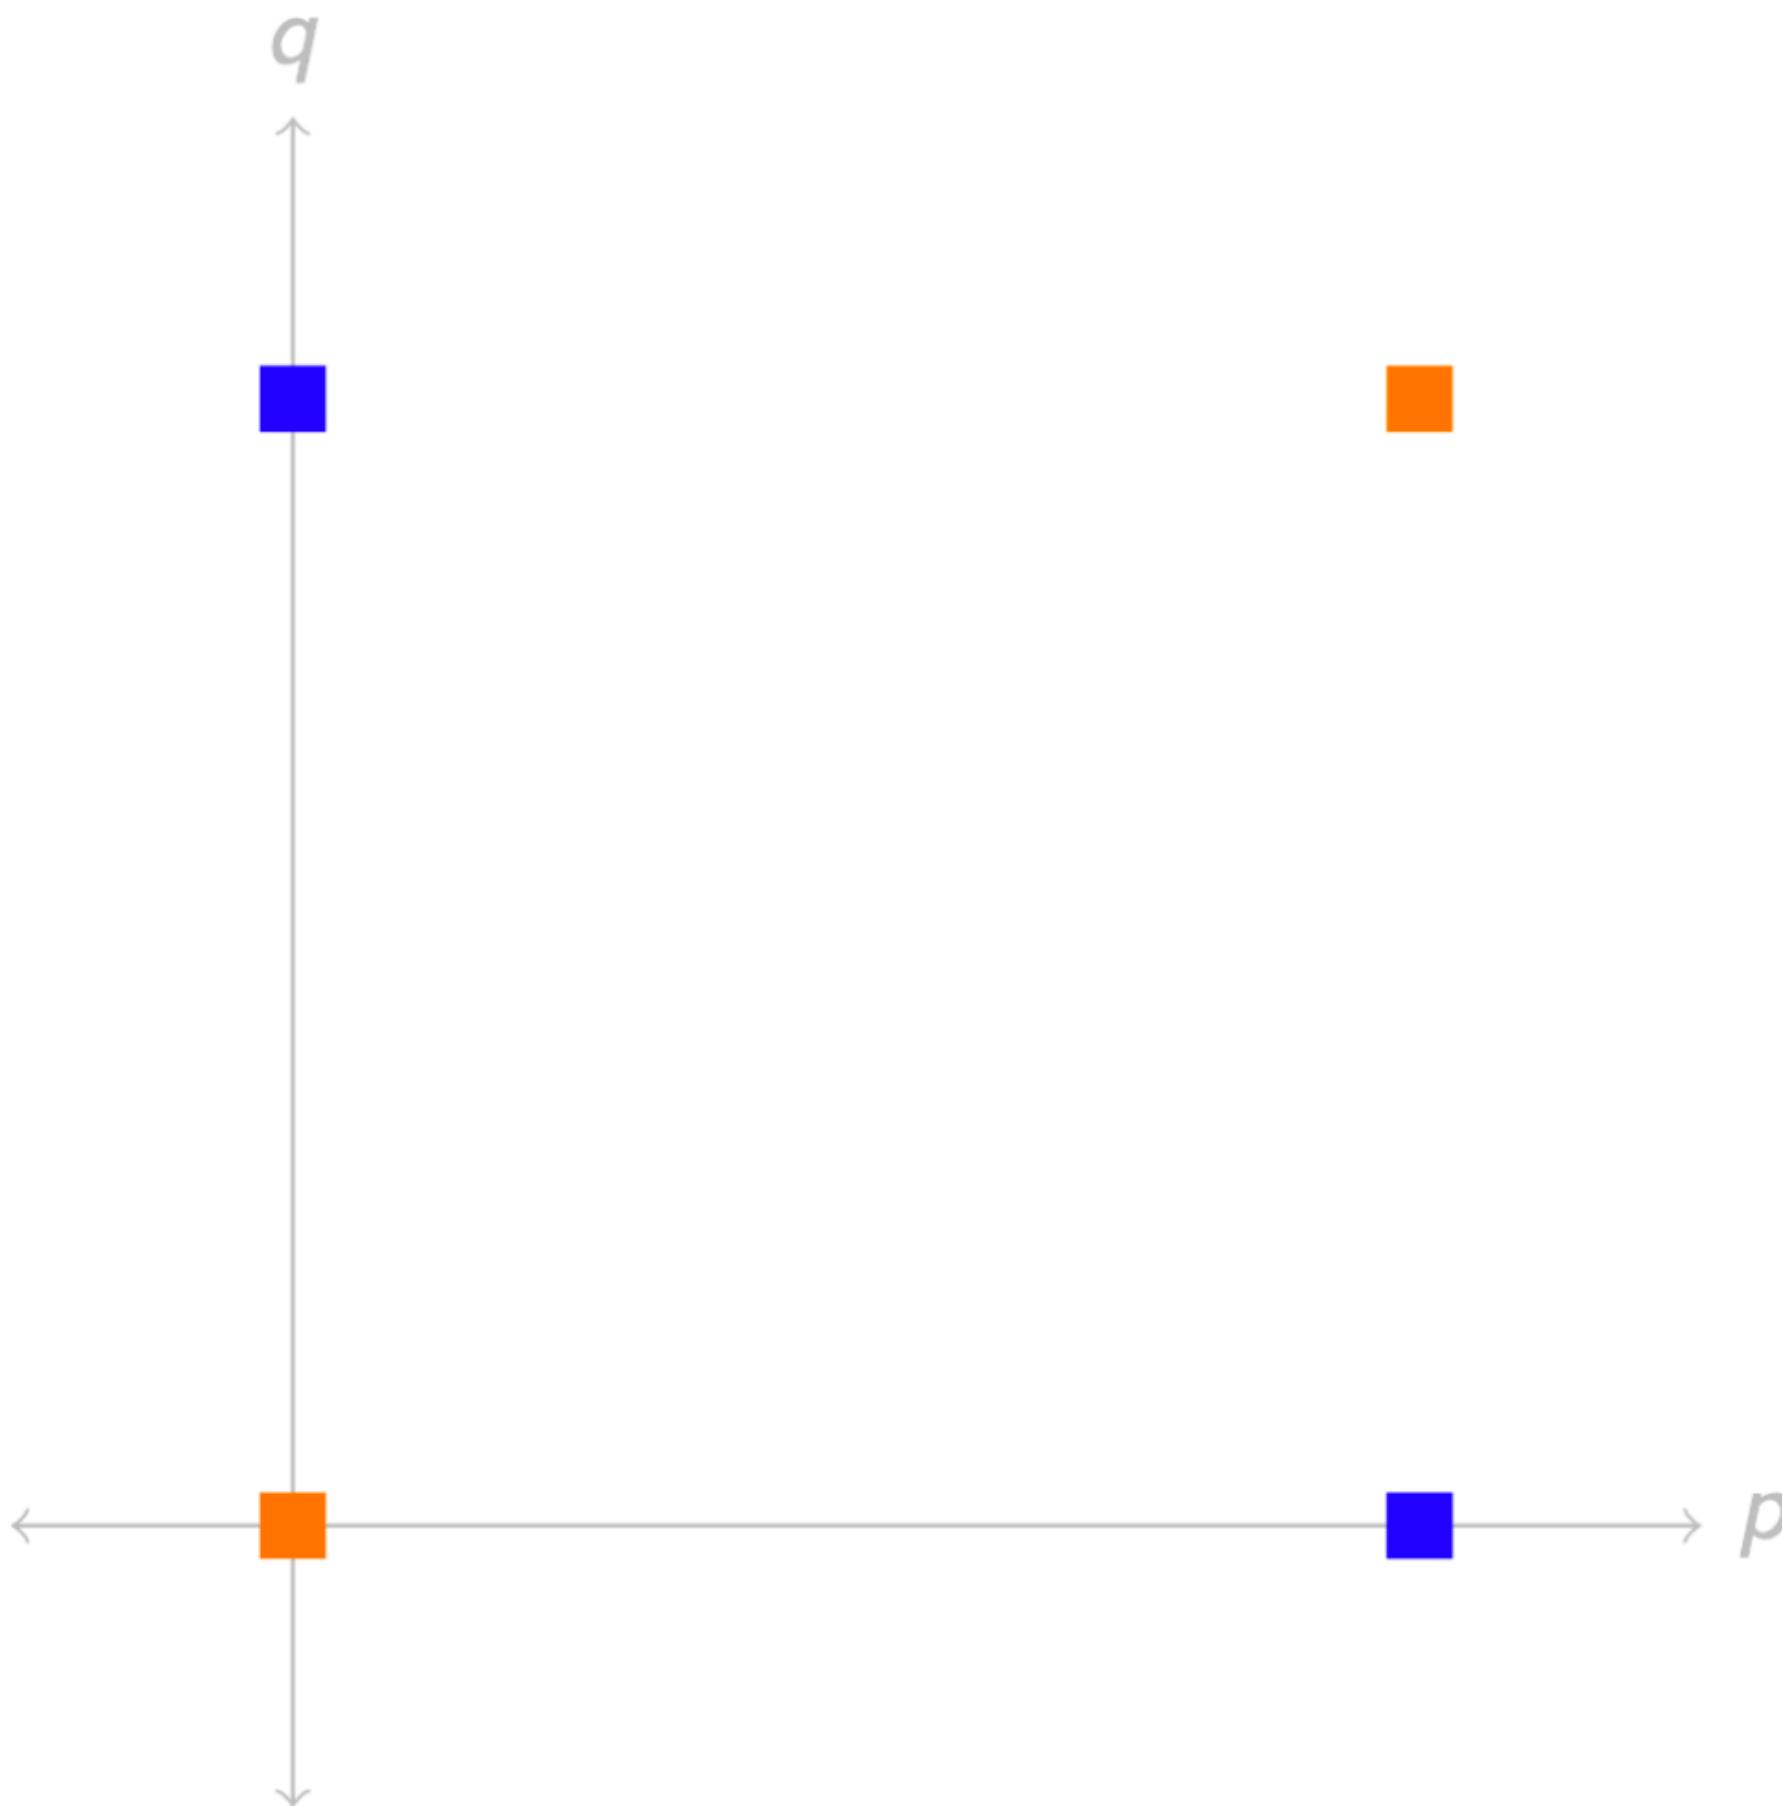
# Computing ‘and’



# The XOR problem

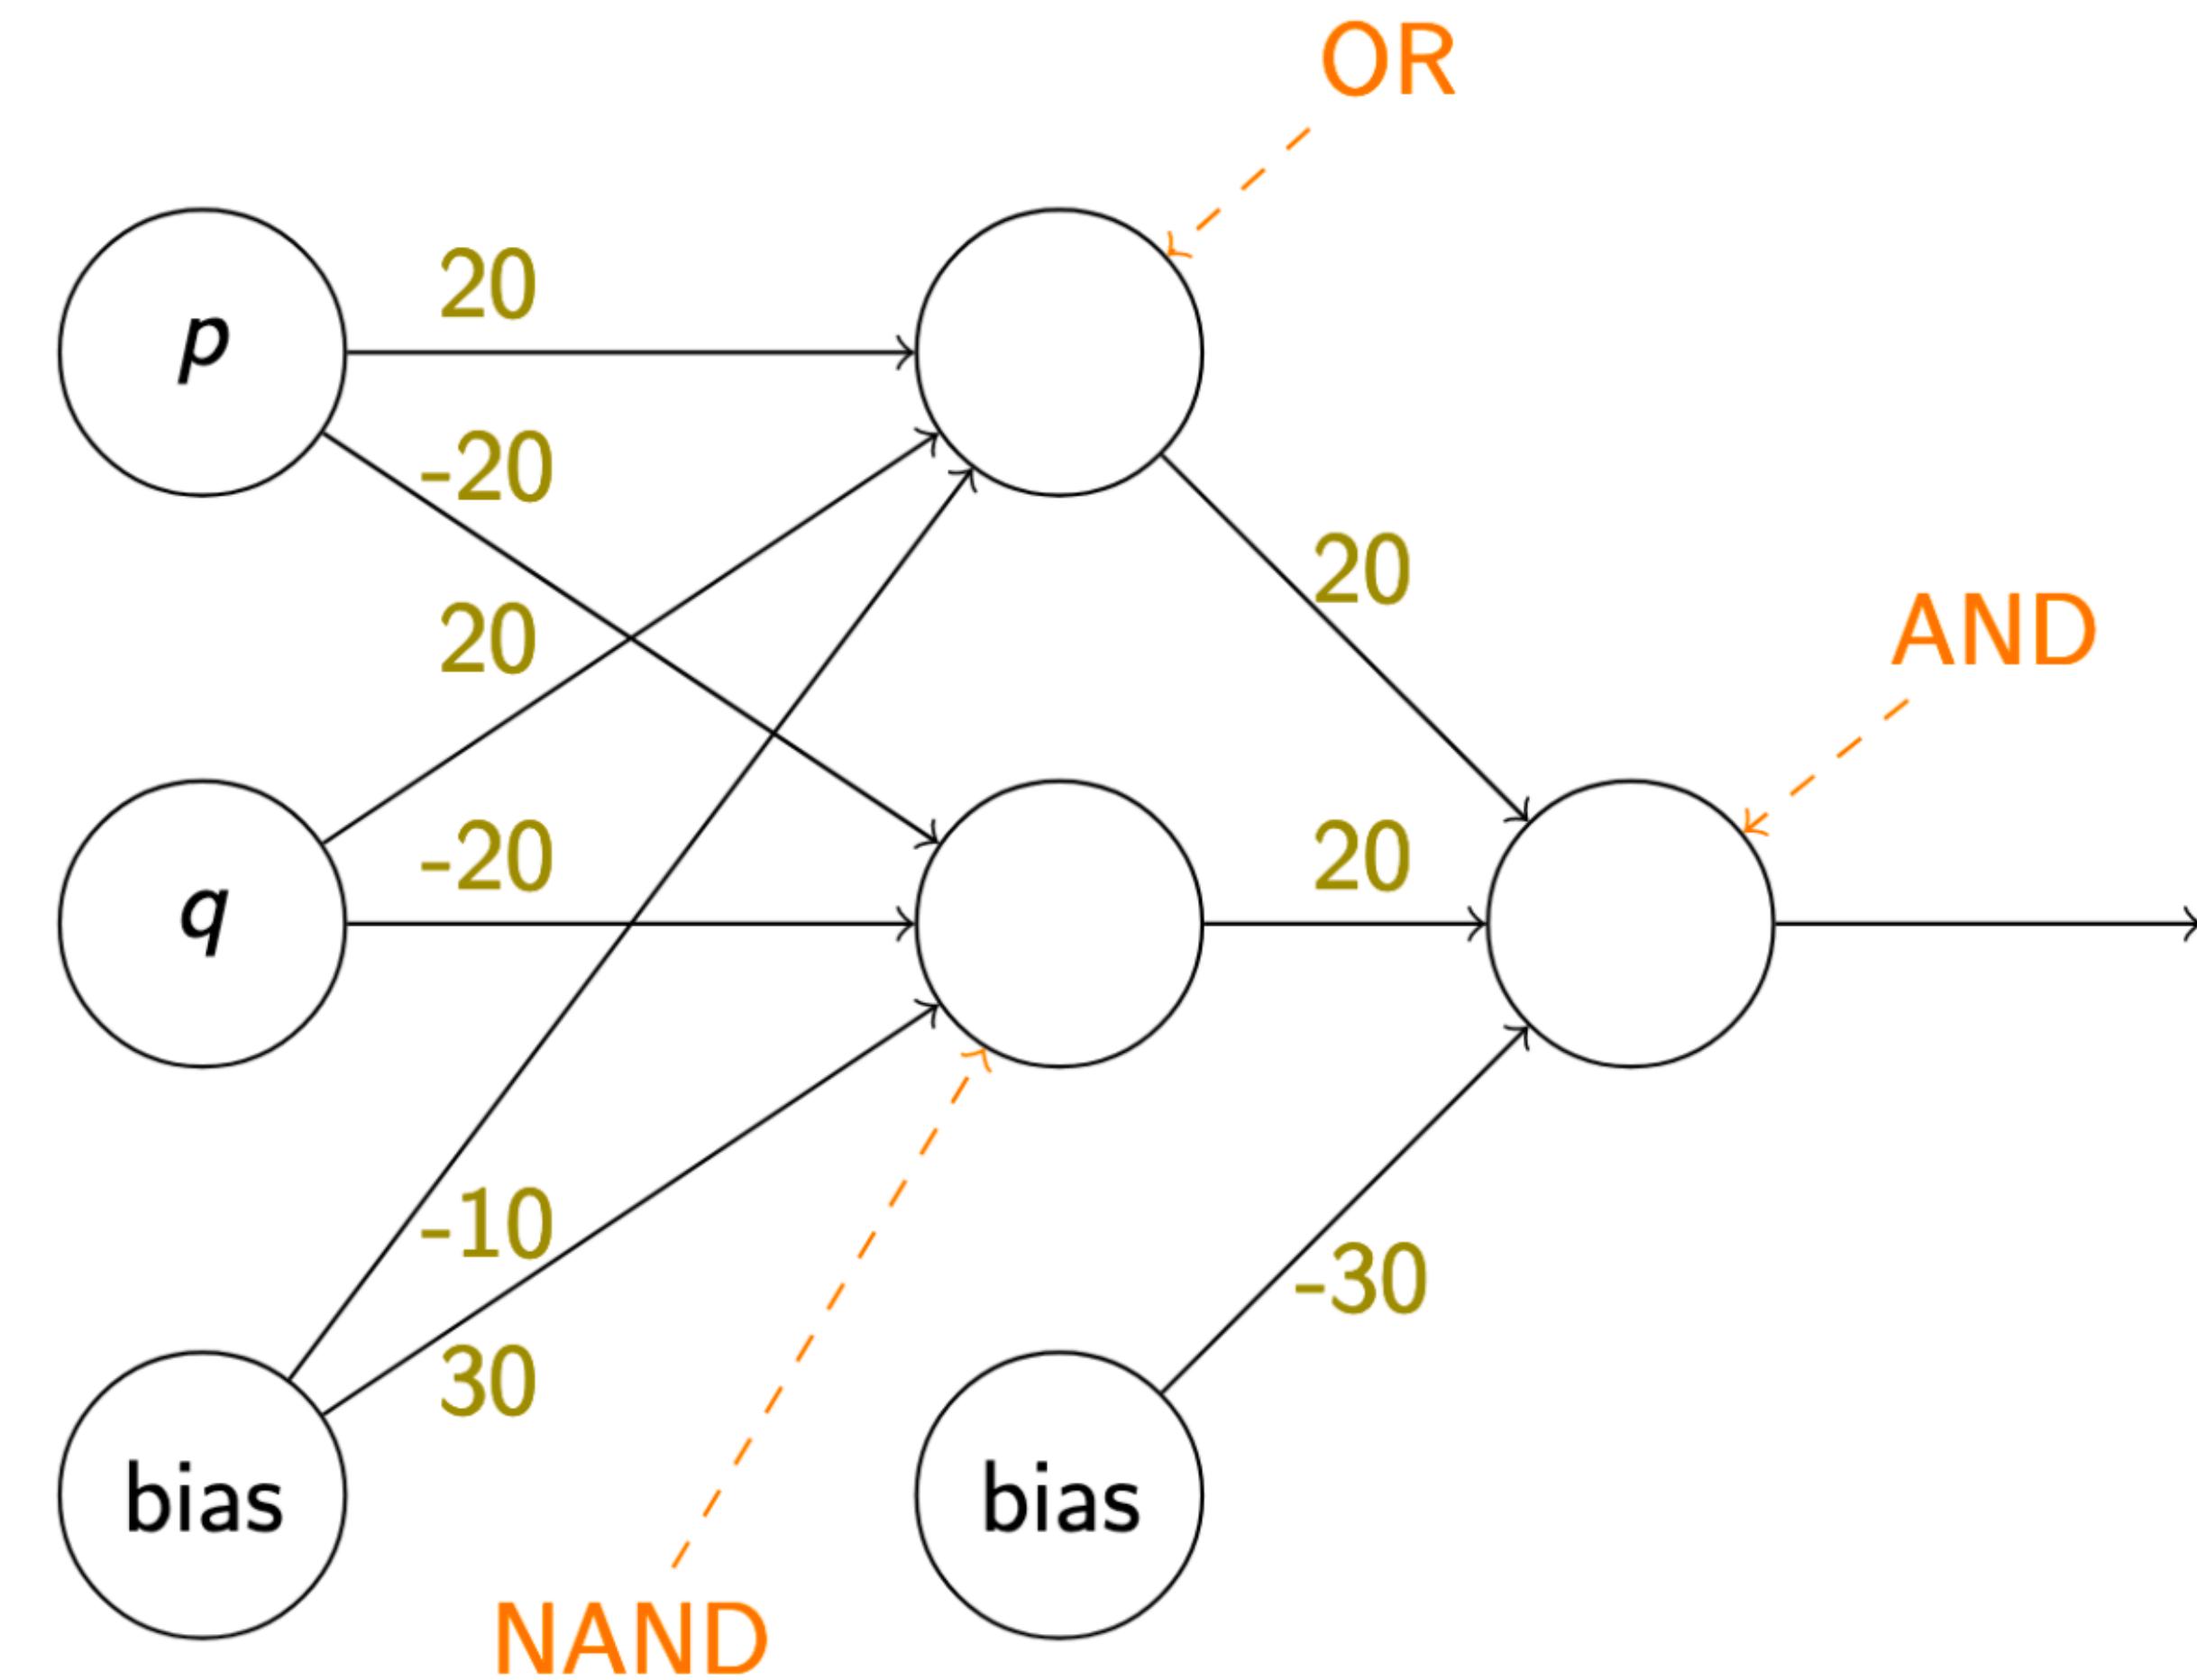


# The XOR problem

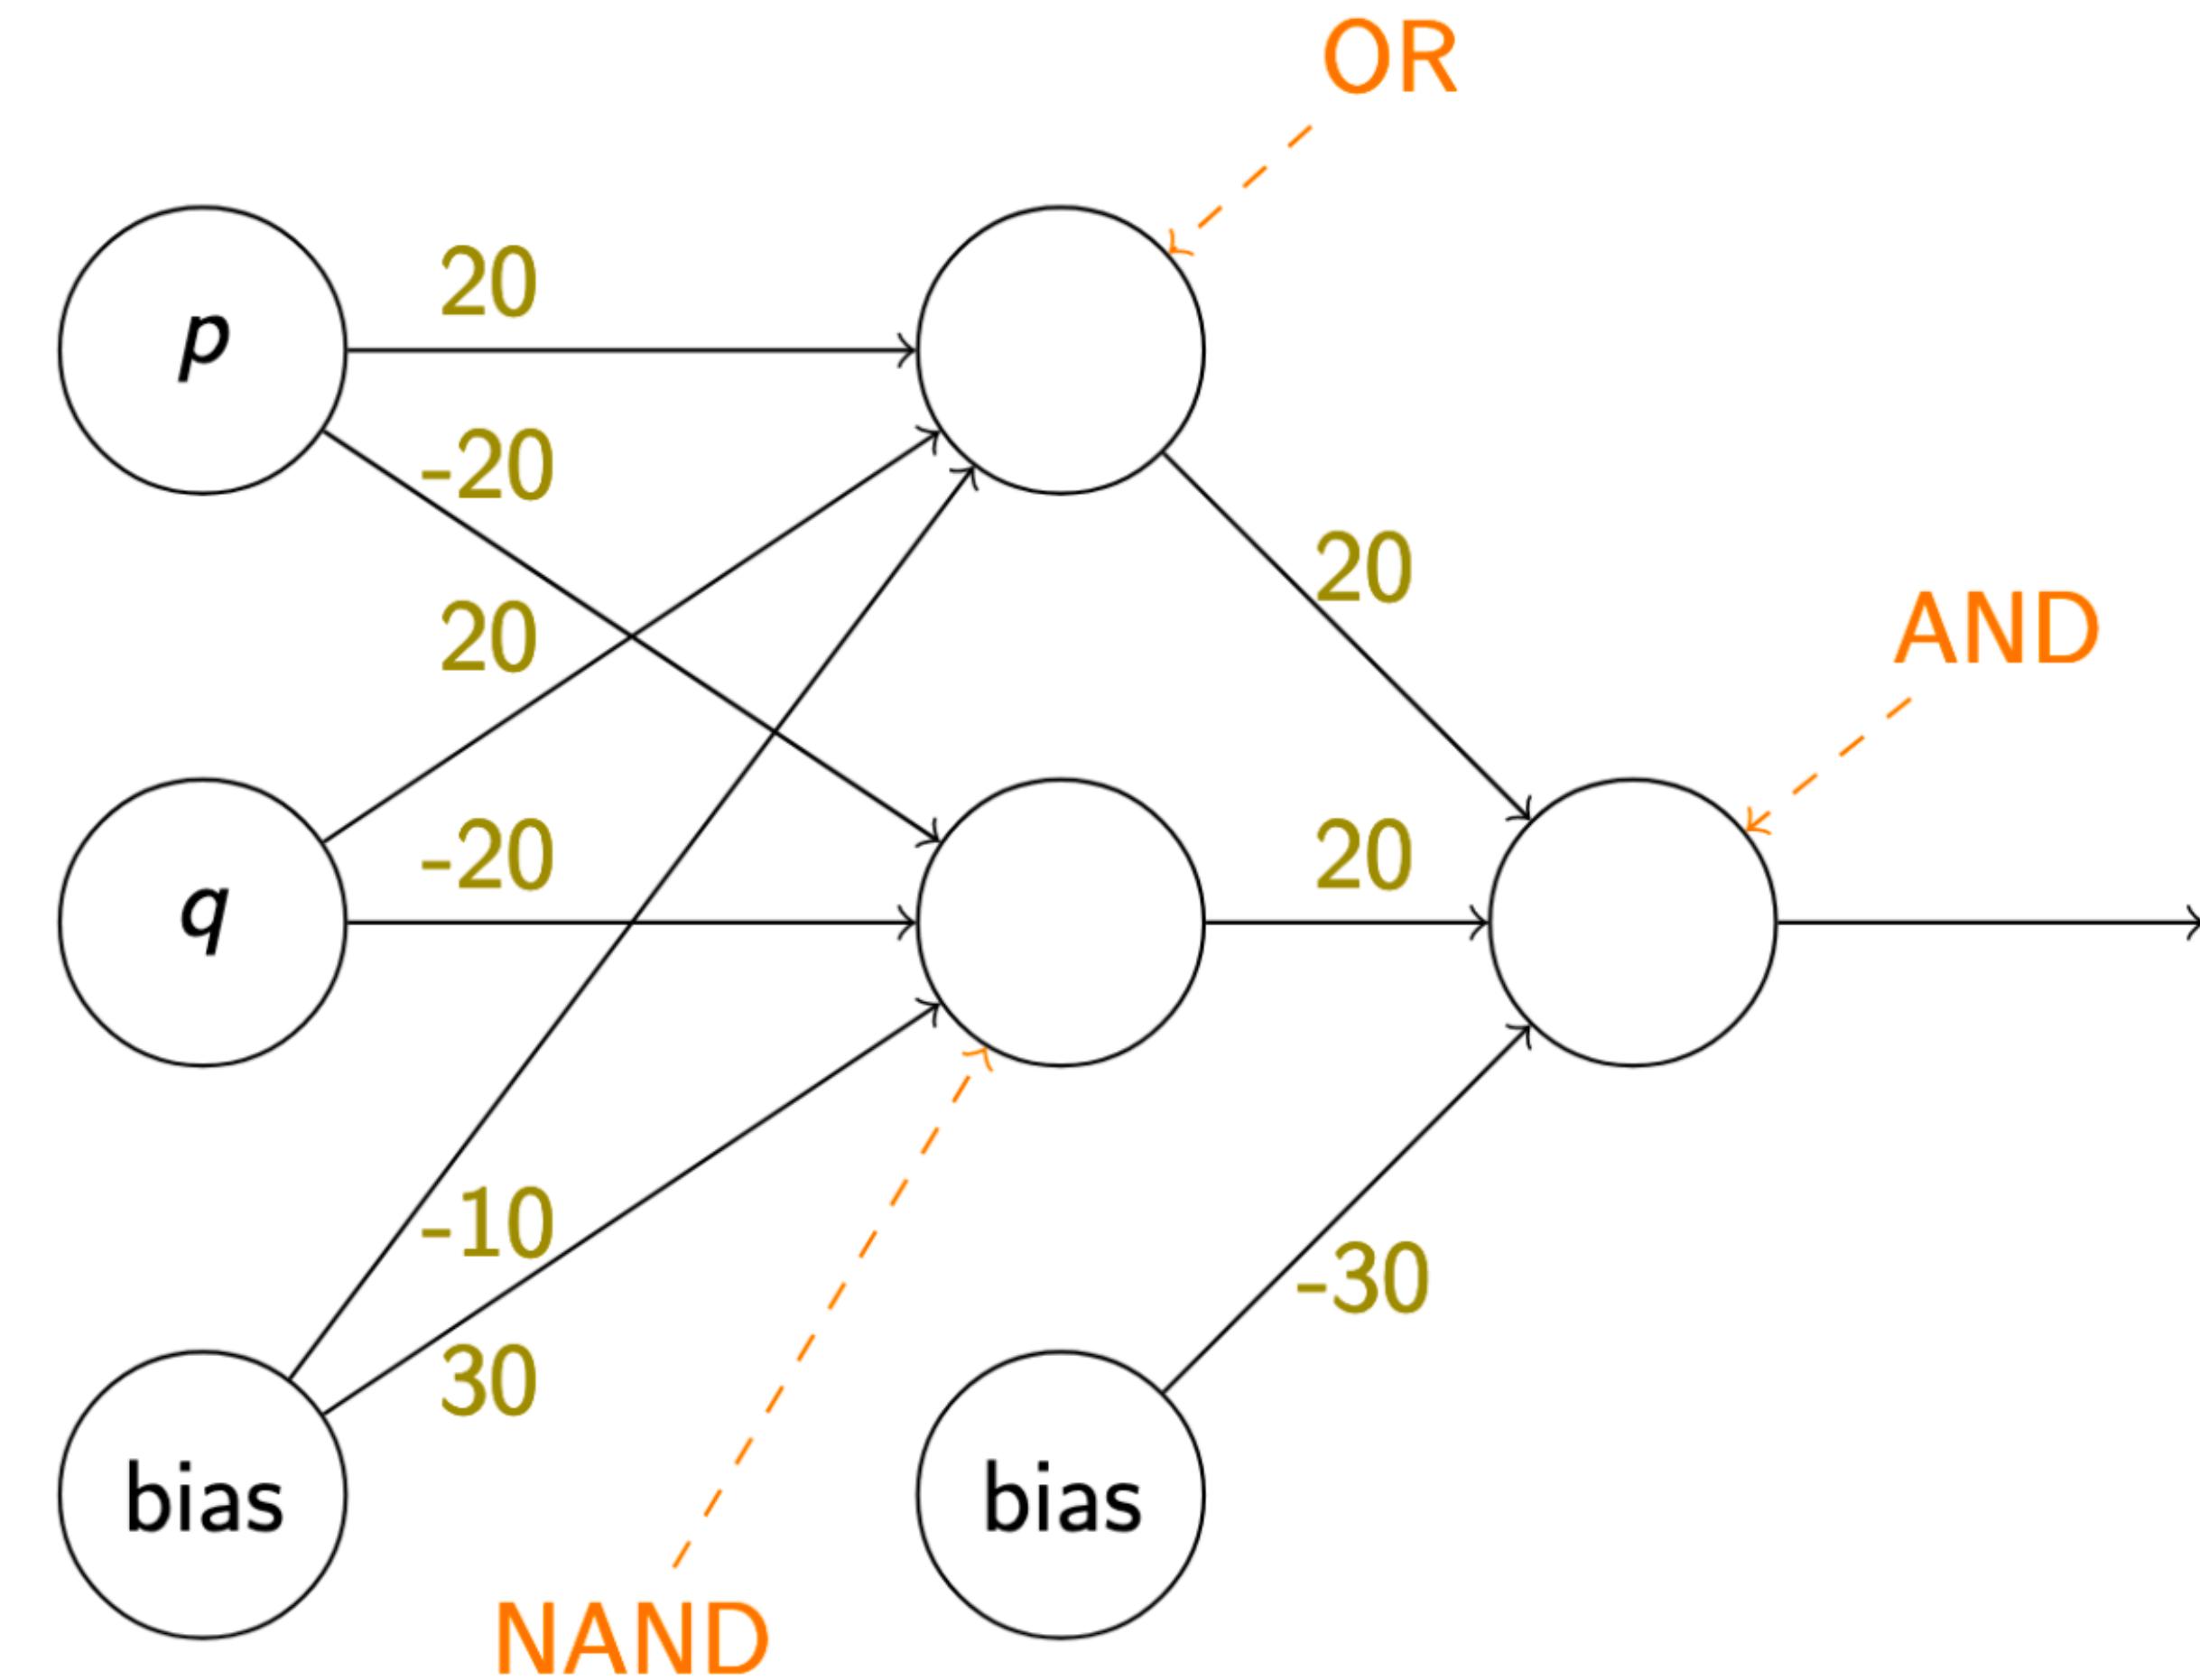


XOR is not linearly separable

# Computing XOR

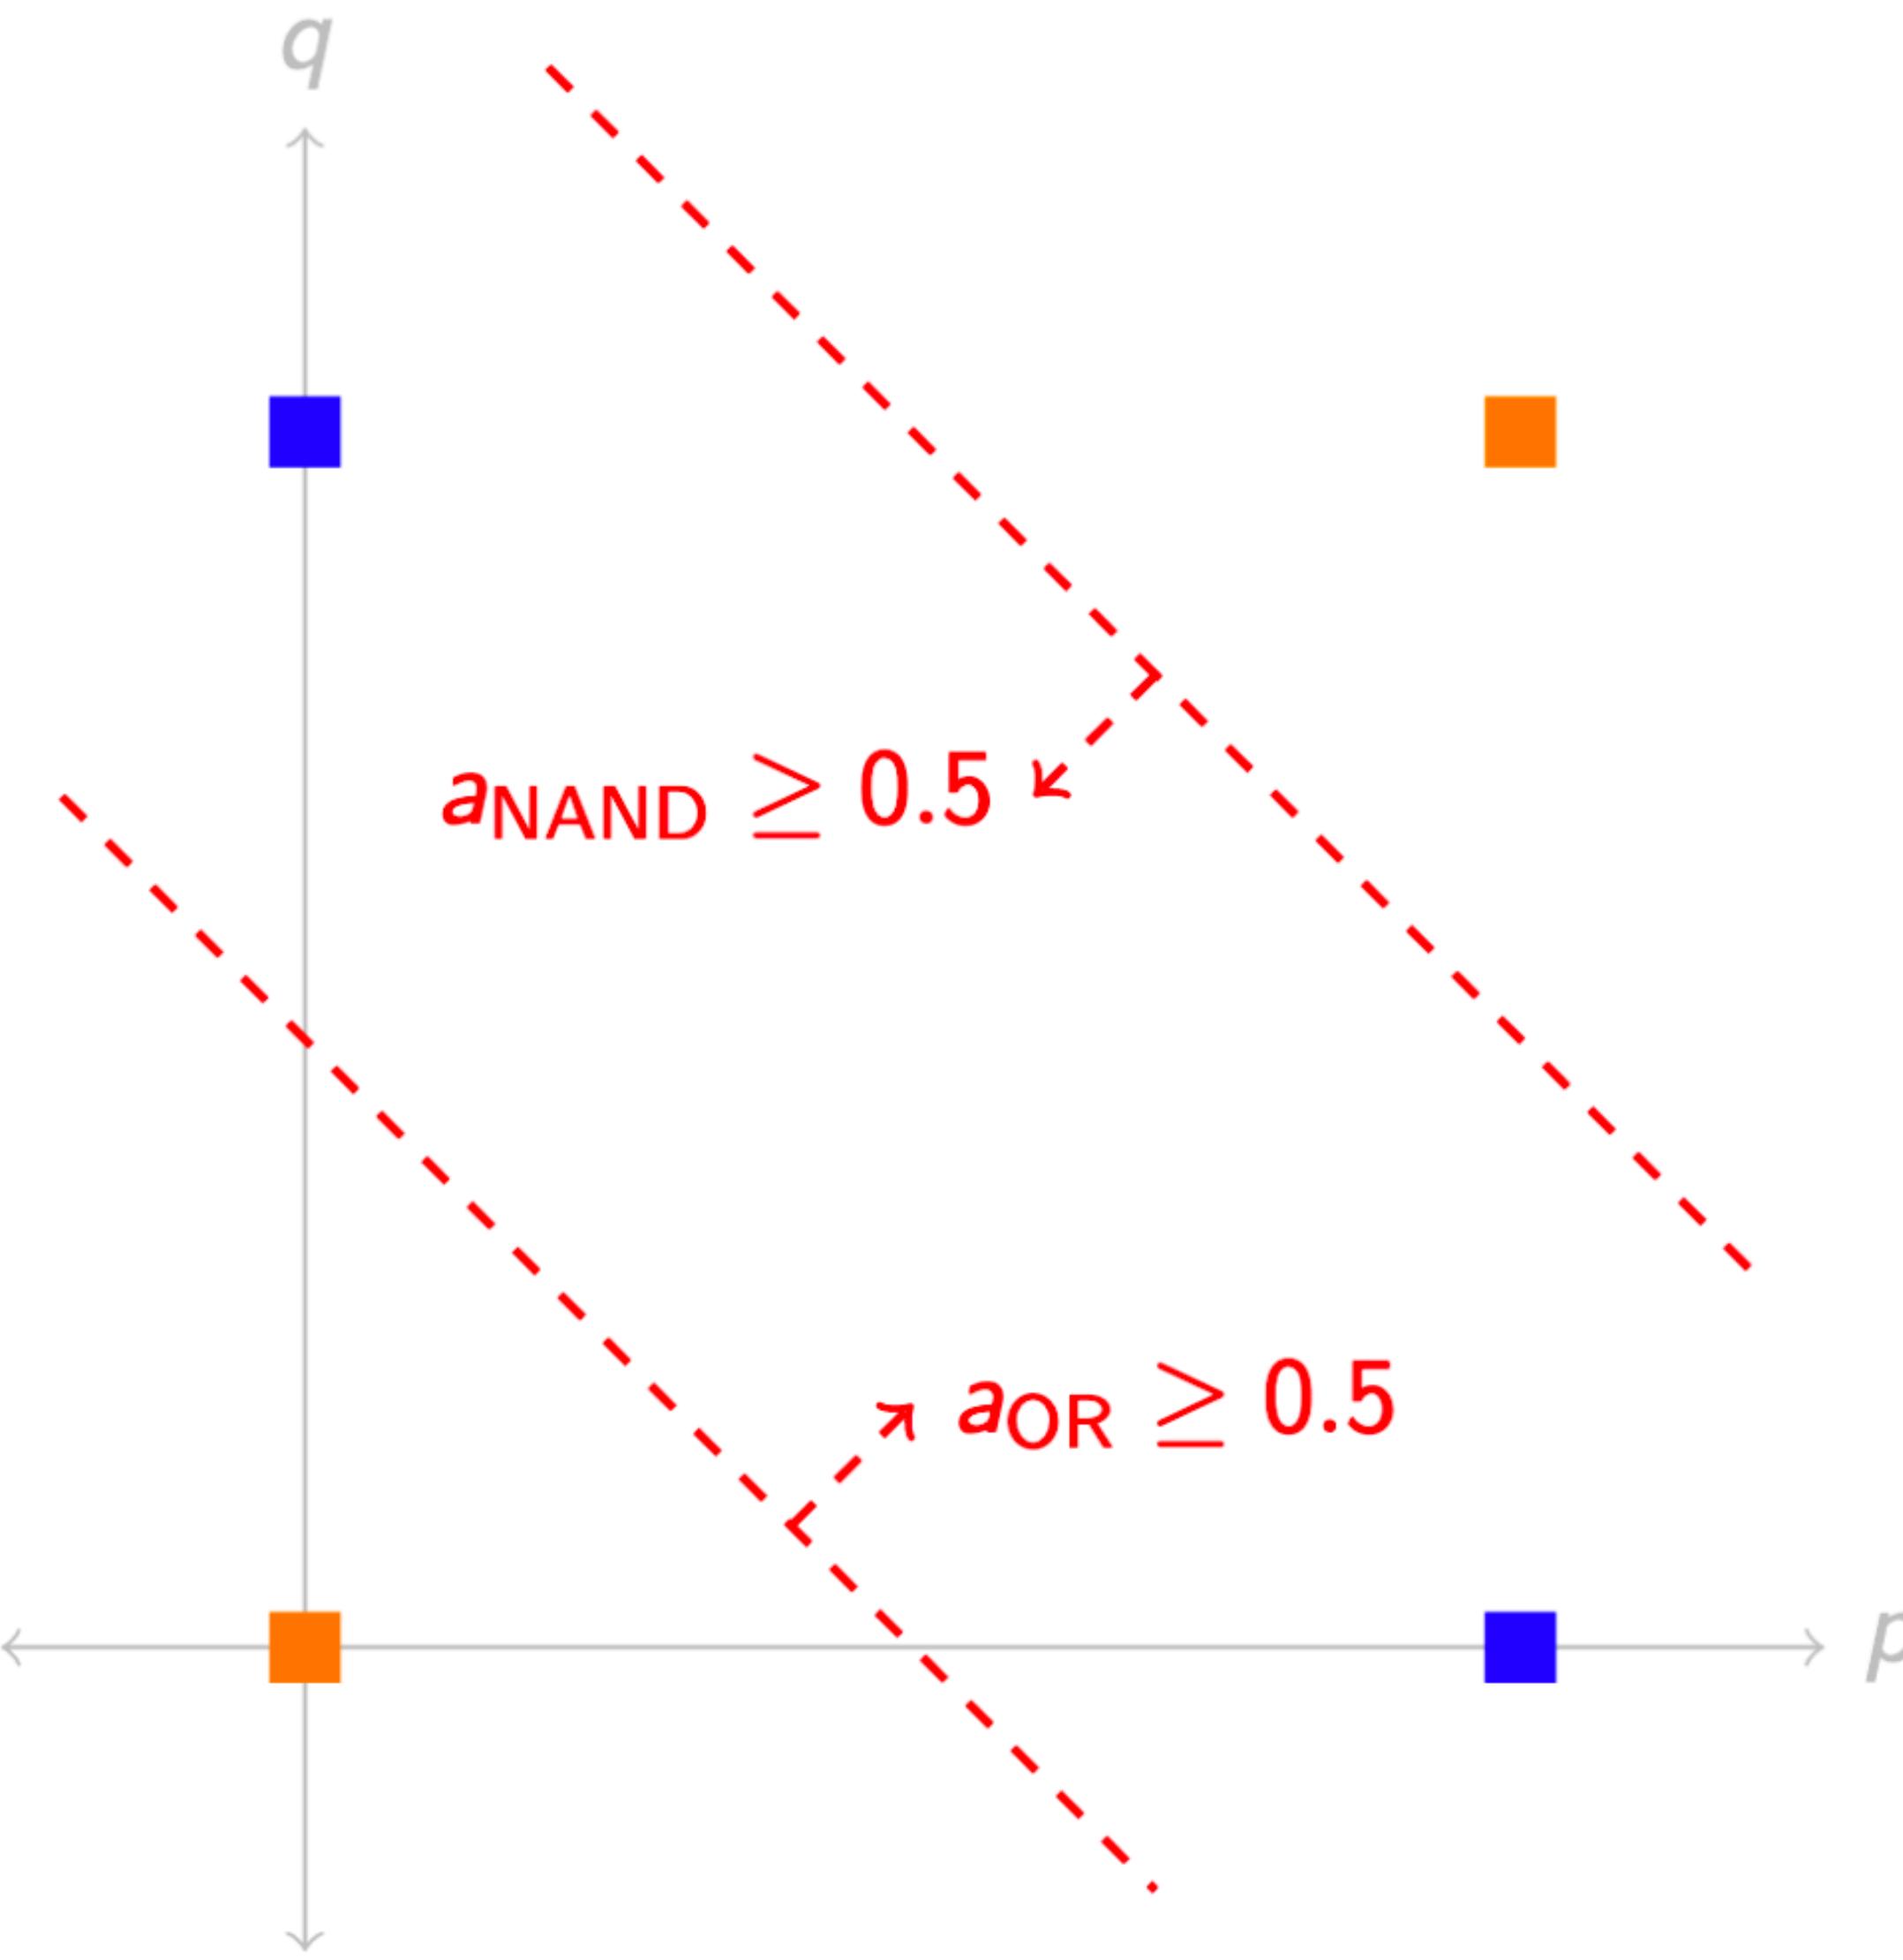


# Computing XOR



Exercise: show that  
NAND behaves as described.

# Computing XOR



# Key Ideas

- **Hidden layers:** intermediate layers of representation that are **not directly used** as outputs
  - Compute **high-level / abstract features** of the input
  - Via training, will **learn which features are helpful** for a given task
  - Caveat: doesn't always learn much more than shallow features
- Adding hidden layers **increases the expressive power** of a neural network
  - Strictly more functions can be computed with hidden layers than without

# Expressive Power

# Expressive Power

- Neural networks with **one hidden layer** are **universal function approximators**

# Expressive Power

- Neural networks with **one hidden layer** are **universal function approximators**
- Let  $f: [0,1]^m \rightarrow \mathbb{R}$  be continuous and  $\epsilon > 0$ . Then there is a one-hidden-layer neural network  $g$  with sigmoid activation such that  $|f(x) - g(x)| < \epsilon$  for all  $x \in [0,1]^m$ .

# Expressive Power

- Neural networks with **one hidden layer** are **universal function approximators**
- Let  $f: [0,1]^m \rightarrow \mathbb{R}$  be continuous and  $\epsilon > 0$ . Then there is a one-hidden-layer neural network  $g$  with sigmoid activation such that  $|f(x) - g(x)| < \epsilon$  for all  $x \in [0,1]^m$ .
- Generalizations (diff activation functions, less bounded, etc.) exist.

# Expressive Power

- Neural networks with **one hidden layer** are **universal function approximators**
- Let  $f: [0,1]^m \rightarrow \mathbb{R}$  be continuous and  $\epsilon > 0$ . Then there is a one-hidden-layer neural network  $g$  with sigmoid activation such that  $|f(x) - g(x)| < \epsilon$  for all  $x \in [0,1]^m$ .
- Generalizations (diff activation functions, less bounded, etc.) exist.
- But:
  - Size of the hidden layer is must be **exponential in  $m$**
  - How does one find/learn such a good approximation?

# Expressive Power

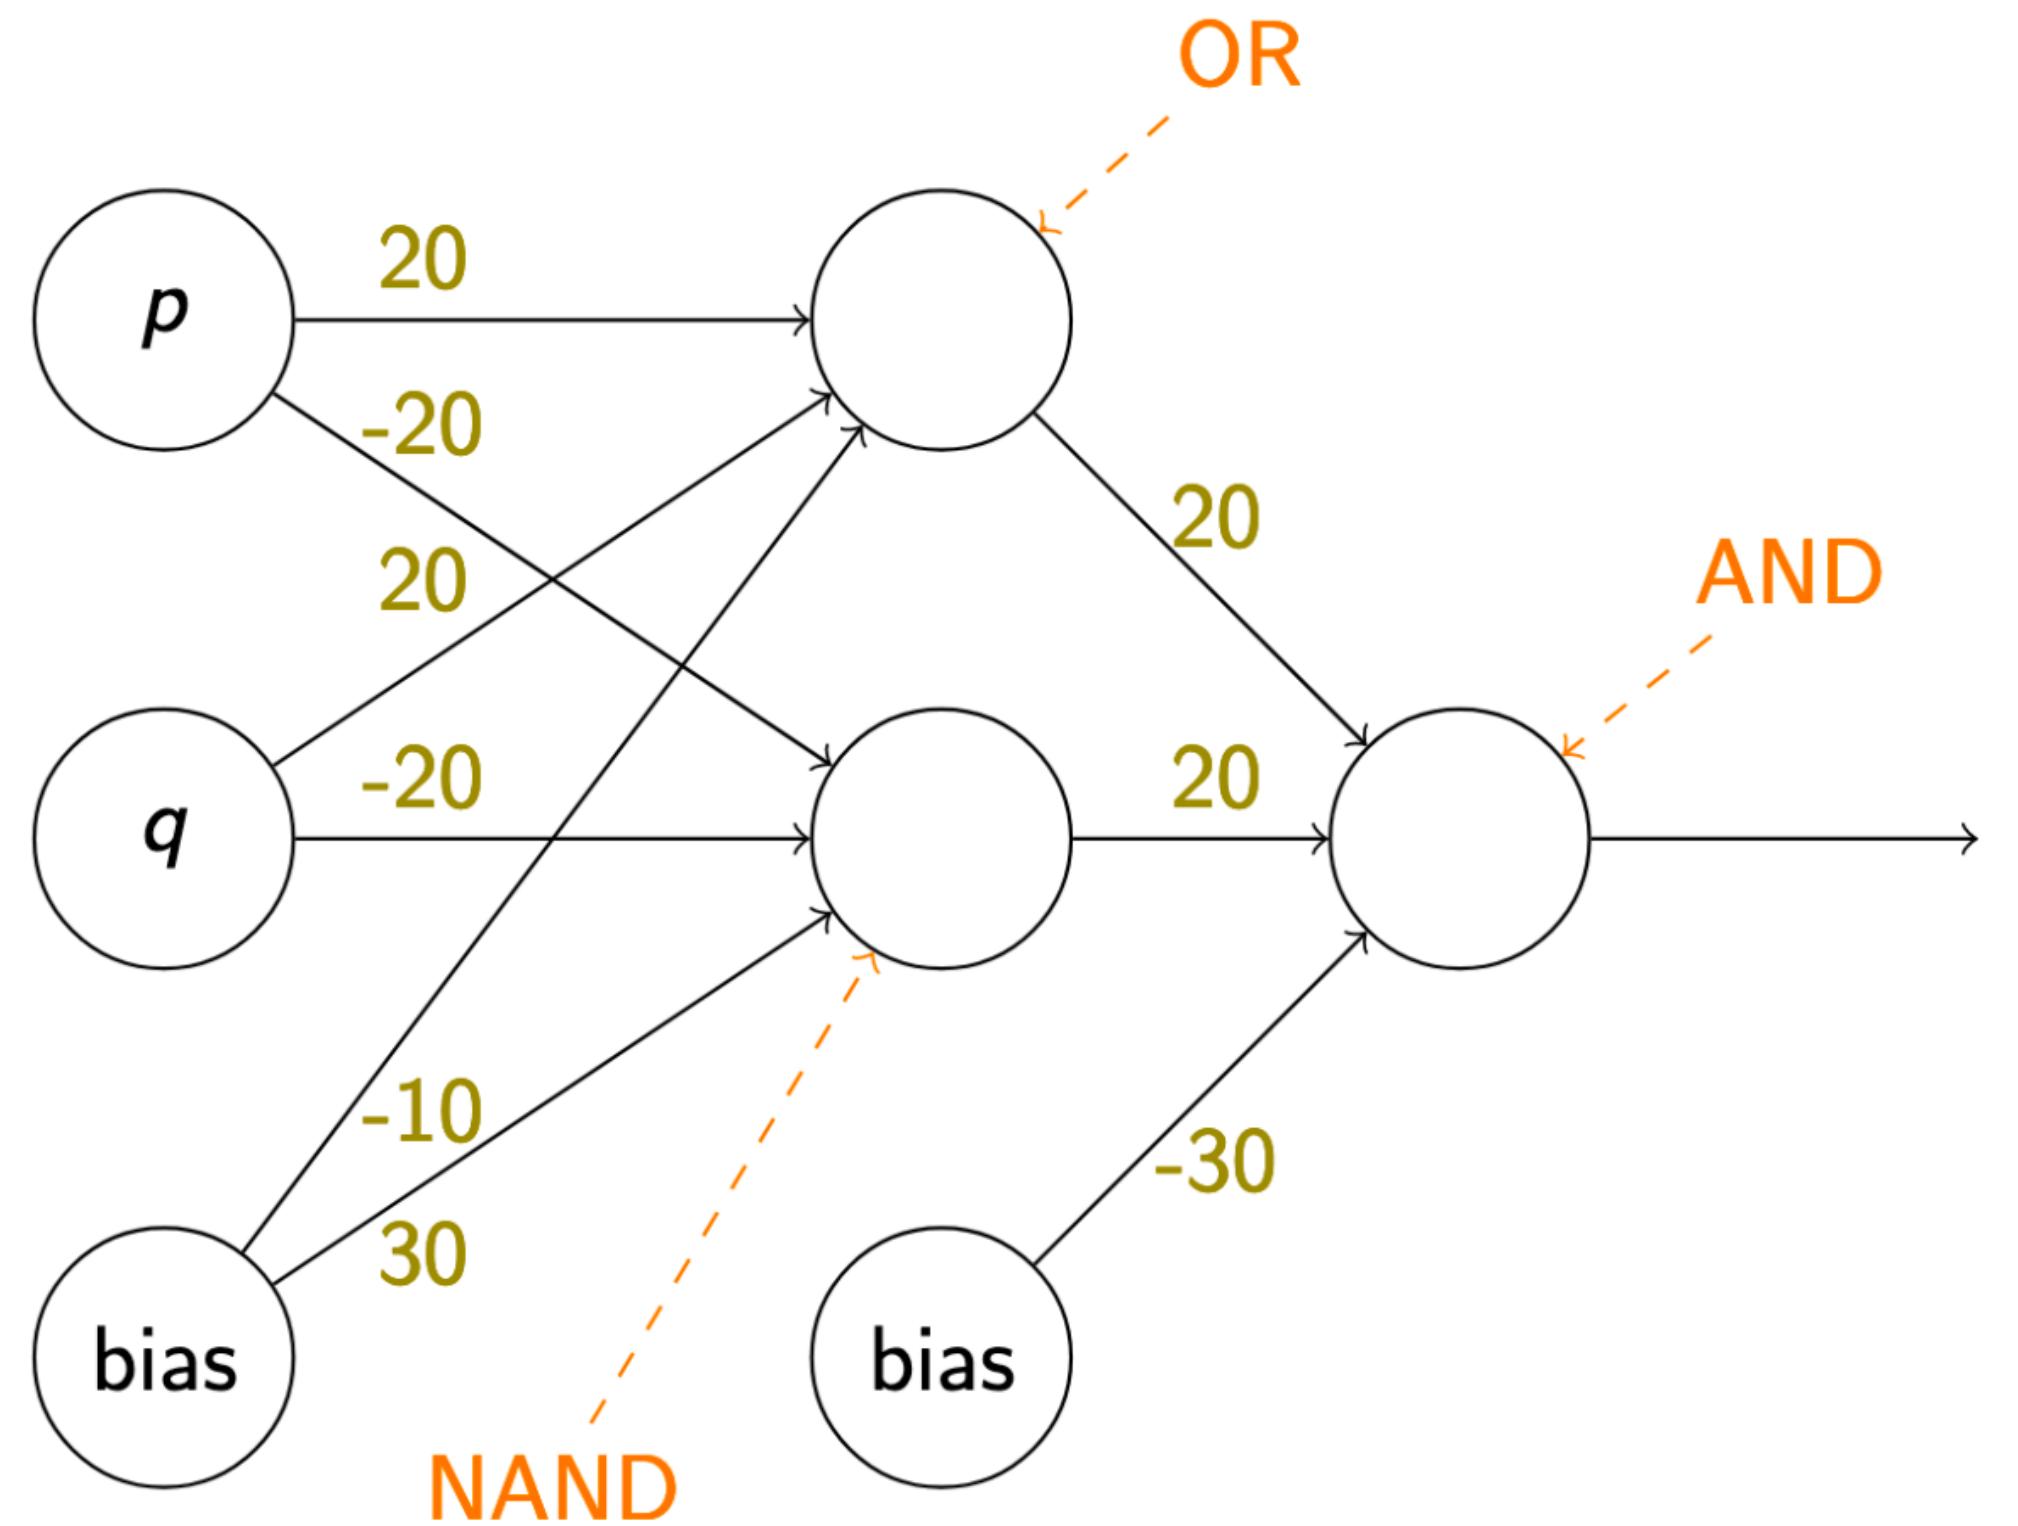
- Neural networks with **one hidden layer** are **universal function approximators**
- Let  $f: [0,1]^m \rightarrow \mathbb{R}$  be continuous and  $\epsilon > 0$ . Then there is a one-hidden-layer neural network  $g$  with sigmoid activation such that  $|f(x) - g(x)| < \epsilon$  for all  $x \in [0,1]^m$ .
- Generalizations (diff activation functions, less bounded, etc.) exist.
- But:
  - Size of the hidden layer is must be **exponential in  $m$**
  - How does one find/learn such a good approximation?
- Nice walkthrough: <http://neuralnetworksanddeeplearning.com/chap4.html>

# Expressive Power

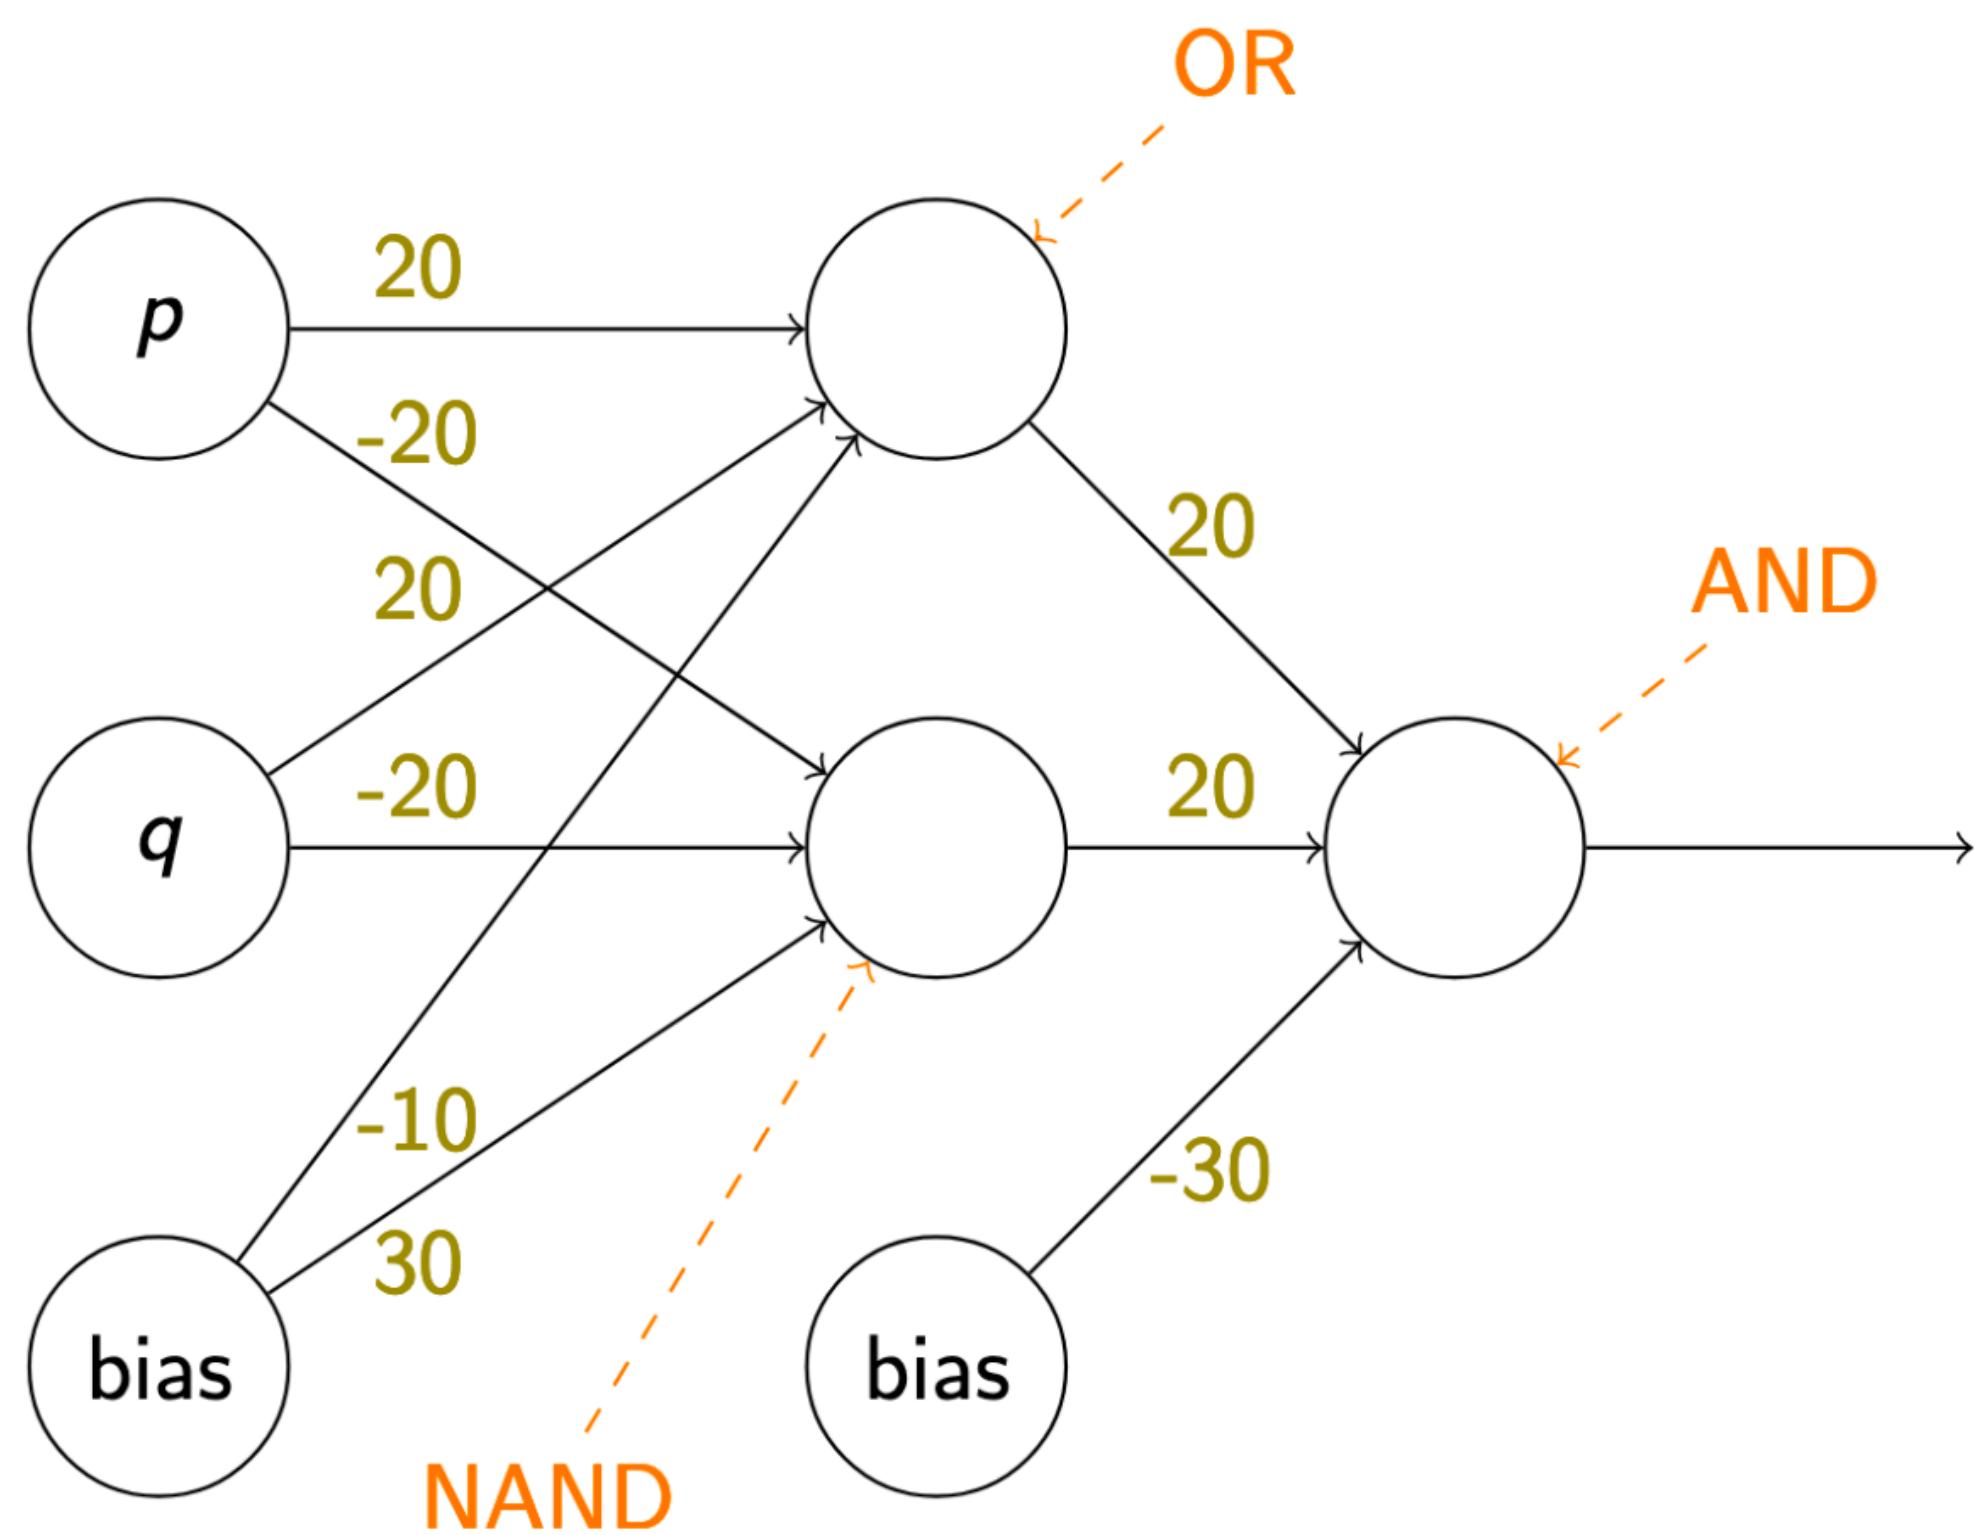
- Neural networks with **one hidden layer** are **universal function approximators**
- Let  $f: [0,1]^m \rightarrow \mathbb{R}$  be continuous and  $\epsilon > 0$ . Then there is a one-hidden-layer neural network  $g$  with sigmoid activation such that  $|f(x) - g(x)| < \epsilon$  for all  $x \in [0,1]^m$ .
- Generalizations (diff activation functions, less bounded, etc.) exist.
- But:
  - Size of the hidden layer is must be **exponential in  $m$**
  - How does one find/learn such a good approximation?
- Nice walkthrough: <http://neuralnetworksanddeeplearning.com/chap4.html>
- See also [GBC](#) 6.4.1 for more references, generalizations, discussion

# Feed-forward networks aka Multi-Layer Perceptrons (MLP)

# XOR Network

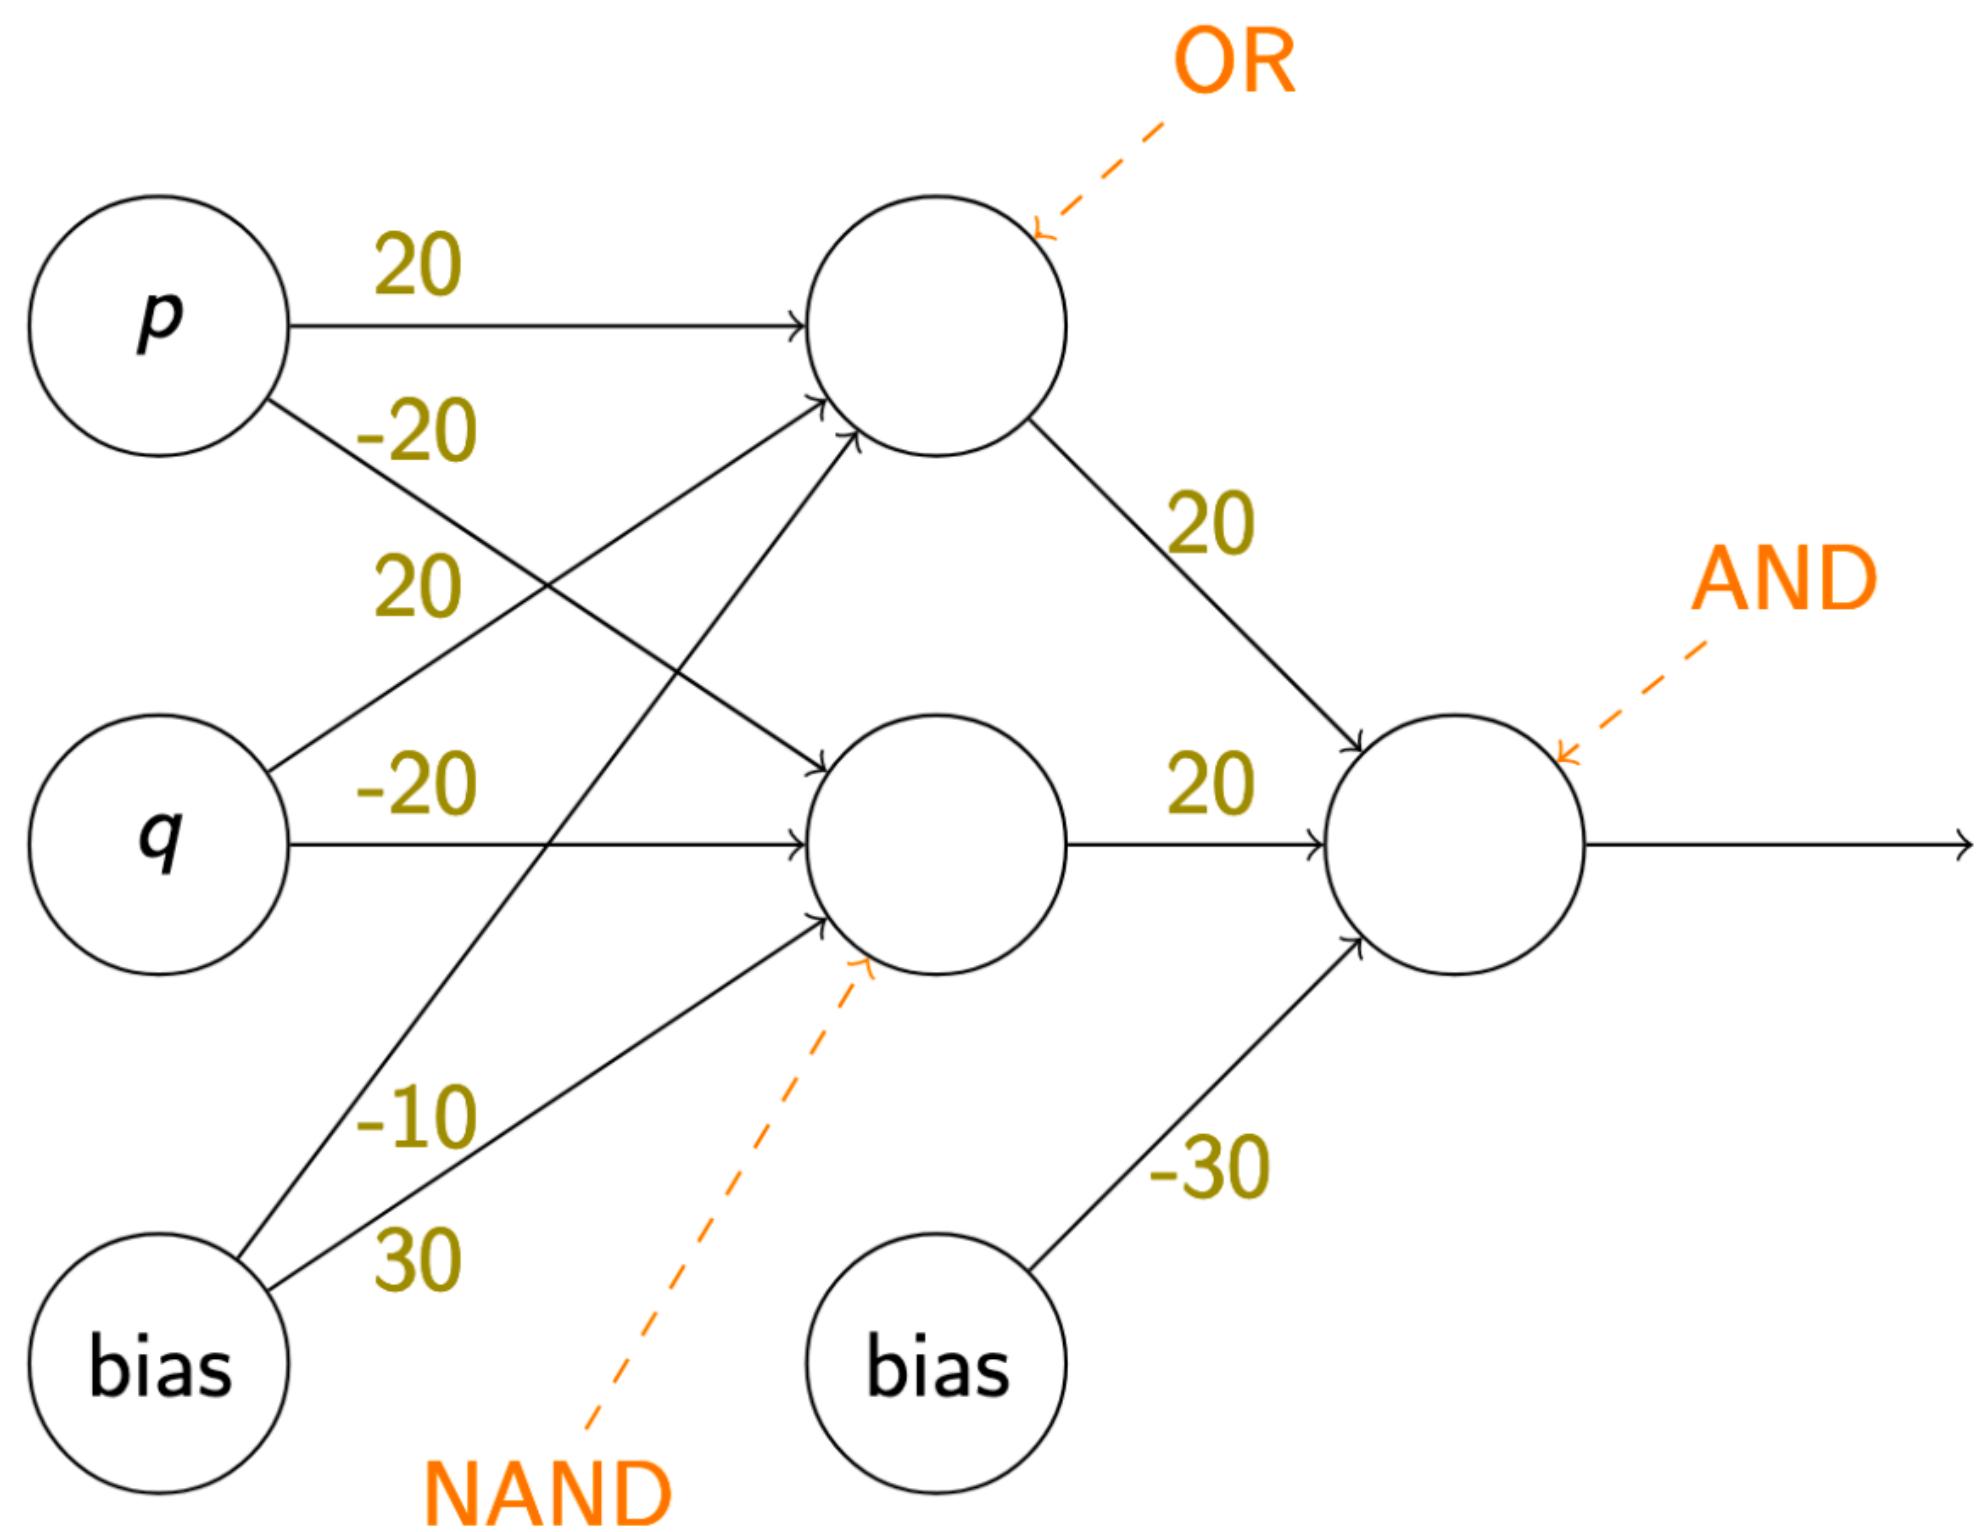


# XOR Network



$$a_{\text{or}} = \sigma \left( [w_p^{\text{or}} \quad w_q^{\text{or}}] \begin{bmatrix} p \\ q \end{bmatrix} + b^{\text{or}} \right)$$

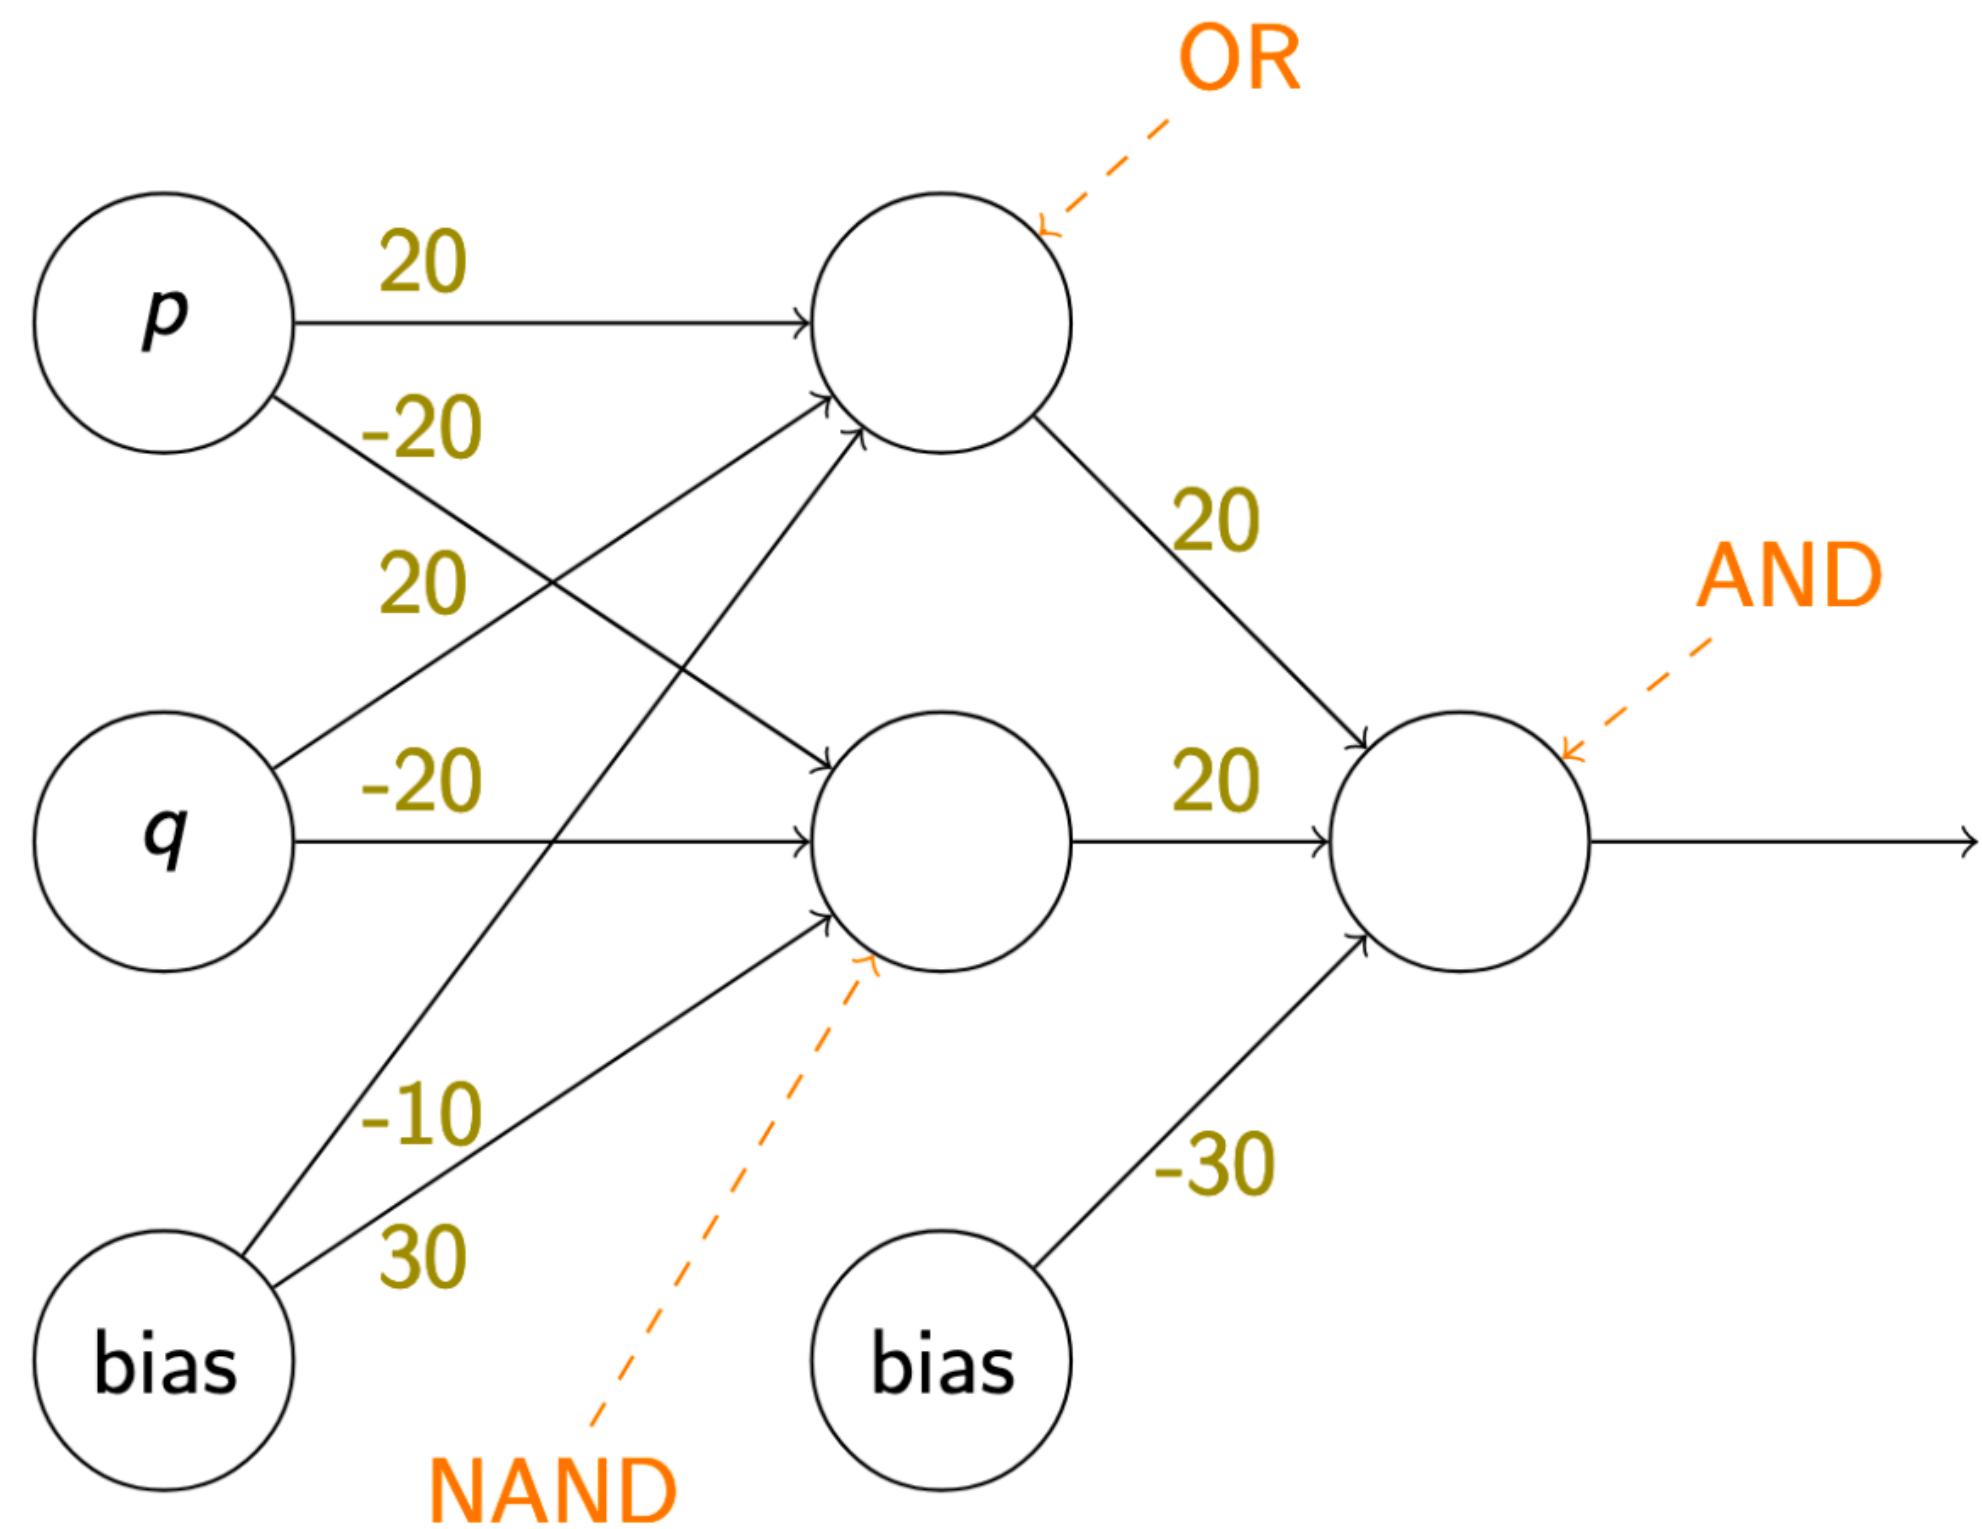
# XOR Network



$$a_{\text{or}} = \sigma \left( [w_p^{\text{or}} \quad w_q^{\text{or}}] \begin{bmatrix} p \\ q \end{bmatrix} + b^{\text{or}} \right)$$

$$a_{\text{nand}} = \sigma \left( [w_p^{\text{nand}} \quad w_q^{\text{nand}}] \begin{bmatrix} p \\ q \end{bmatrix} + b^{\text{nand}} \right)$$

# XOR Network

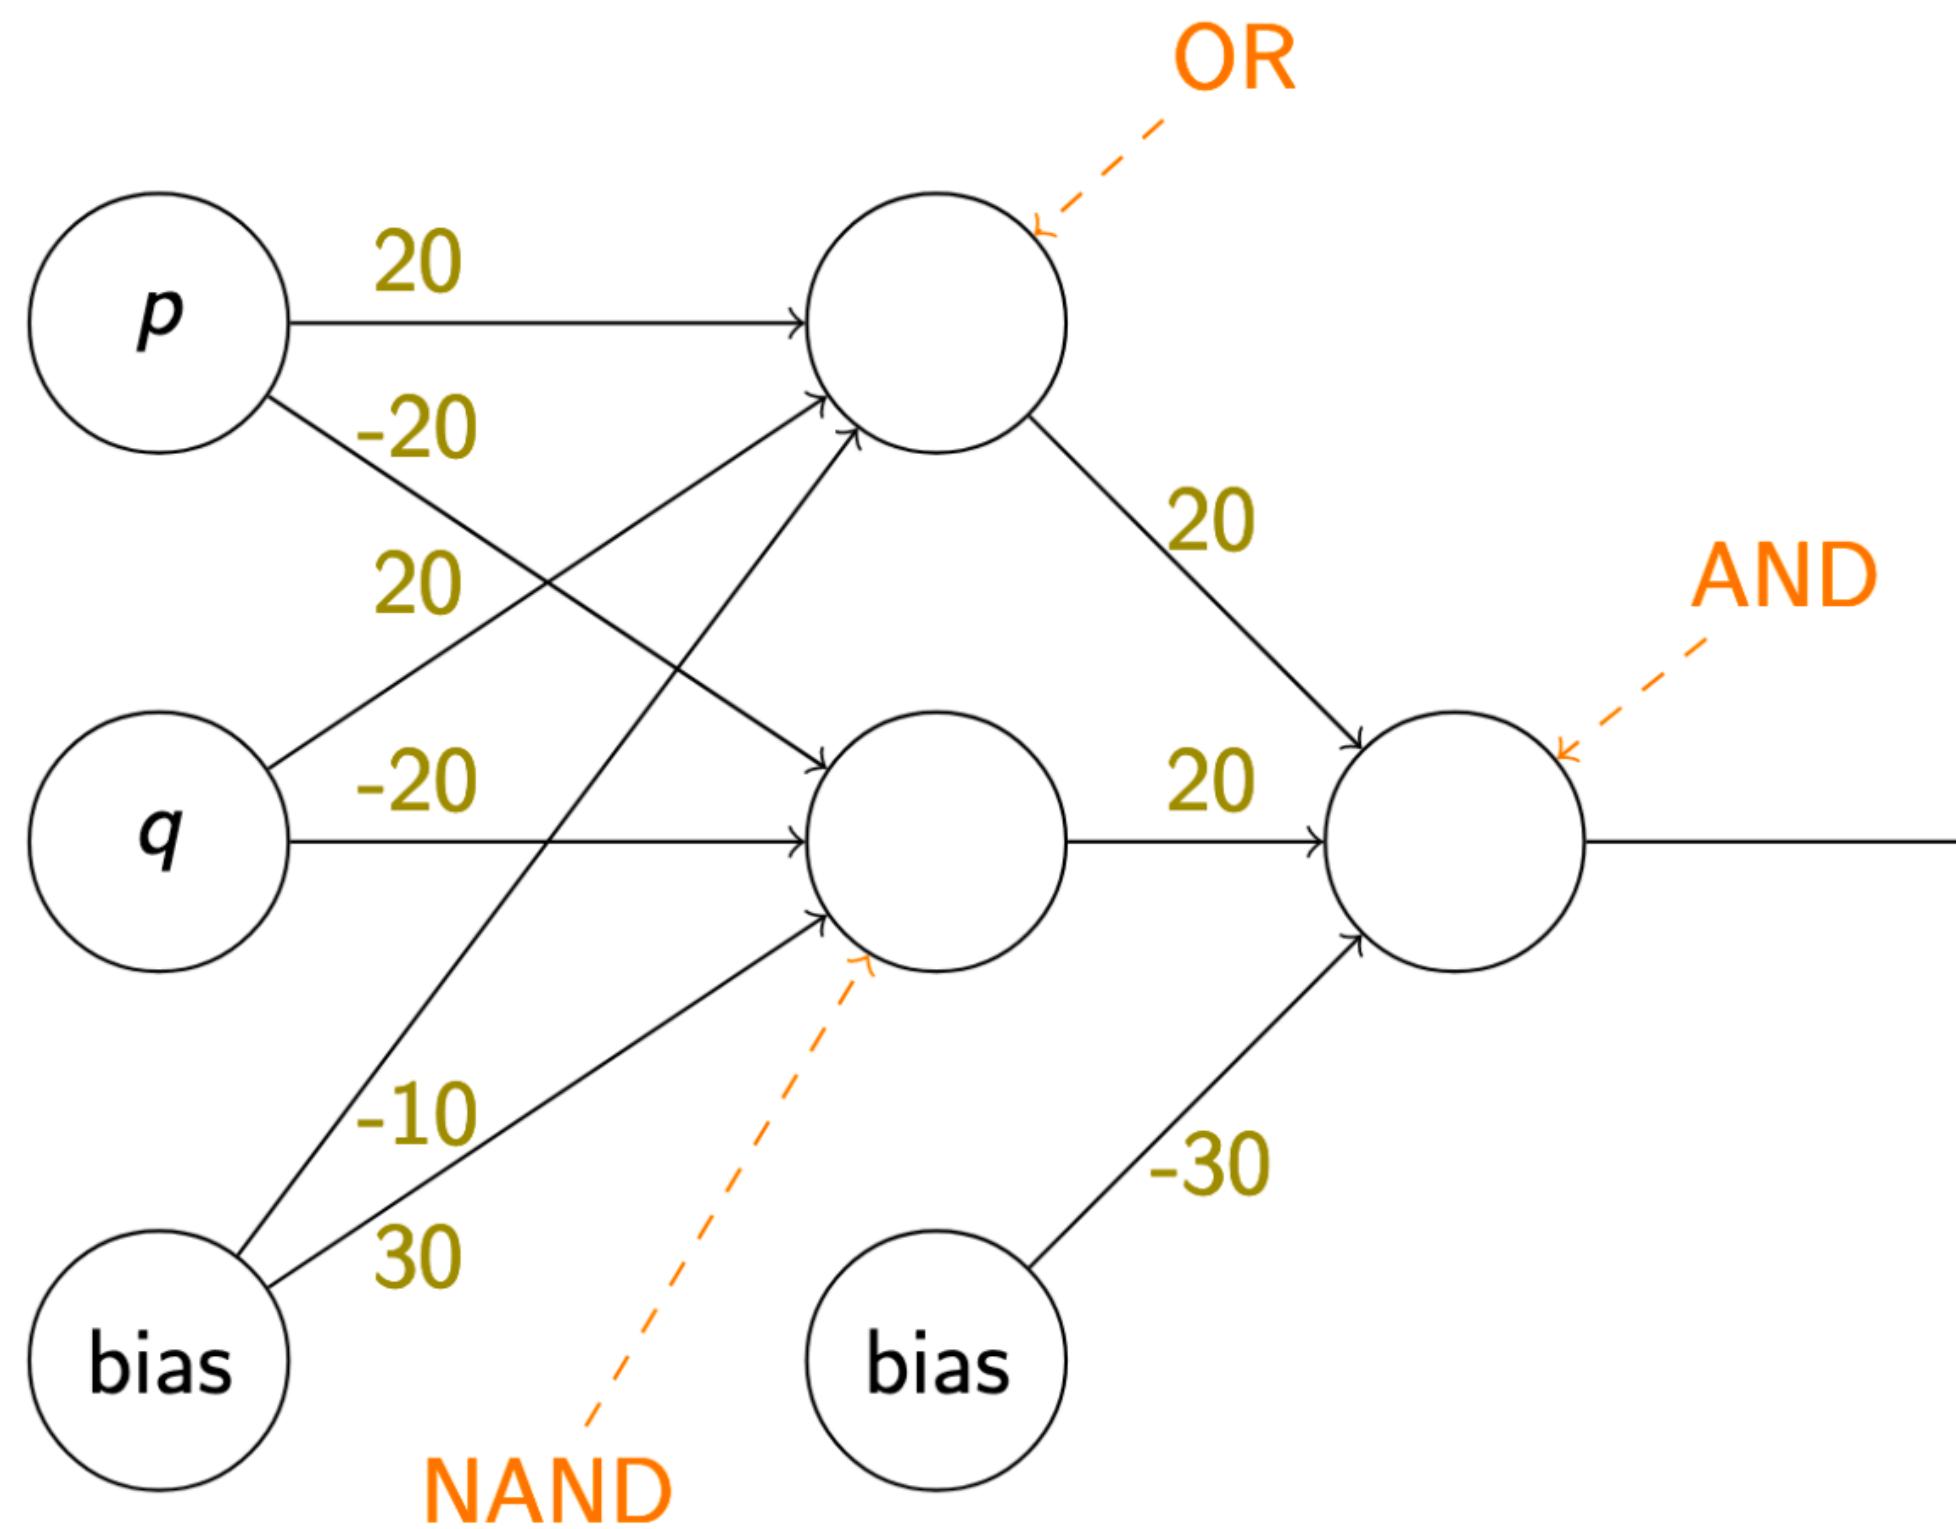


$$a_{\text{or}} = \sigma \left( [w_p^{\text{or}} \quad w_q^{\text{or}}] \begin{bmatrix} p \\ q \end{bmatrix} + b^{\text{or}} \right)$$

$$a_{\text{nand}} = \sigma \left( [w_p^{\text{nand}} \quad w_q^{\text{nand}}] \begin{bmatrix} p \\ q \end{bmatrix} + b^{\text{nand}} \right)$$

$$\begin{bmatrix} a_{\text{or}} \\ a_{\text{nand}} \end{bmatrix} = \sigma \left( \begin{bmatrix} w_p^{\text{or}} & w_q^{\text{or}} \\ w_p^{\text{nand}} & w_q^{\text{nand}} \end{bmatrix} \begin{bmatrix} p \\ q \end{bmatrix} + \begin{bmatrix} b^{\text{or}} \\ b^{\text{nand}} \end{bmatrix} \right)$$

# XOR Network

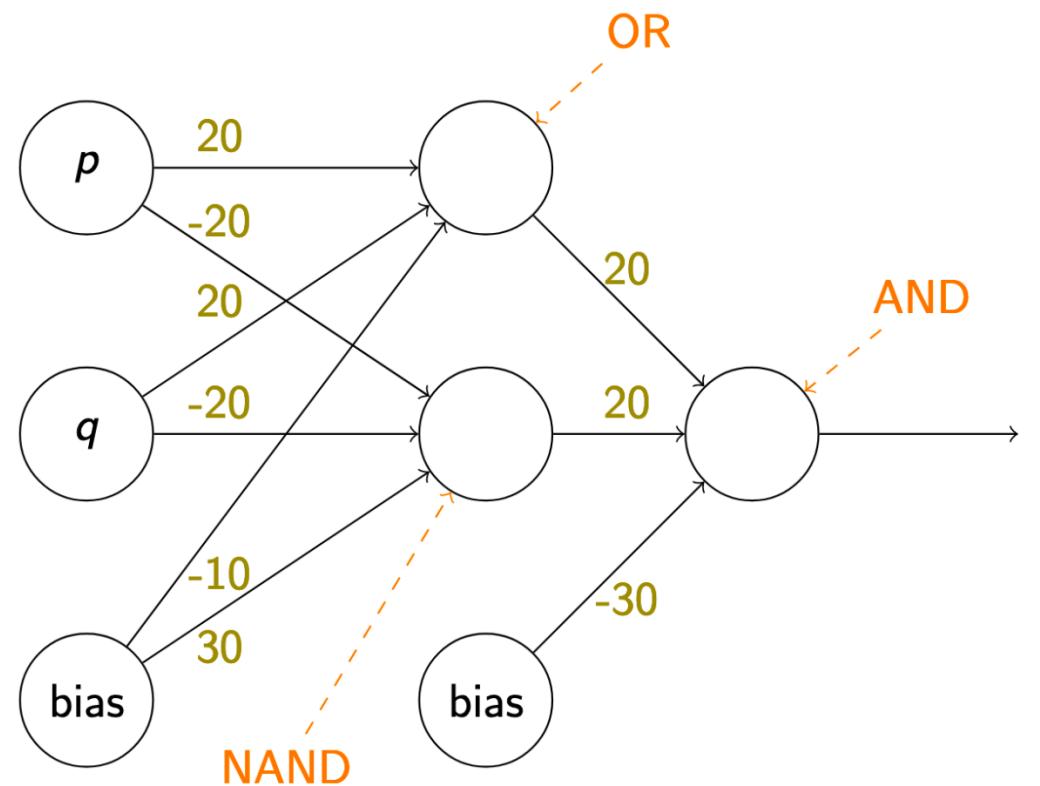


$$a_{\text{or}} = \sigma \left( [w_p^{\text{or}} \quad w_q^{\text{or}}] \begin{bmatrix} p \\ q \end{bmatrix} + b^{\text{or}} \right)$$

$$a_{\text{nand}} = \sigma \left( [w_p^{\text{nand}} \quad w_q^{\text{nand}}] \begin{bmatrix} p \\ q \end{bmatrix} + b^{\text{nand}} \right)$$

$$\begin{bmatrix} a_{\text{or}} \\ a_{\text{nand}} \end{bmatrix} = \sigma \left( \begin{bmatrix} w_p^{\text{or}} & w_q^{\text{or}} \\ w_p^{\text{nand}} & w_q^{\text{nand}} \end{bmatrix} \begin{bmatrix} p \\ q \end{bmatrix} + \begin{bmatrix} b^{\text{or}} \\ b^{\text{nand}} \end{bmatrix} \right)$$

$$a_{\text{and}} = \sigma \left( [w_{\text{or}}^{\text{and}} \quad w_{\text{nand}}^{\text{and}}] \begin{bmatrix} a_{\text{or}} \\ a_{\text{nand}} \end{bmatrix} + b^{\text{and}} \right)$$



# XOR Network

$$a_{\text{and}} = \sigma \left( \begin{bmatrix} w_{\text{or}}^{\text{and}} & w_{\text{nand}}^{\text{and}} \end{bmatrix} \sigma \begin{bmatrix} w_p^{\text{or}} & w_q^{\text{or}} \\ w_p^{\text{nand}} & w_q^{\text{nand}} \end{bmatrix} \begin{bmatrix} a_p \\ a_q \end{bmatrix} + \begin{bmatrix} b^{\text{or}} \\ b^{\text{nand}} \end{bmatrix} \right) + b^{\text{and}}$$

# Generalizing

$$a_{\text{and}} = \sigma \left( \begin{bmatrix} w_{\text{or}}^{\text{and}} & w_{\text{nand}}^{\text{and}} \end{bmatrix} \sigma \left( \begin{bmatrix} w_p^{\text{or}} & w_q^{\text{or}} \\ w_p^{\text{nand}} & w_q^{\text{nand}} \end{bmatrix} \begin{bmatrix} a_p \\ a_q \end{bmatrix} + \begin{bmatrix} b^{\text{or}} \\ b^{\text{nand}} \end{bmatrix} \right) + b^{\text{and}} \right)$$

# Generalizing

$$a_{\text{and}} = \sigma \left( \begin{bmatrix} w_{\text{or}}^{\text{and}} & w_{\text{nand}}^{\text{and}} \end{bmatrix} \sigma \left( \begin{bmatrix} w_p^{\text{or}} & w_q^{\text{or}} \\ w_p^{\text{nand}} & w_q^{\text{nand}} \end{bmatrix} \begin{bmatrix} a_p \\ a_q \end{bmatrix} + \begin{bmatrix} b^{\text{or}} \\ b^{\text{nand}} \end{bmatrix} \right) + b^{\text{and}} \right)$$

$$\hat{y} = f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right)$$

# Generalizing

$$a_{\text{and}} = \sigma \left( \begin{bmatrix} w_{\text{or}}^{\text{and}} & w_{\text{nand}}^{\text{and}} \\ w_{\text{or}}^{\text{and}} & w_{\text{nand}}^{\text{and}} \end{bmatrix} \sigma \left( \begin{bmatrix} w_p^{\text{or}} & w_q^{\text{or}} \\ w_p^{\text{nand}} & w_q^{\text{nand}} \end{bmatrix} \begin{bmatrix} a_p \\ a_q \end{bmatrix} + \begin{bmatrix} b^{\text{or}} \\ b^{\text{nand}} \end{bmatrix} \right) + b^{\text{and}} \right)$$

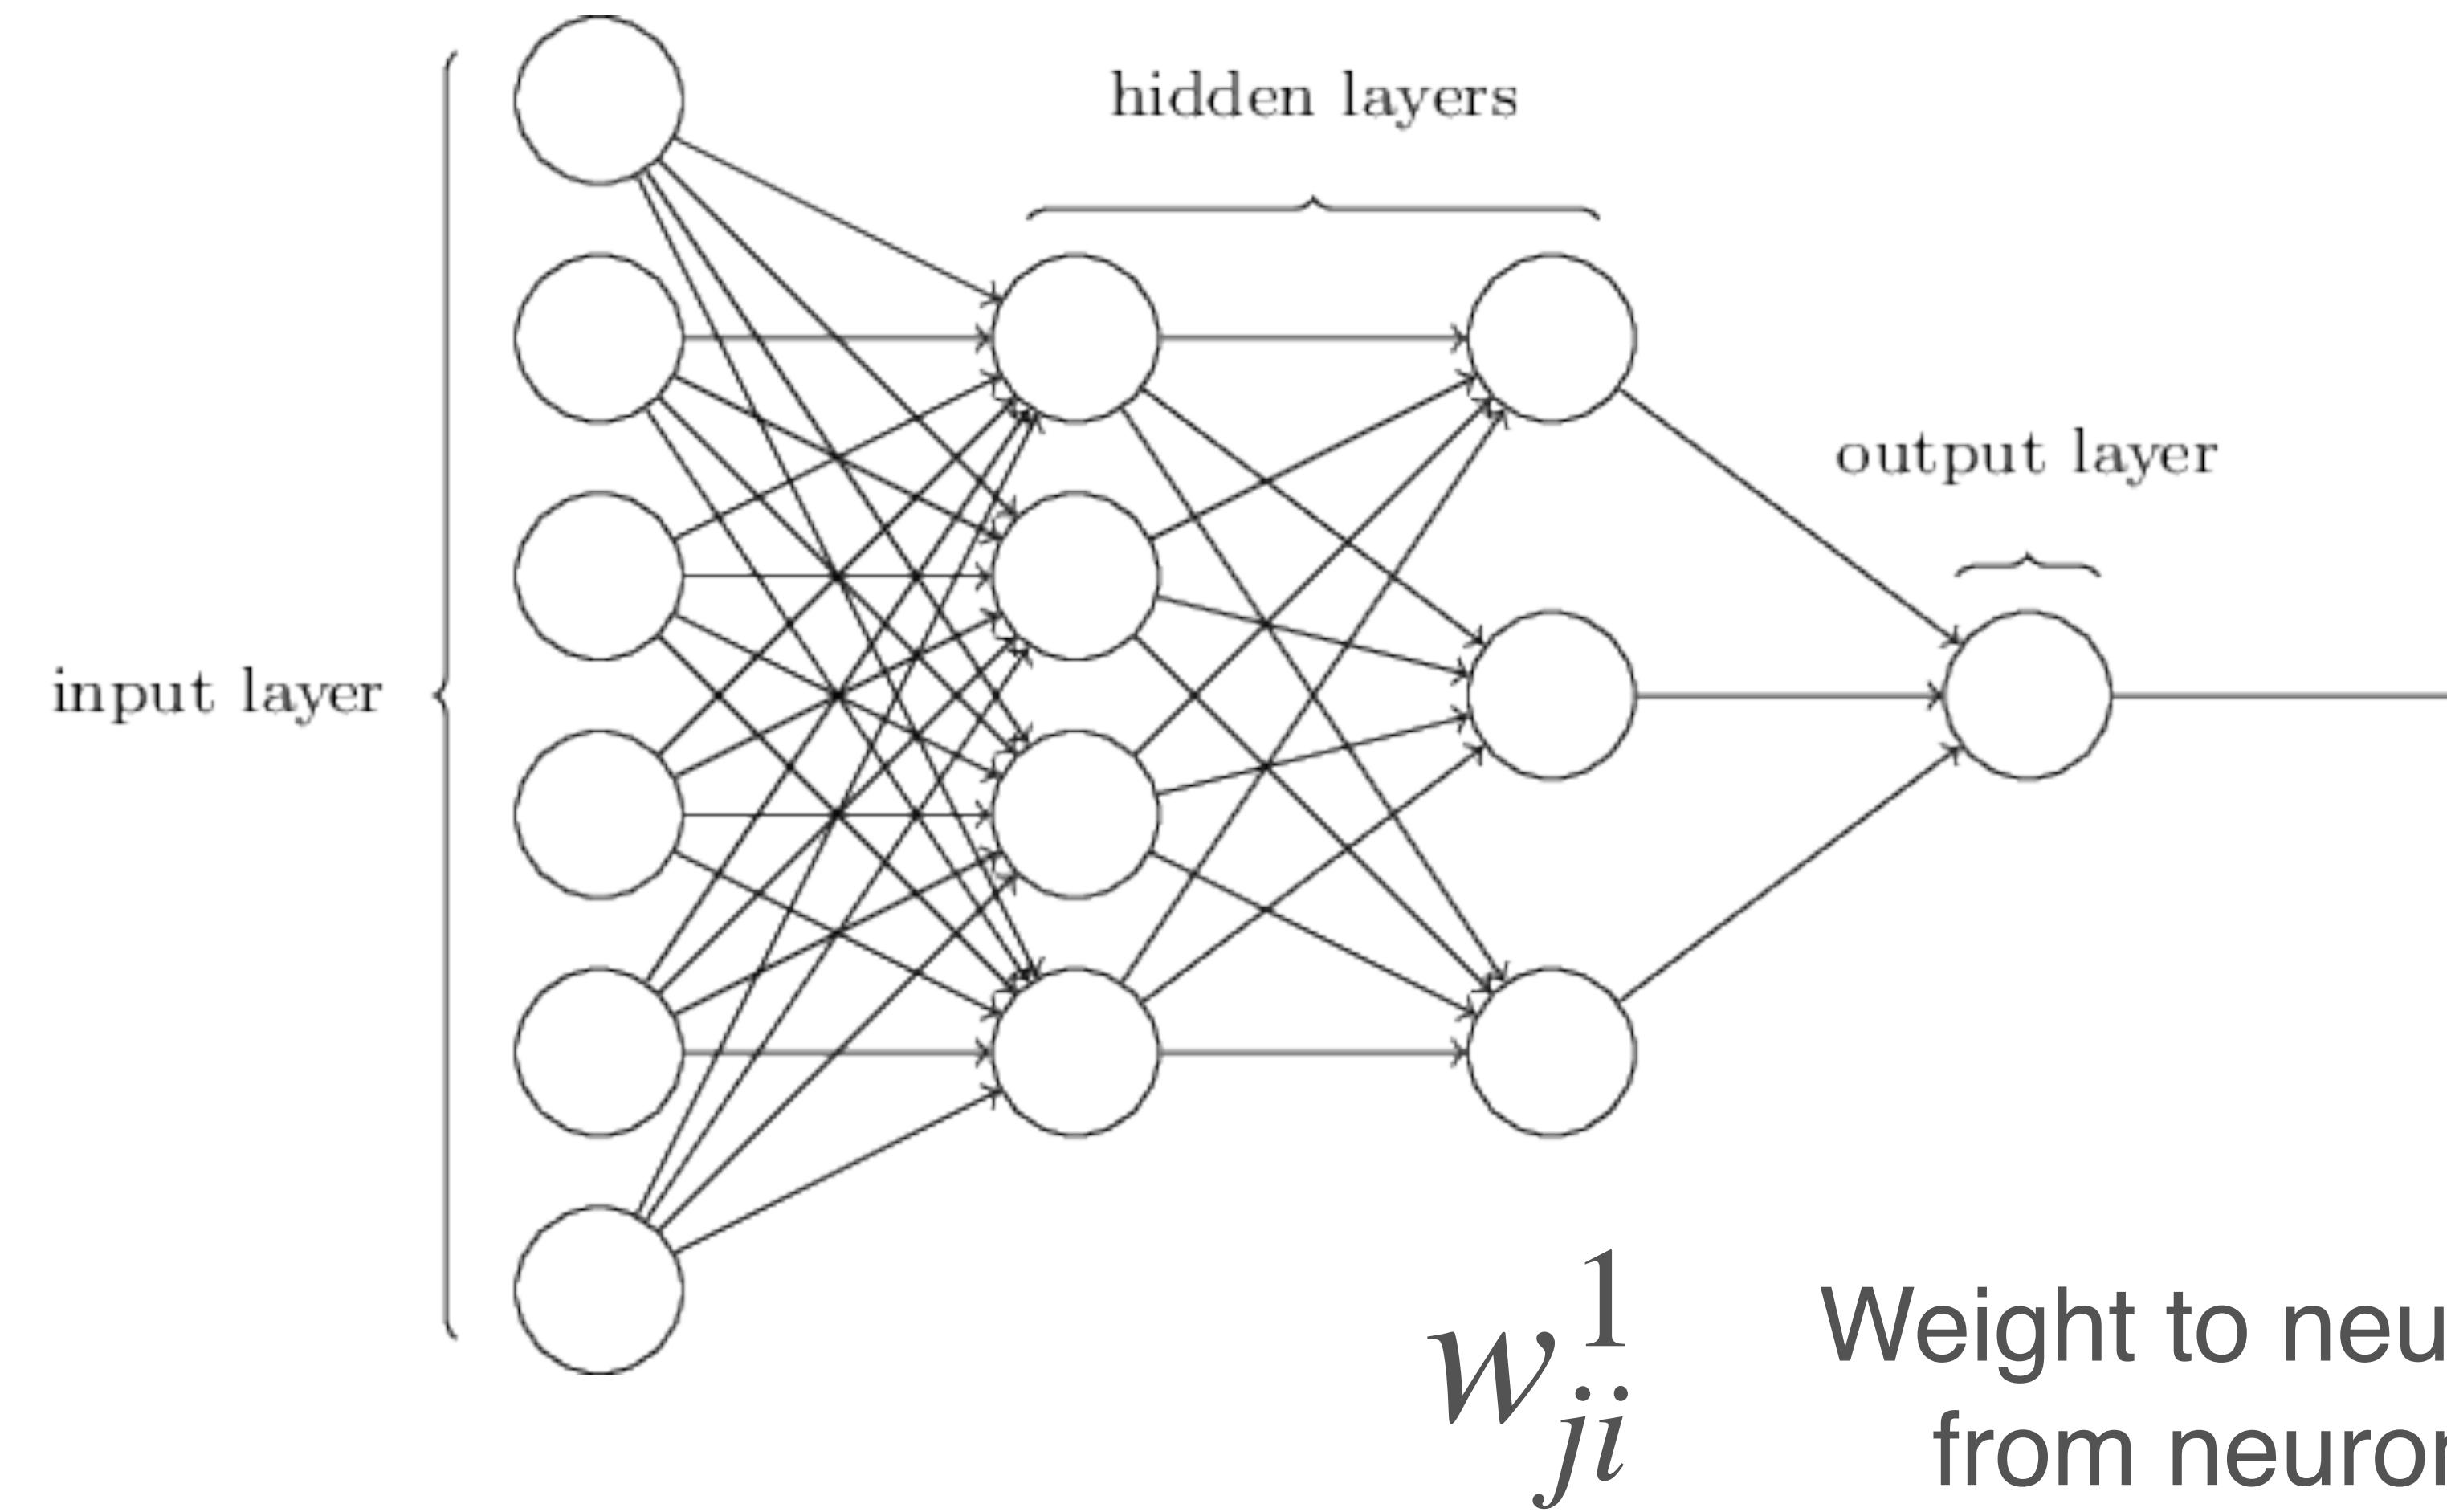
$$\hat{y} = f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right)$$

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

# Some terminology

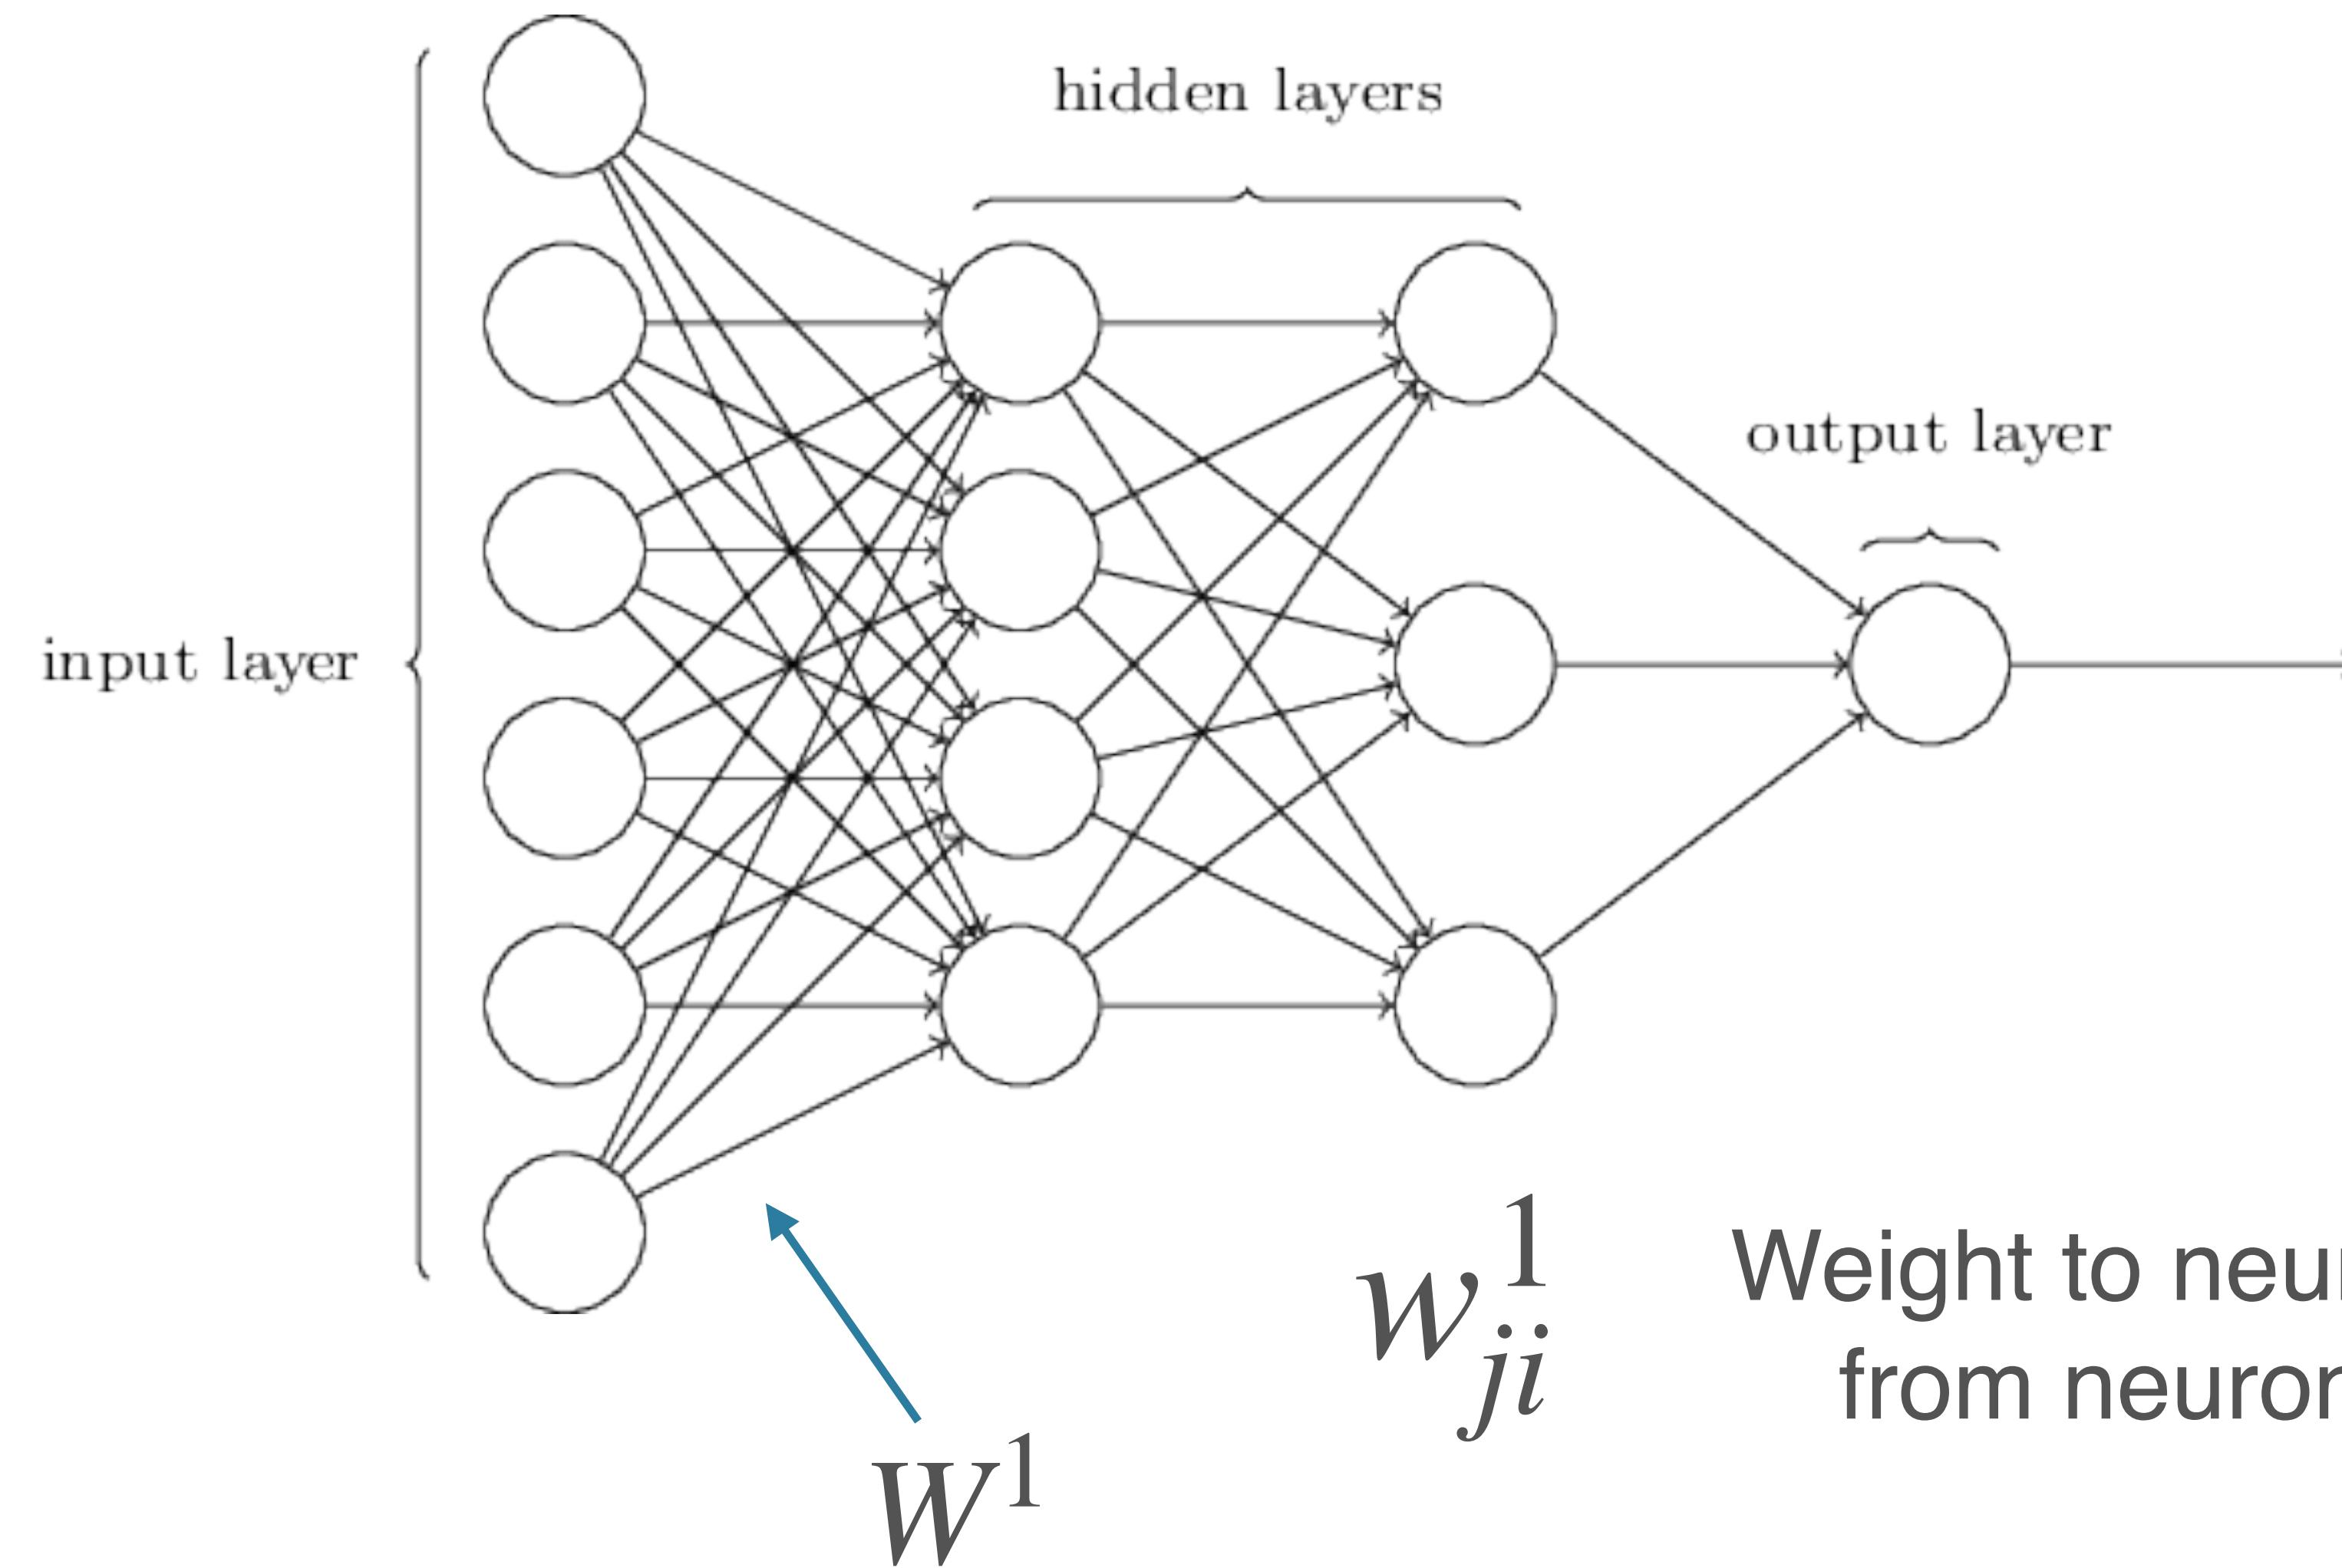
- Our XOR network is a **feed-forward** neural network with **one hidden layer**
  - Aka a multi-layer perceptron (MLP)
- 2 input nodes
- 1 output node
- 1 hidden layer with 2 neurons
- Sigmoid activation function

# General MLP



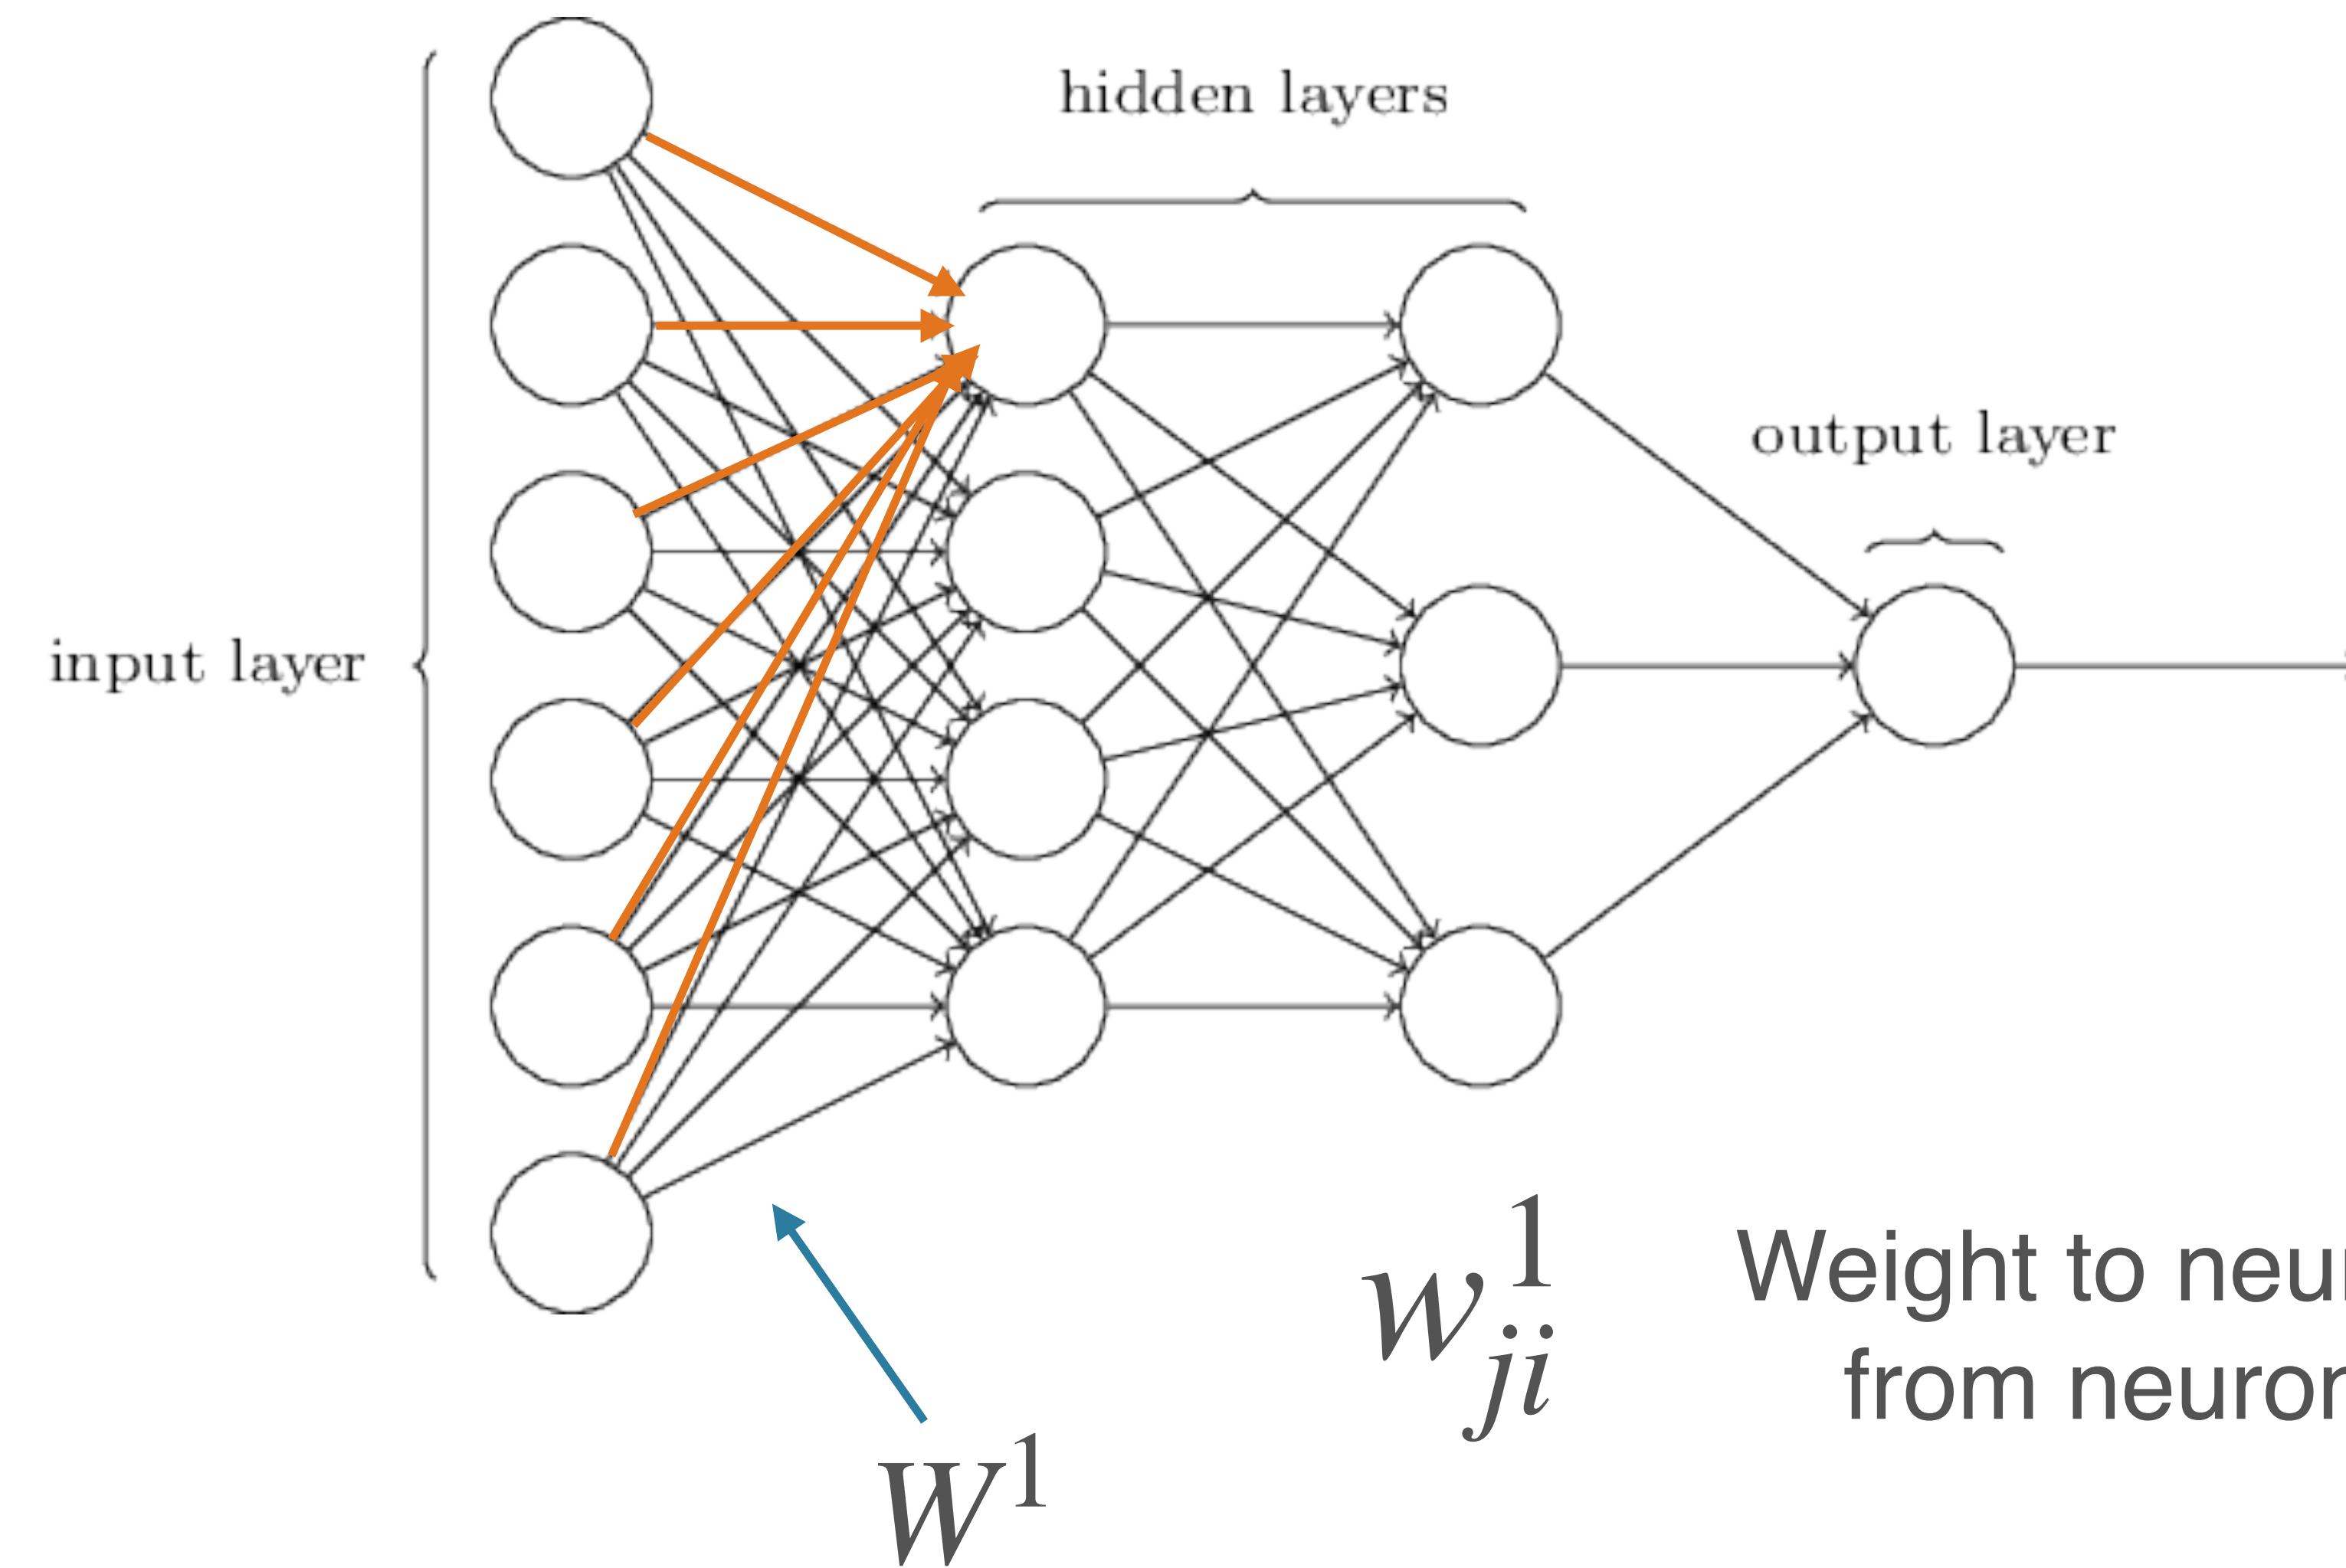
[source](#)

# General MLP



[source](#)

# General MLP



# General MLP

# General MLP

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

# General MLP

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n_0} \end{bmatrix}$$

shape:  $(n_0, 1)$

# General MLP

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n_0} \end{bmatrix}$$

shape:  $(n_0, 1)$

$$W^1 = \begin{bmatrix} w_{00} & w_{10} & \cdots & w_{0n_0} \\ w_{10} & w_{11} & \cdots & w_{1n_0} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n_1 0} & w_{n_1 1} & \cdots & w_{n_1 n_0} \end{bmatrix}$$

shape:  $(n_1, n_0)$

$n_0$ : dimension of input (layer 0)

$n_1$ : output dimension of layer 1

# General MLP

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n_0} \end{bmatrix}$$

shape:  $(n_0, 1)$

$$W^1 = \begin{bmatrix} w_{00} & w_{10} & \cdots & w_{0n_0} \\ w_{10} & w_{11} & \cdots & w_{1n_0} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n_1 0} & w_{n_1 1} & \cdots & w_{n_1 n_0} \end{bmatrix}$$

shape:  $(n_1, n_0)$

$n_0$ : dimension of input (layer 0)  
 $n_1$ : output dimension of layer 1

$$b^1 = \begin{bmatrix} b_0 \\ b_1 \\ \vdots \\ b_{n_1} \end{bmatrix}$$

shape:  $(n_1, 1)$

# Parameters of an MLP

- Weights and biases
  - For **each layer**  $l$ :  $n_l(n_{l-1} + 1)$
  - $n_l n_{l-1}$  weights;  $n_l$  biases
- With  $n$  hidden layers (considering the output as a hidden layer):

$$\sum_{i=1}^n n_i(n_{i-1} + 1)$$

# Hyper-parameters of an MLP

# Hyper-parameters of an MLP

- **Input & output size**

- Usually **fixed by your problem / dataset**
- Input: image size, vocab size; number of “raw” features in general
- Output: 1 for binary classification or simple regression, number of labels for classification, ...

# Hyper-parameters of an MLP

- **Input & output size**
  - Usually **fixed by your problem / dataset**
  - Input: image size, vocab size; number of “raw” features in general
  - Output: 1 for binary classification or simple regression, number of labels for classification, ...
- **Number of hidden layers**

# Hyper-parameters of an MLP

- **Input & output size**
  - Usually **fixed by your problem / dataset**
  - Input: image size, vocab size; number of “raw” features in general
  - Output: 1 for binary classification or simple regression, number of labels for classification, ...
- **Number of hidden layers**
- For each hidden layer:
  - **Size**
  - **Activation function**

# Hyper-parameters of an MLP

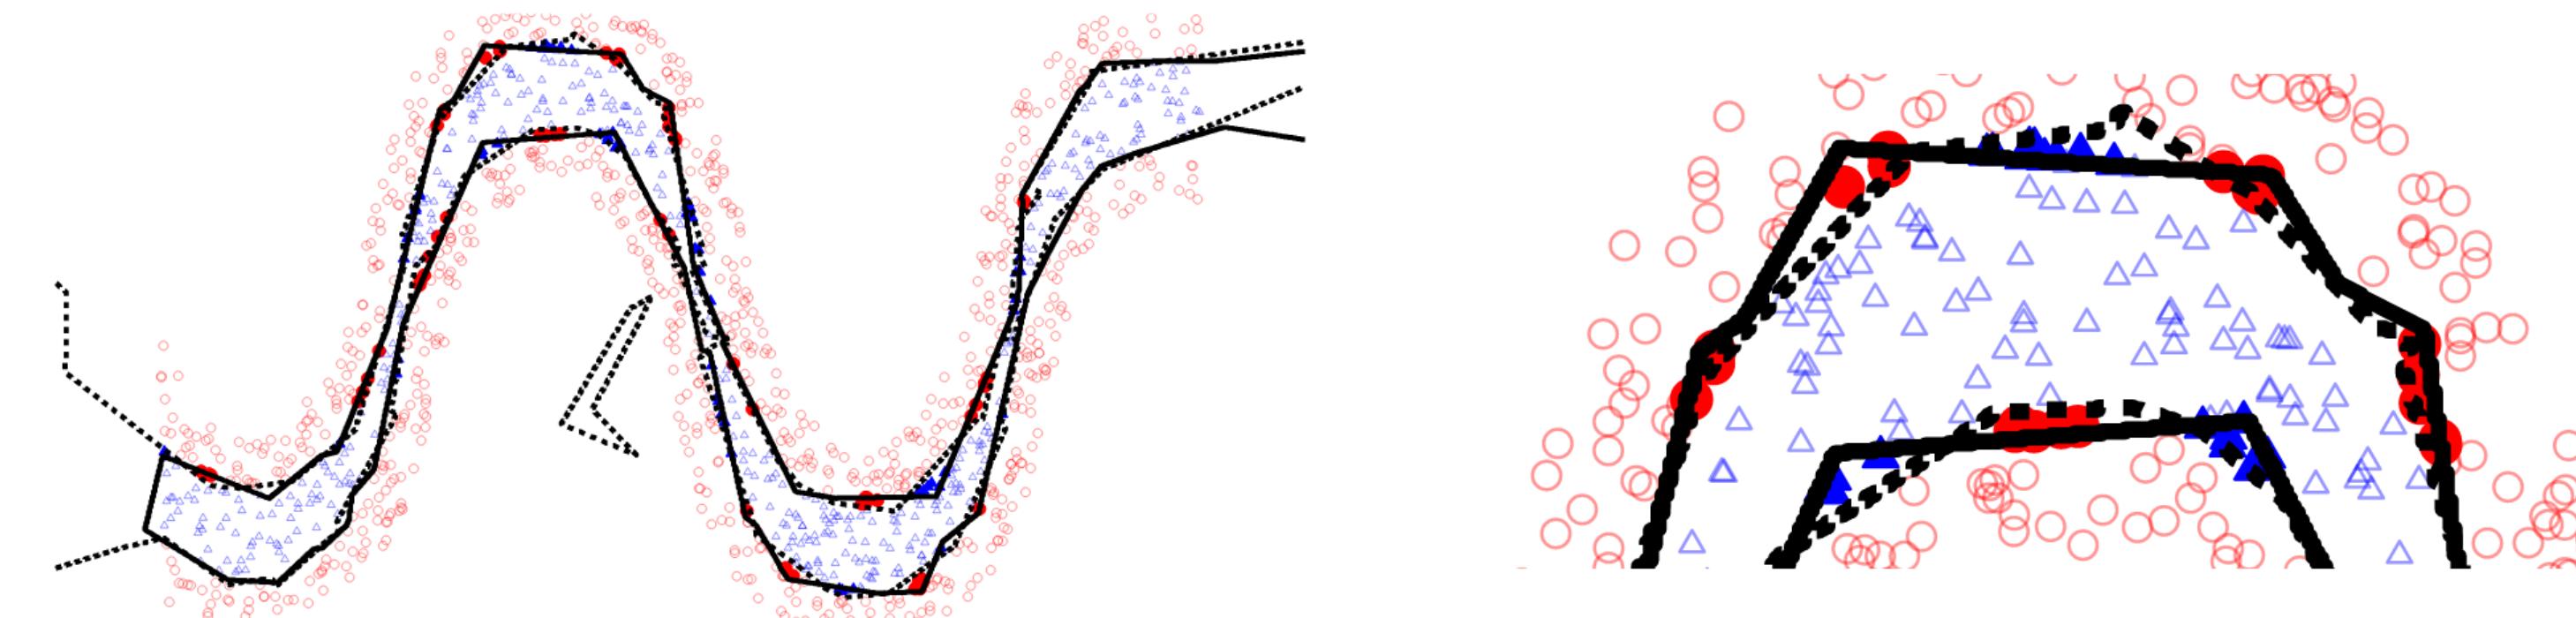
- **Input & output size**
  - Usually **fixed by your problem / dataset**
  - Input: image size, vocab size; number of “raw” features in general
  - Output: 1 for binary classification or simple regression, number of labels for classification, ...
- **Number of hidden layers**
- For each hidden layer:
  - **Size**
  - **Activation function**
- Others: initialization, regularization (and associated values), learning rate / training, ...

# The Deep in Deep Learning

- The Universal Approximation Theorem says that **one hidden layer suffices** for arbitrarily-closely approximating a given function
- Empirical drawbacks: Super-exponentially many neurons; hard to discover
- “Deep and narrow” >> “Shallow and wide” (some theoretical analysis)
  - In principle allows hierarchical features to be learned
  - More well-behaved w/r/t optimization

# The Deep in Deep Learning

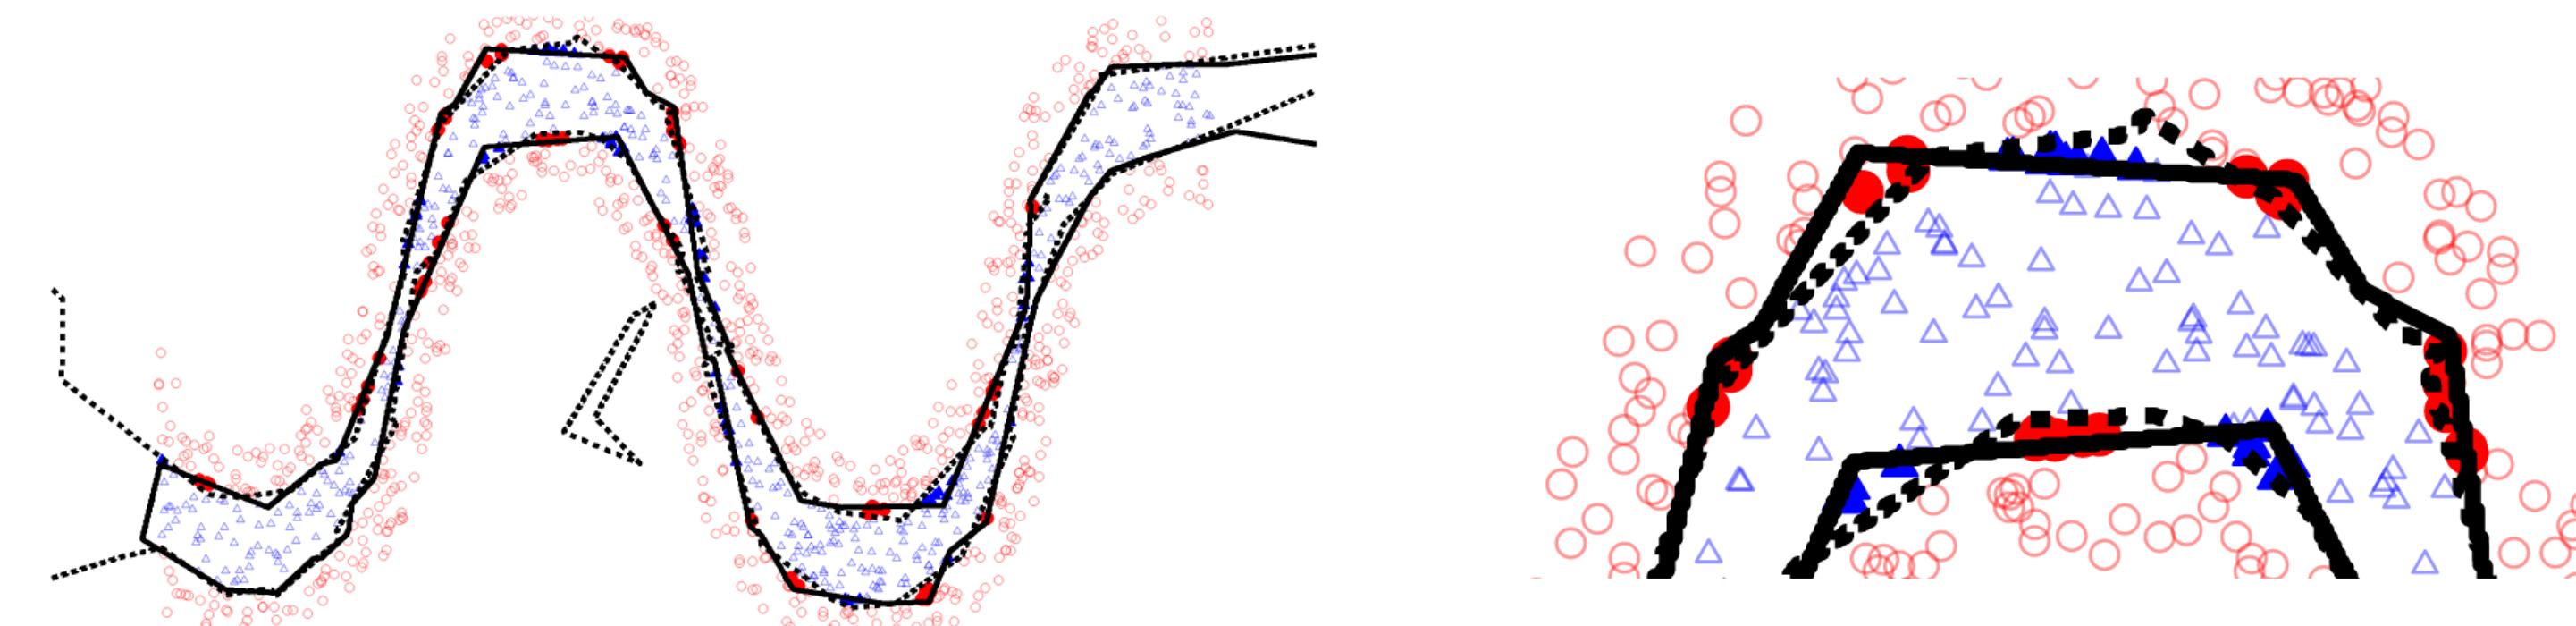
- The Universal Approximation Theorem says that **one hidden layer suffices** for arbitrarily-closely approximating a given function
- Empirical di
- “Deep and i
- In principle
- More well-behaved w/r/t optimization



# The Deep in Deep Learning

- The Universal Approximation Theorem says that **one hidden layer suffices** for arbitrarily-closely approximating a given function

- Empirical di

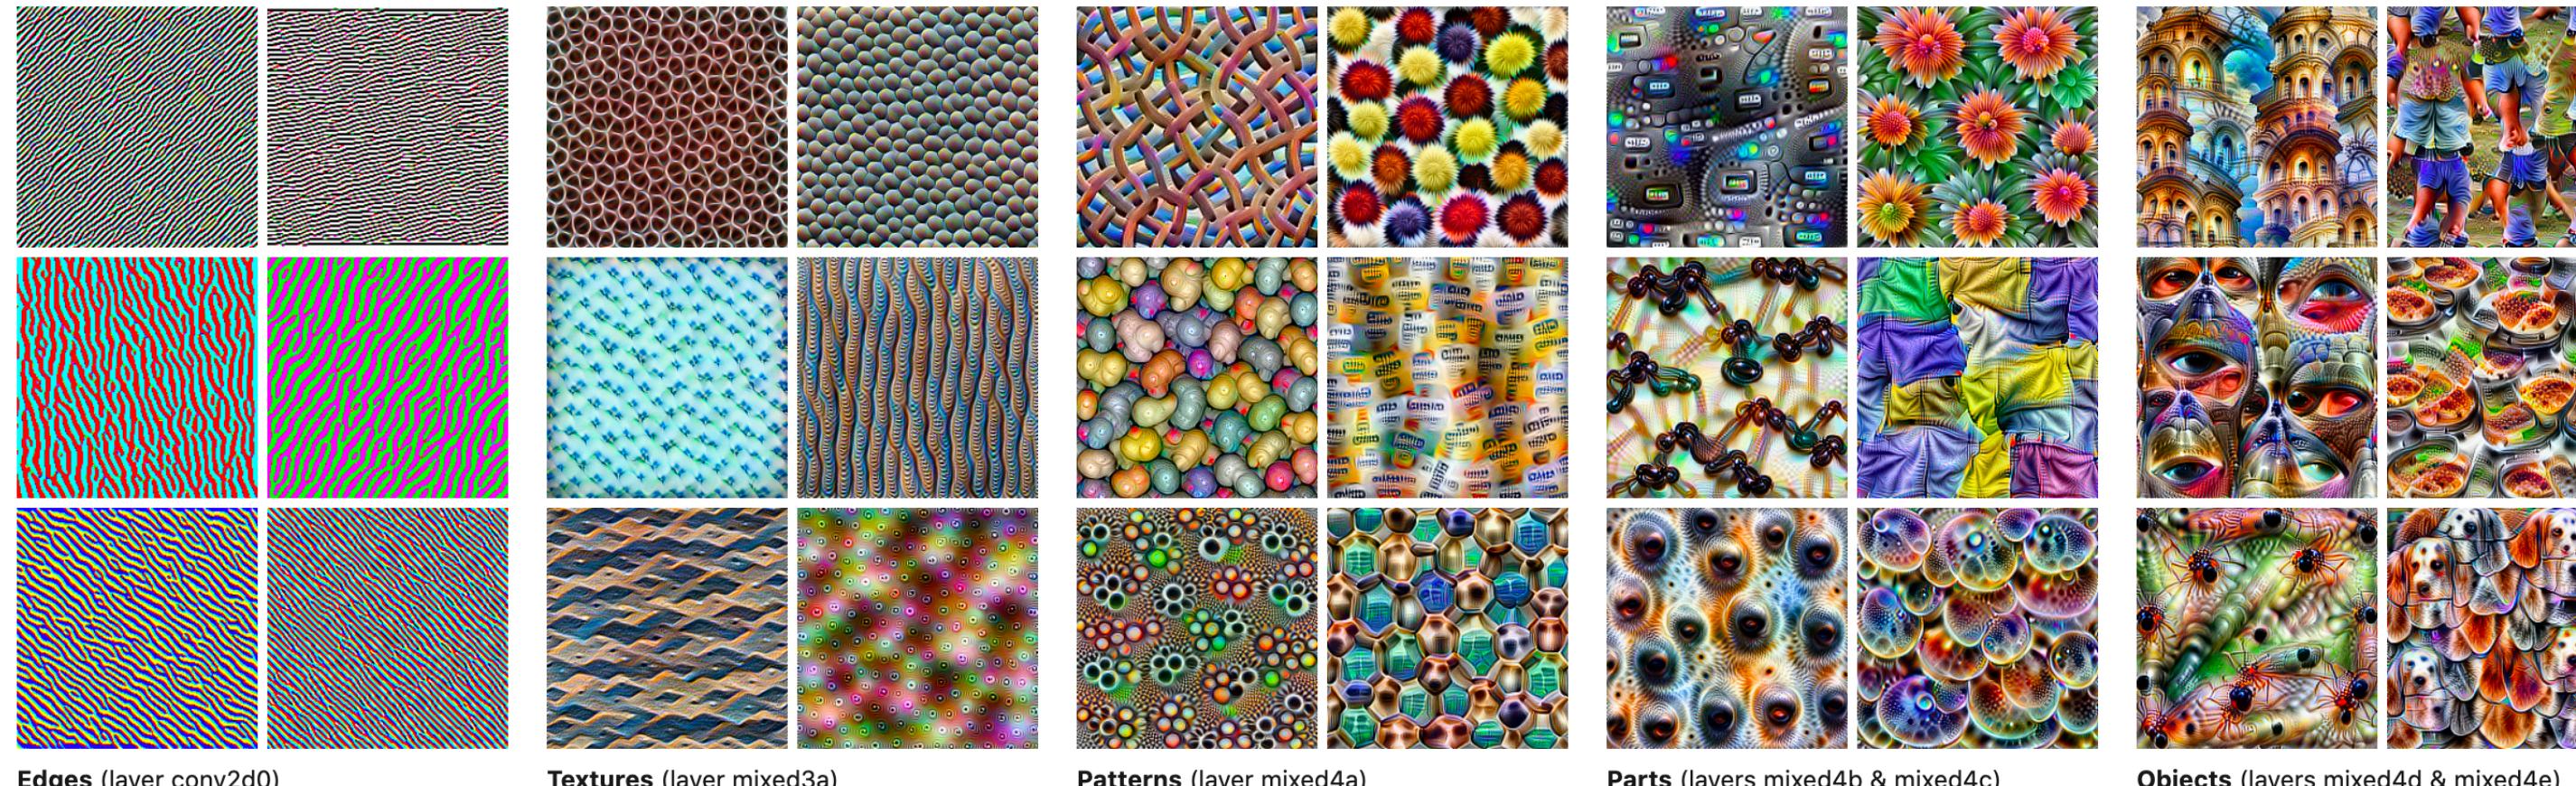


discover  
source  
sis)

- “Deep and i

- In principle

- More well-behaved w/r/t optimization



source

# Activation Functions

- **Non-linear** activation functions are essential
- MLP: linear transformation, followed by a point-wise non-linearity, repeated several times over
- Without the non-linearity, would just have several linear transformations
  - **Composition** of linear transformations is **also linear!**

# Activation Functions

- **Non-linear** activation functions are essential
- MLP: linear transformation, followed by a point-wise non-linearity, repeated several times over
- Without the non-linearity, would just have several linear transformations
  - **Composition** of linear transformations is **also linear!**

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

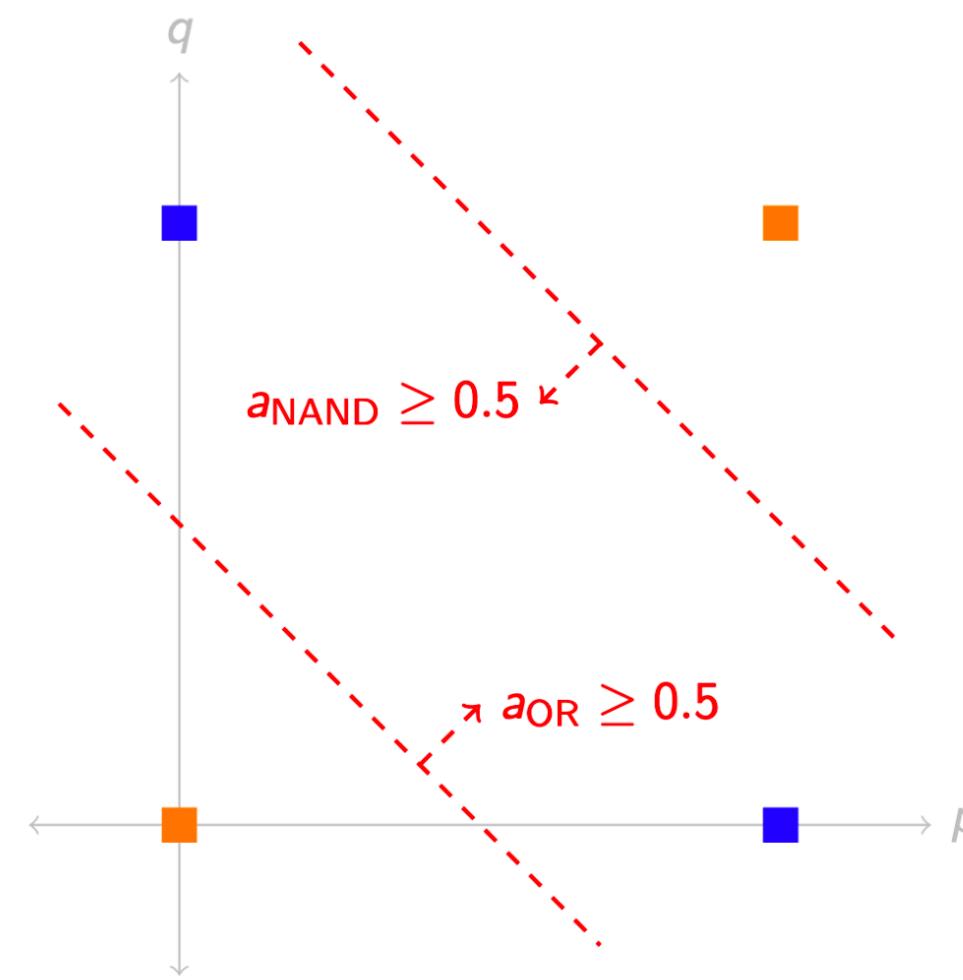
# Non-linearity, cont.

# Non-linearity, cont.

- Recall: XOR was not computable by a single neuron because the latter can only compute *linearly separable* functions

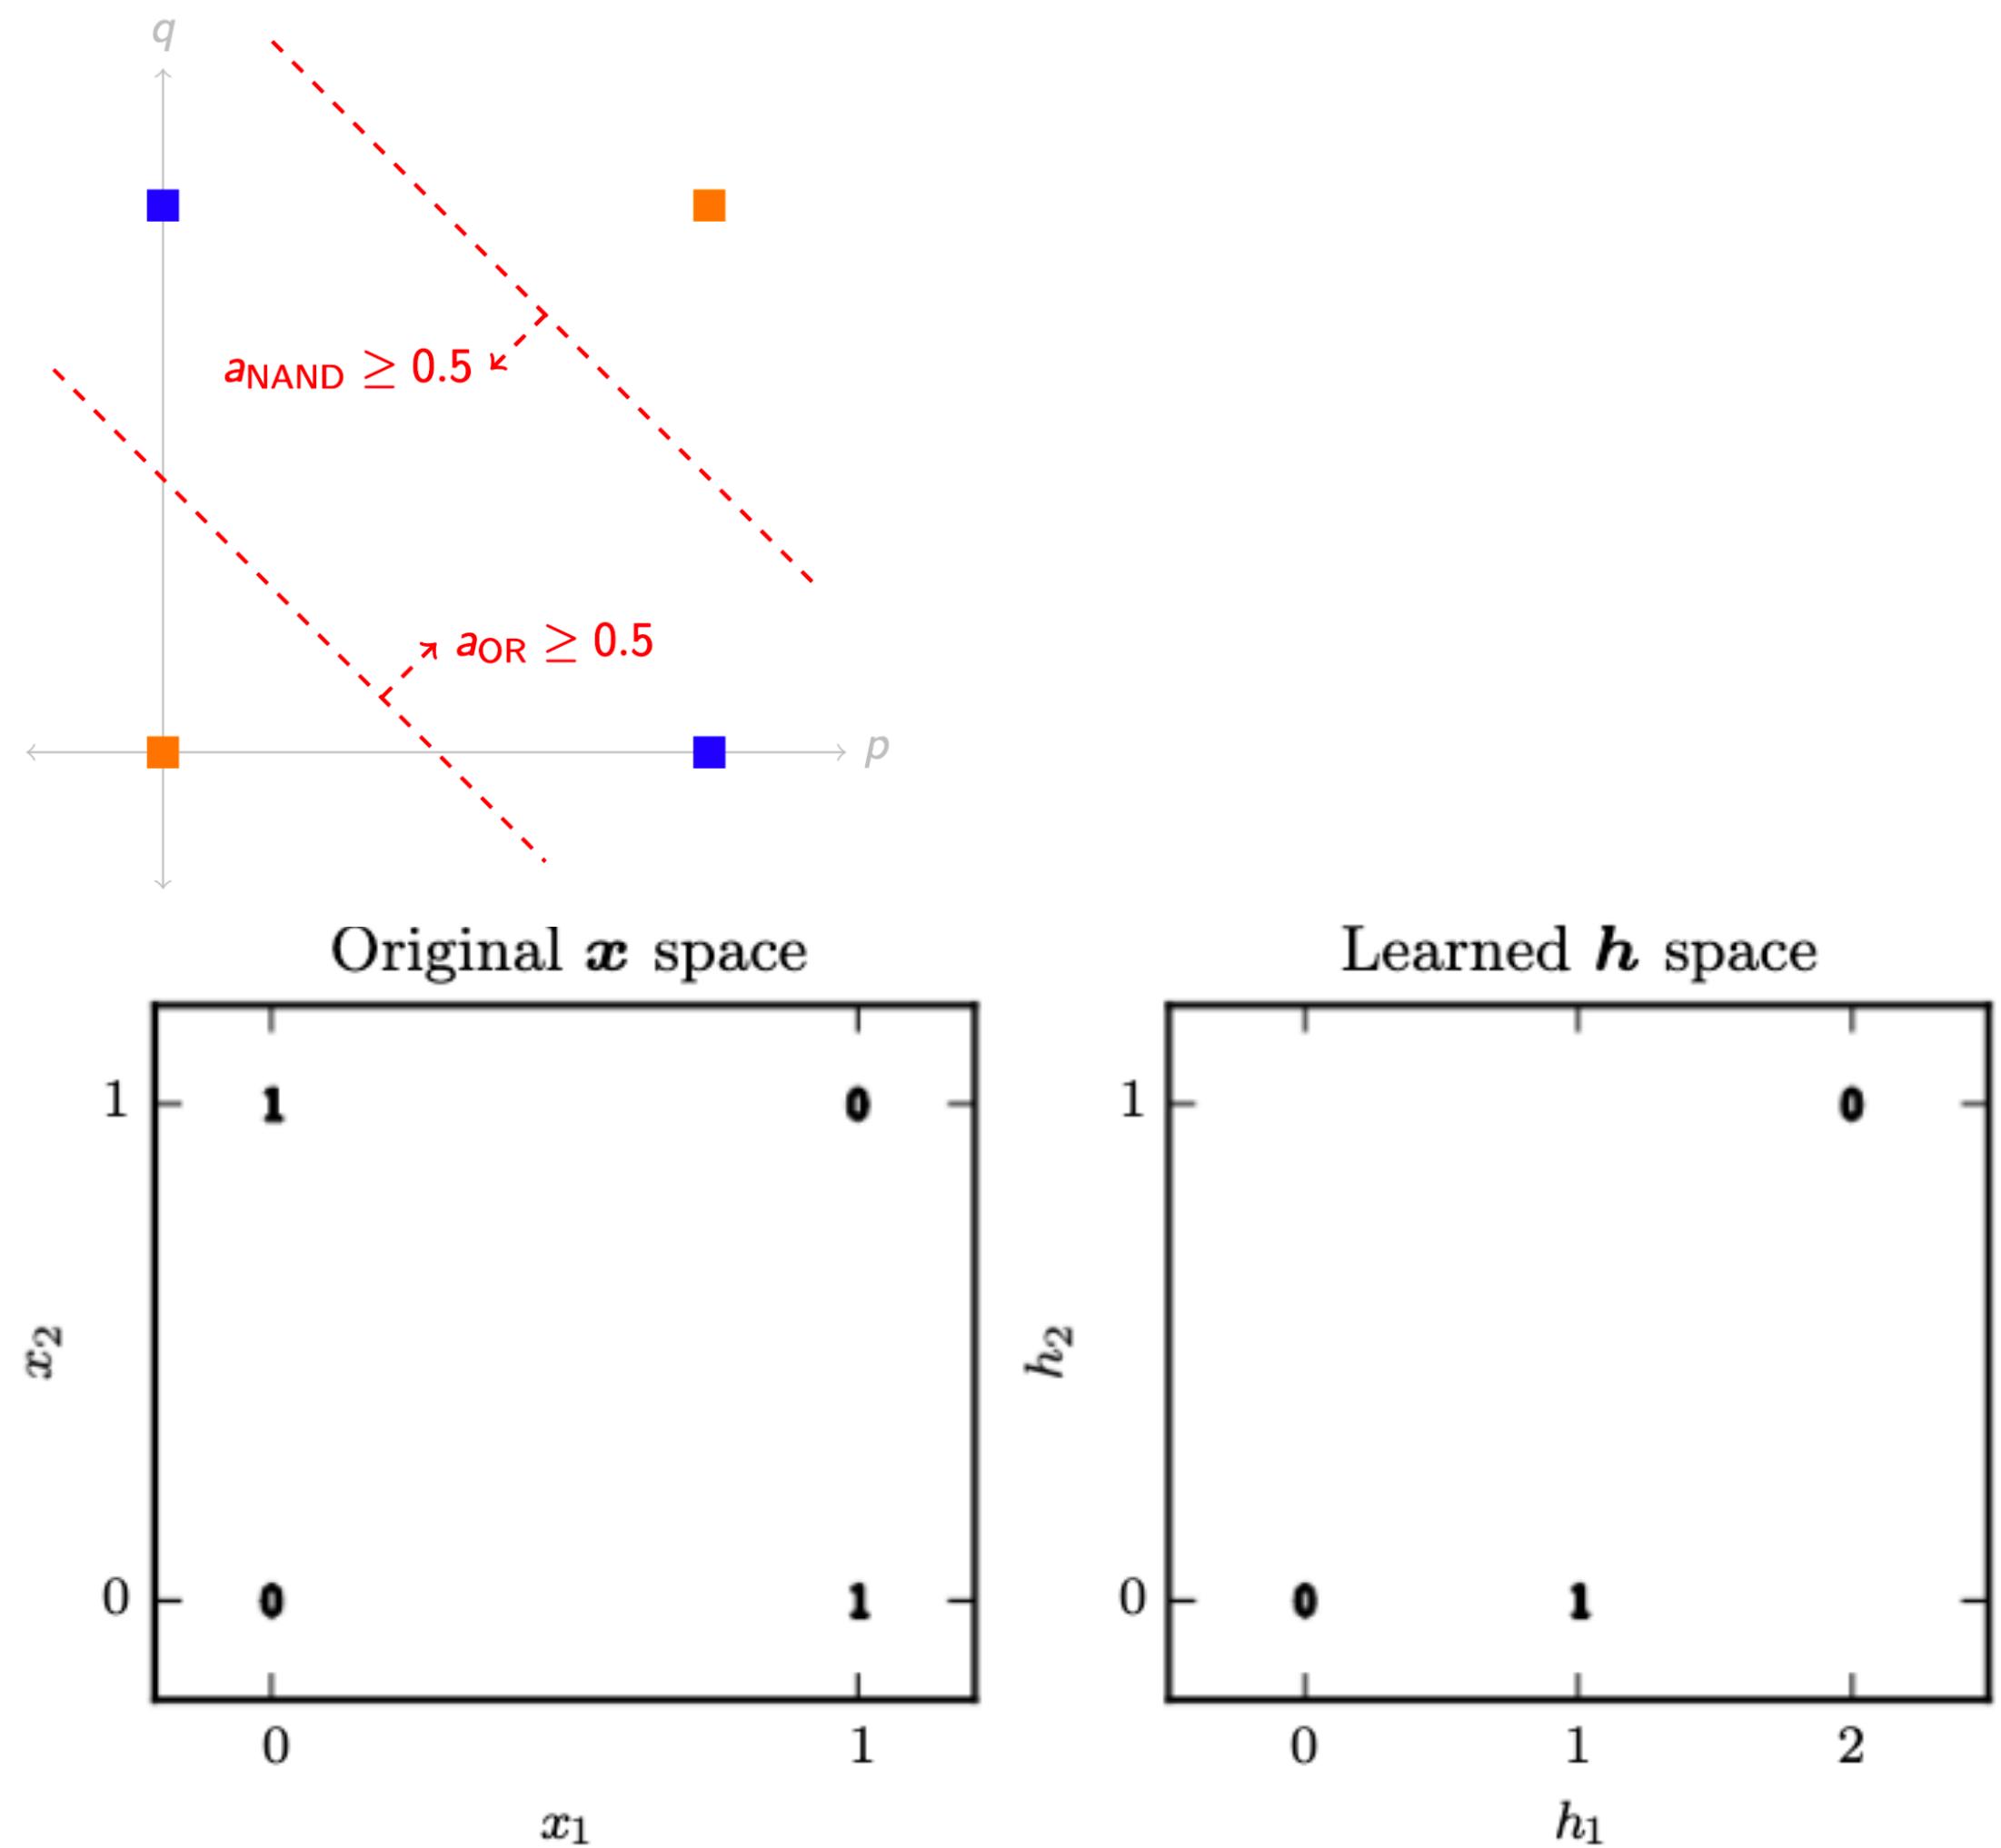
# Non-linearity, cont.

- Recall: XOR was not computable by a single neuron because the latter can only compute *linearly separable* functions
- One perspective: integrating extracted features



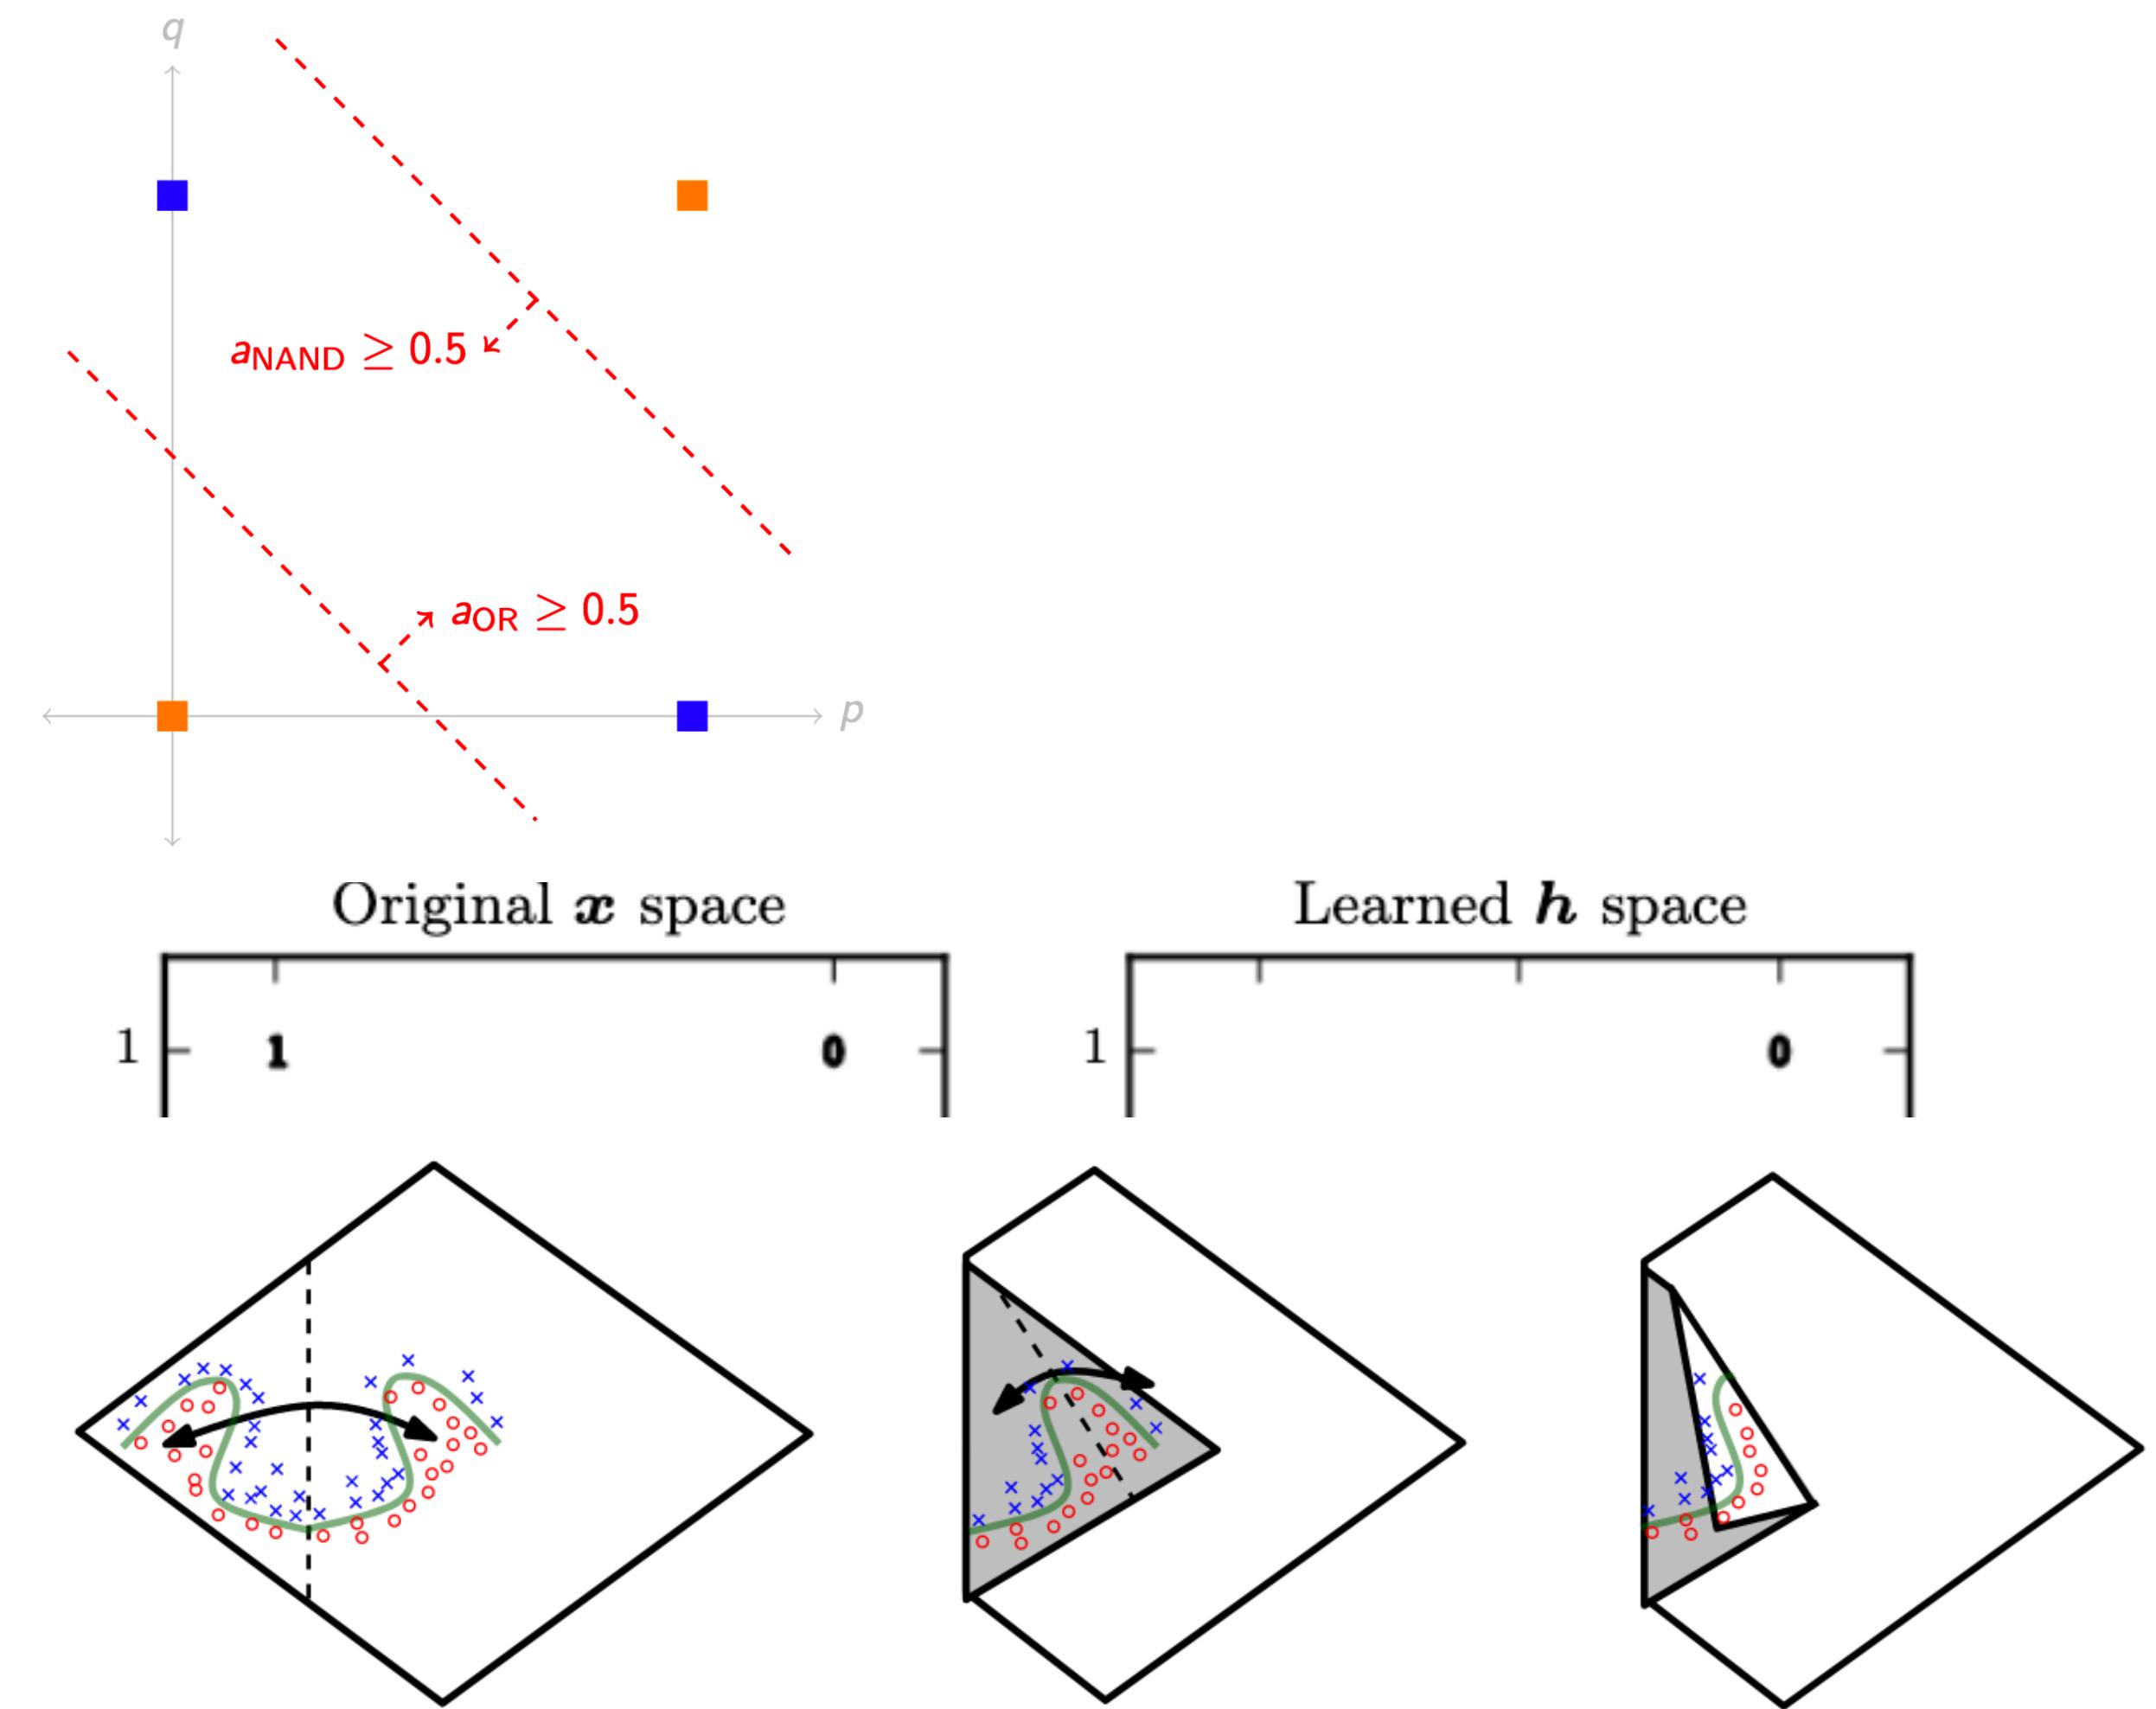
# Non-linearity, cont.

- Recall: XOR was not computable by a single neuron because the latter can only compute *linearly separable* functions
- One perspective: integrating extracted features
- An equivalent perspective:
  - Transforming the input space ([source](#); p. 169)
  - This is a *non-linear* transformation
  - Space folding intuition more generally (also GBC sec 6.4.1)



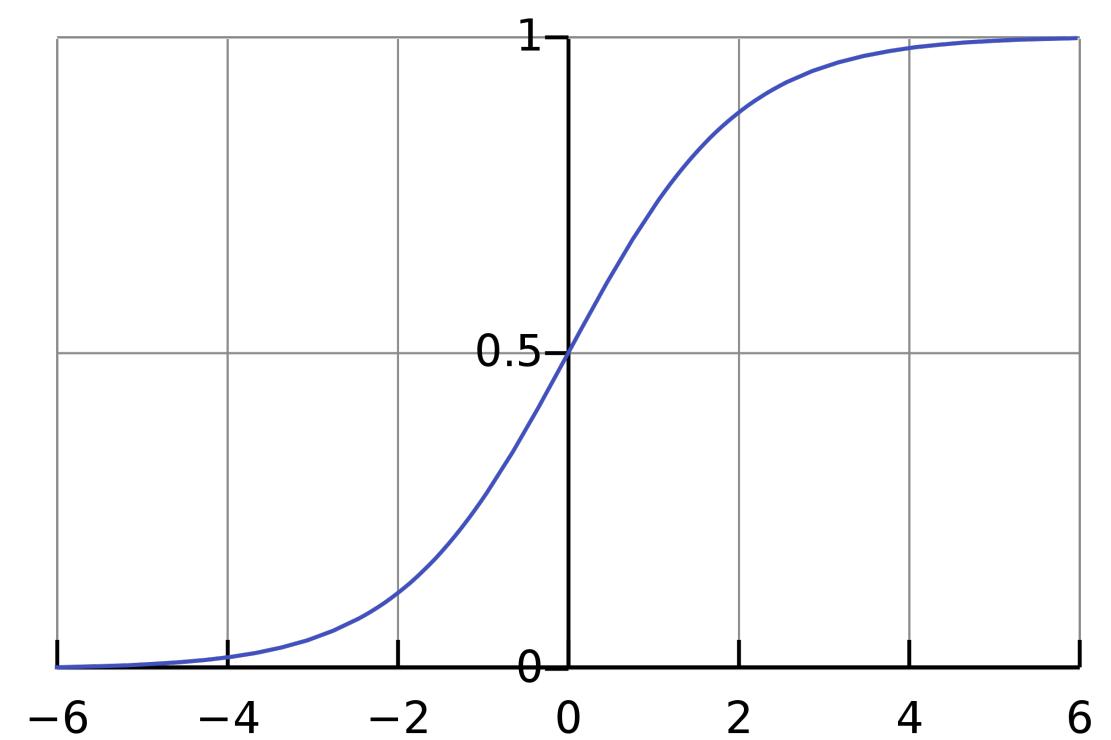
# Non-linearity, cont.

- Recall: XOR was not computable by a single neuron because the latter can only compute *linearly separable* functions
- One perspective: integrating extracted features
- An equivalent perspective:
  - Transforming the input space ([source](#); p. 169)
  - This is a *non-linear* transformation
  - Space folding intuition more generally (also GBC sec 6.4.1)

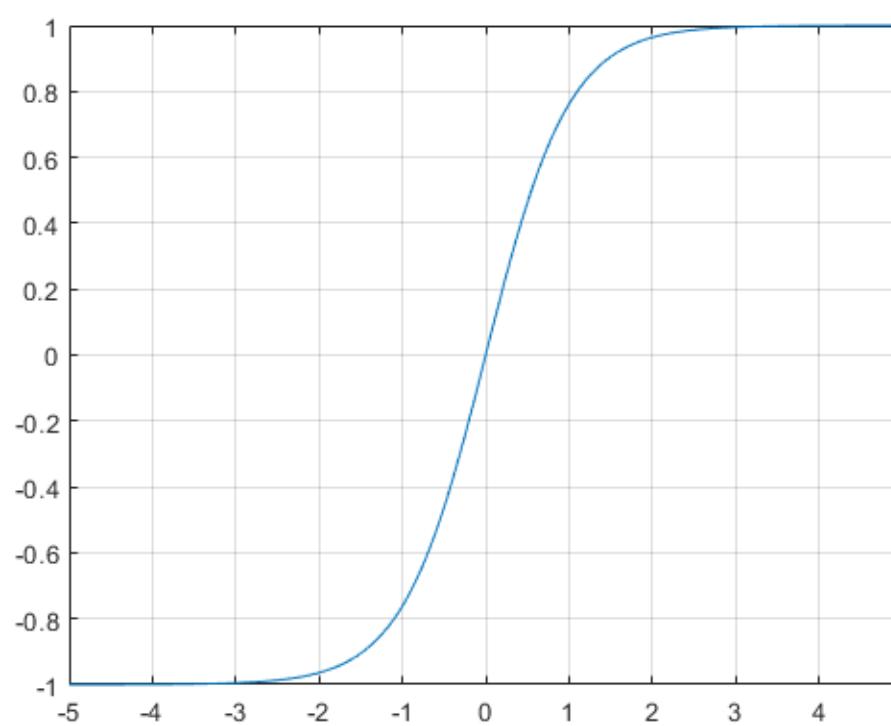


# Activation Functions: Hidden Layer

sigmoid



tanh

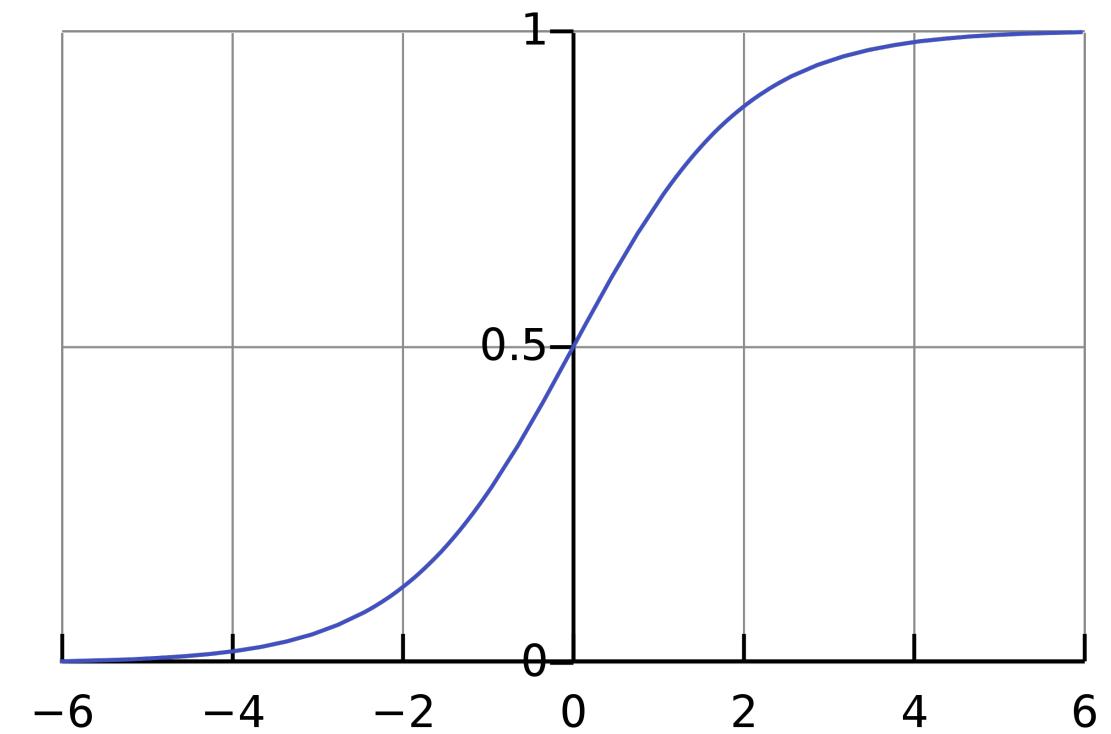


$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

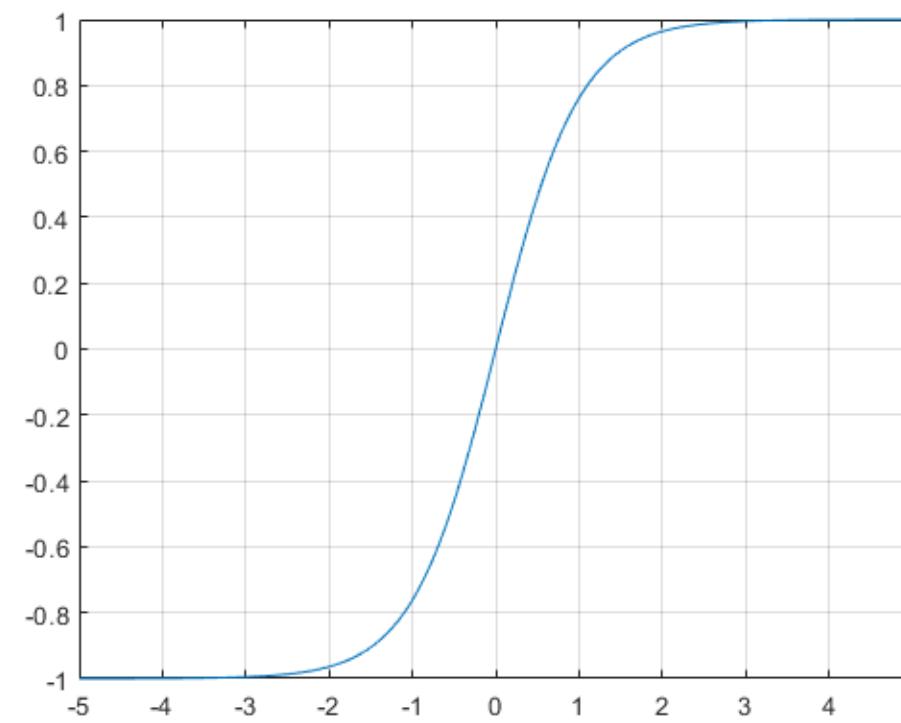
$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2\sigma(2x) - 1$$

# Activation Functions: Hidden Layer

sigmoid



tanh



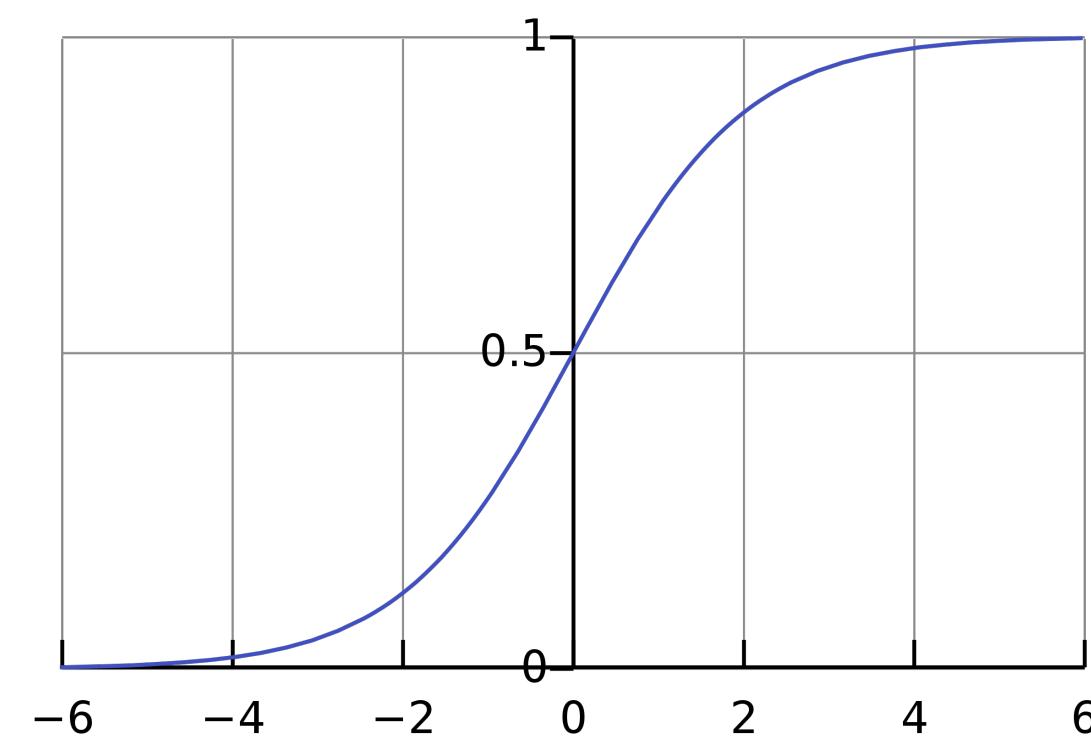
$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2\sigma(2x) - 1$$

Problem with these two: derivative  
“saturates” (nearly 0) everywhere except  
near origin

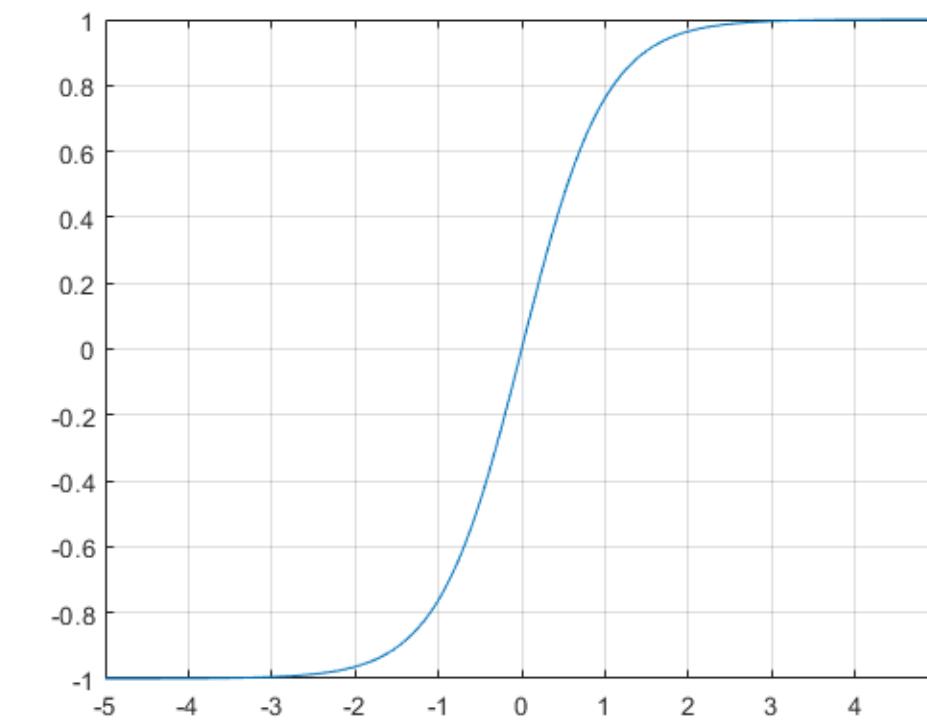
# Activation Functions: Hidden Layer

sigmoid



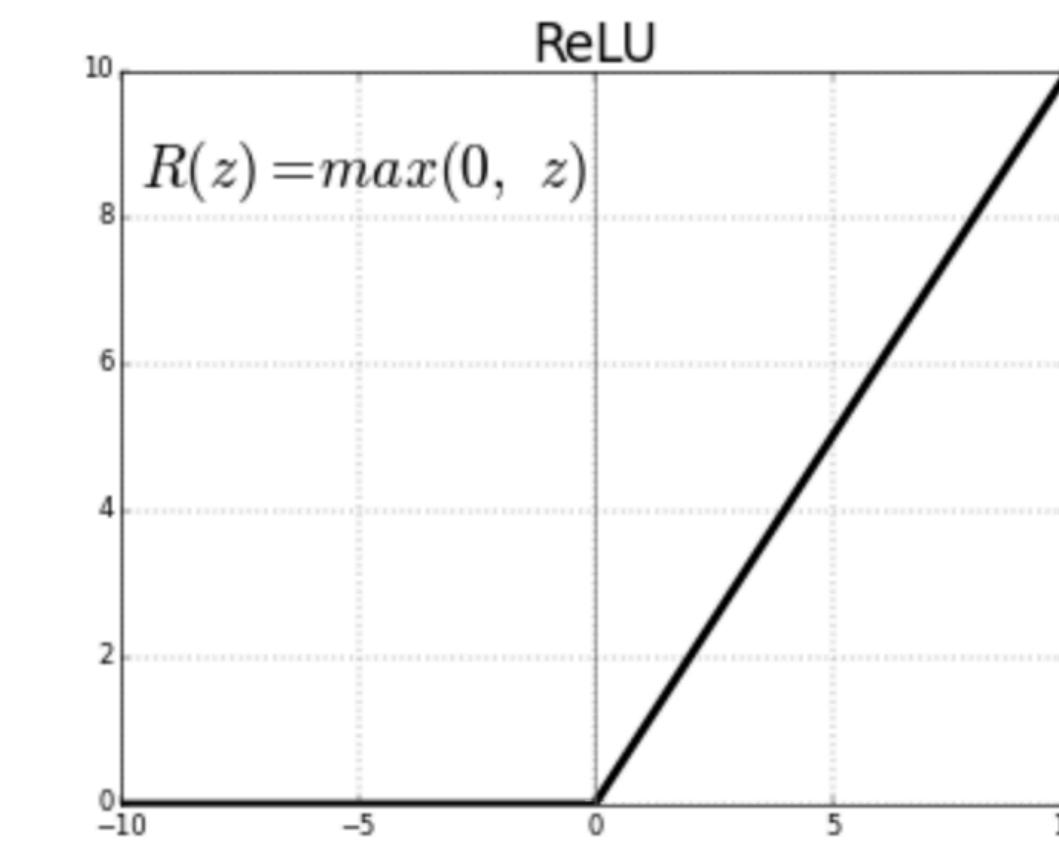
$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

tanh



$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2\sigma(2x) - 1$$

Problem with these two: derivative “saturates” (nearly 0) everywhere except near origin



ReLU does not saturate. Good “default”

# Activation Functions: Output Layer

- Depends on the task!
- Regression (continuous output(s)): **none!**
  - Just use final linear transformation
- Binary classification: **sigmoid**
  - Also for *multi-label* classification
- Multi-class classification: **softmax**
  - Terminology: the inputs to a softmax are called **logits**
  - (there are sometimes other uses of the term, so beware)

$$\text{softmax}(x)_i = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

# Mini-batch computation

# Computing with a Single Input

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$



# Computing with a Single Input

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n_0} \end{bmatrix}$$

Shape:  $(n_0, 1)$

# Computing with a Single Input

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \cdots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \cdots \right) + b^n \right)$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n_0} \end{bmatrix}$$

Shape:  $(n_0, 1)$

$$W^1 = \begin{bmatrix} w_{00} & w_{10} & \cdots & w_{0n_0} \\ w_{10} & w_{11} & \cdots & w_{1n_0} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n_1 0} & w_{n_1 1} & \cdots & w_{n_1 n_0} \end{bmatrix}$$

Shape:  $(n_1, n_0)$

$n_0$ : dimension of input (layer 0)

$n_1$ : output dimension of layer 1

# Computing with a Single Input

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \cdots f_2 \left( W^2 \cdot f_1 \left( W^1 x + b^1 \right) + b^2 \right) \cdots \right) + b^n \right)$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n_0} \end{bmatrix}$$

Shape:  $(n_0, 1)$

$$W^1 = \begin{bmatrix} w_{00} & w_{10} & \cdots & w_{0n_0} \\ w_{10} & w_{11} & \cdots & w_{1n_0} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n_1 0} & w_{n_1 1} & \cdots & w_{n_1 n_0} \end{bmatrix}$$

Shape:  $(n_1, n_0)$

$n_0$ : dimension of input (layer 0)  
 $n_1$ : output dimension of layer 1

$$b^1 = \begin{bmatrix} b_0 \\ b_1 \\ \vdots \\ b_{n_1} \end{bmatrix}$$

Shape:  $(n_1, 1)$

# Mini-batch Gradient Descent

initialize parameters / build model

for each epoch:

```
data = shuffle(data)
batches = make_batches(data)
```

for each batch in batches:

```
outputs = model(batch)
loss = loss_fn(outputs, true_outputs)
compute gradients
update parameters
```

# Computing with Mini-batches

- Bad idea:

```
for each batch in batches:  
    for each datum in batch:  
        outputs = model(datum)  
        loss = loss_fn(outputs, true_outputs)  
        compute gradients  
        update parameters
```

# Computing with a Batch of Inputs

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 X + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

# Computing with a Batch of Inputs

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \dots f_2 \left( W^2 \cdot f_1 \left( W^1 X + b^1 \right) + b^2 \right) \dots \right) + b^n \right)$$

$$X = \begin{bmatrix} x_0^0 & x_0^1 & \dots & x_0^k \\ x_1^0 & x_1^1 & \dots & x_1^k \\ \vdots & \vdots & \ddots & \vdots \\ x_{n_0}^0 & x_{n_0}^1 & \dots & x_{n_0}^k \end{bmatrix}$$

Shape:  $(n_0, k)$   
k: batch\_size

# Computing with a Batch of Inputs

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \cdots f_2 \left( W^2 \cdot f_1 \left( W^1 X + b^1 \right) + b^2 \right) \cdots \right) + b^n \right)$$

$$X = \begin{bmatrix} x_0^0 & x_0^1 & \dots & x_0^k \\ x_1^0 & x_1^1 & \dots & x_1^k \\ \vdots & \vdots & \ddots & \vdots \\ x_{n_0}^0 & x_{n_0}^1 & \dots & x_{n_0}^k \end{bmatrix} \quad W^1 = \begin{bmatrix} w_{00} & w_{01} & \cdots & w_{0n_1} \\ w_{10} & w_{11} & \cdots & w_{1n_1} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n_00} & w_{n_01} & \cdots & w_{n_0n_1} \end{bmatrix}$$

Shape:  $(n_0, k)$   
k: batch\_size

# Computing with a Batch of Inputs

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \cdots f_2 \left( W^2 \cdot f_1 \left( W^1 X + b^1 \right) + b^2 \right) \cdots \right) + b^n \right)$$

$$X = \begin{bmatrix} x_0^0 & x_0^1 & \dots & x_0^k \\ x_1^0 & x_1^1 & \dots & x_1^k \\ \vdots & \vdots & \ddots & \vdots \\ x_{n_0}^0 & x_{n_0}^1 & \dots & x_{n_0}^k \end{bmatrix} \quad W^1 = \begin{bmatrix} w_{00} & w_{01} & \cdots & w_{0n_1} \\ w_{10} & w_{11} & \cdots & w_{1n_1} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n_00} & w_{n_01} & \cdots & w_{n_0n_1} \end{bmatrix}$$

Shape:  $(n_0, k)$   
k: batch\_size

Shape:  $(n_1, n_0)$   
 $n_0$ : dimension of input (layer 0)  
 $n_1$ : output dimension of layer 1

# Computing with a Batch of Inputs

$$\hat{y} = f_n \left( W^n \cdot f_{n-1} \left( \cdots f_2 \left( W^2 \cdot f_1 \left( W^1 X + b^1 \right) + b^2 \right) \cdots \right) + b^n \right)$$

$$X = \begin{bmatrix} x_0^0 & x_0^1 & \dots & x_0^k \\ x_1^0 & x_1^1 & \dots & x_1^k \\ \vdots & \vdots & \ddots & \vdots \\ x_{n_0}^0 & x_{n_0}^1 & \dots & x_{n_0}^k \end{bmatrix}$$

Shape:  $(n_0, k)$   
k: batch\_size

$$W^1 = \begin{bmatrix} w_{00} & w_{01} & \cdots & w_{0n_1} \\ w_{10} & w_{11} & \cdots & w_{1n_1} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n_00} & w_{n_01} & \cdots & w_{n_0n_1} \end{bmatrix}$$

Shape:  $(n_1, n_0)$   
 $n_0$ : dimension of input (layer 0)  
 $n_1$ : output dimension of layer 1

$$b^1 = \begin{bmatrix} b_0 \\ b_1 \\ \vdots \\ b_{n_1} \end{bmatrix}$$

Shape:  $(n_1, 1)$   
Added to each col. of  $W^1 X$

# Note on mini-batches and shape

# Note on mini-batches and shape

- Most modern neural net libraries (e.g. PyTorch) expect the **first dimension** of matrices/tensors to be a **batch size**
  - Produce a sequence of representations, *for each item* in the batch
  - e.g.  $(\text{batch\_size}, \text{input\_size}) \rightarrow (\text{batch\_size}, \text{hidden\_size}) \rightarrow (\text{batch\_size}, \text{output\_size})$

# Note on mini-batches and shape

- Most modern neural net libraries (e.g. PyTorch) expect the **first dimension** of matrices/tensors to be a **batch size**
  - Produce a sequence of representations, *for each item* in the batch
  - e.g. (batch\_size, input\_size) —> (batch\_size, hidden\_size) —> (batch\_size, output\_size)
- In principle, can be **higher than 2-dimensional** (tensor rather than just matrix)
  - Images: (batch\_size, width, height, 3)
  - Sequences: (batch\_size, seq\_len, representation\_size)

# Note on mini-batches and shape

- Most modern neural net libraries (e.g. PyTorch) expect the **first dimension** of matrices/tensors to be a **batch size**
  - Produce a sequence of representations, *for each item* in the batch
  - e.g. (batch\_size, input\_size) —> (batch\_size, hidden\_size) —> (batch\_size, output\_size)
- In principle, can be **higher than 2-dimensional** (tensor rather than just matrix)
  - Images: (batch\_size, width, height, 3)
  - Sequences: (batch\_size, seq\_len, representation\_size)
- Two comments:
  - In your code, **annotate every tensor** with a comment showing **intended shape**
  - When debugging, look at shapes early on!!

# Note on mini-batches and shape

# Note on mini-batches and shape

- **Warning:** PyTorch / NN libraries typically multiply matrices in the **opposite direction** as Linear Algebra notation

# Note on mini-batches and shape

- **Warning:** PyTorch / NN libraries typically multiply matrices in the **opposite direction** as Linear Algebra notation
- e.g. the input to an MLP should be (batch\_size, embedding\_size) rather than (embedding\_size, batch\_size)

# Note on mini-batches and shape

- **Warning:** PyTorch / NN libraries typically multiply matrices in the **opposite direction** as Linear Algebra notation
- e.g. the input to an MLP should be (batch\_size, embedding\_size) rather than (embedding\_size, batch\_size)
- The **last dimension of the input** should match the **first dimension of the weights**

# Note on mini-batches and shape

- **Warning:** PyTorch / NN libraries typically multiply matrices in the **opposite direction** as Linear Algebra notation
- e.g. the input to an MLP should be (batch\_size, embedding\_size) rather than (embedding\_size, batch\_size)
- The **last dimension of the input** should match the **first dimension of the weights**
- You can think of it as these libraries preferring  $x^T W^T$  to  $Wx$

# Note on mini-batches and shape

- **Warning:** PyTorch / NN libraries typically multiply matrices in the **opposite direction** as Linear Algebra notation
- e.g. the input to an MLP should be (batch\_size, embedding\_size) rather than (embedding\_size, batch\_size)
- The **last dimension of the input** should match the **first dimension of the weights**
- You can think of it as these libraries preferring  $x^T W^T$  to  $Wx$ 
  - (The result of this multiplication is the same, just transposed)

# Next Time

- Feedforward models for language
  - “Deep Averaging Network”
  - Feedforward language model
- Training regularization
- Language model quality metrics