

Neural Networks for Speech

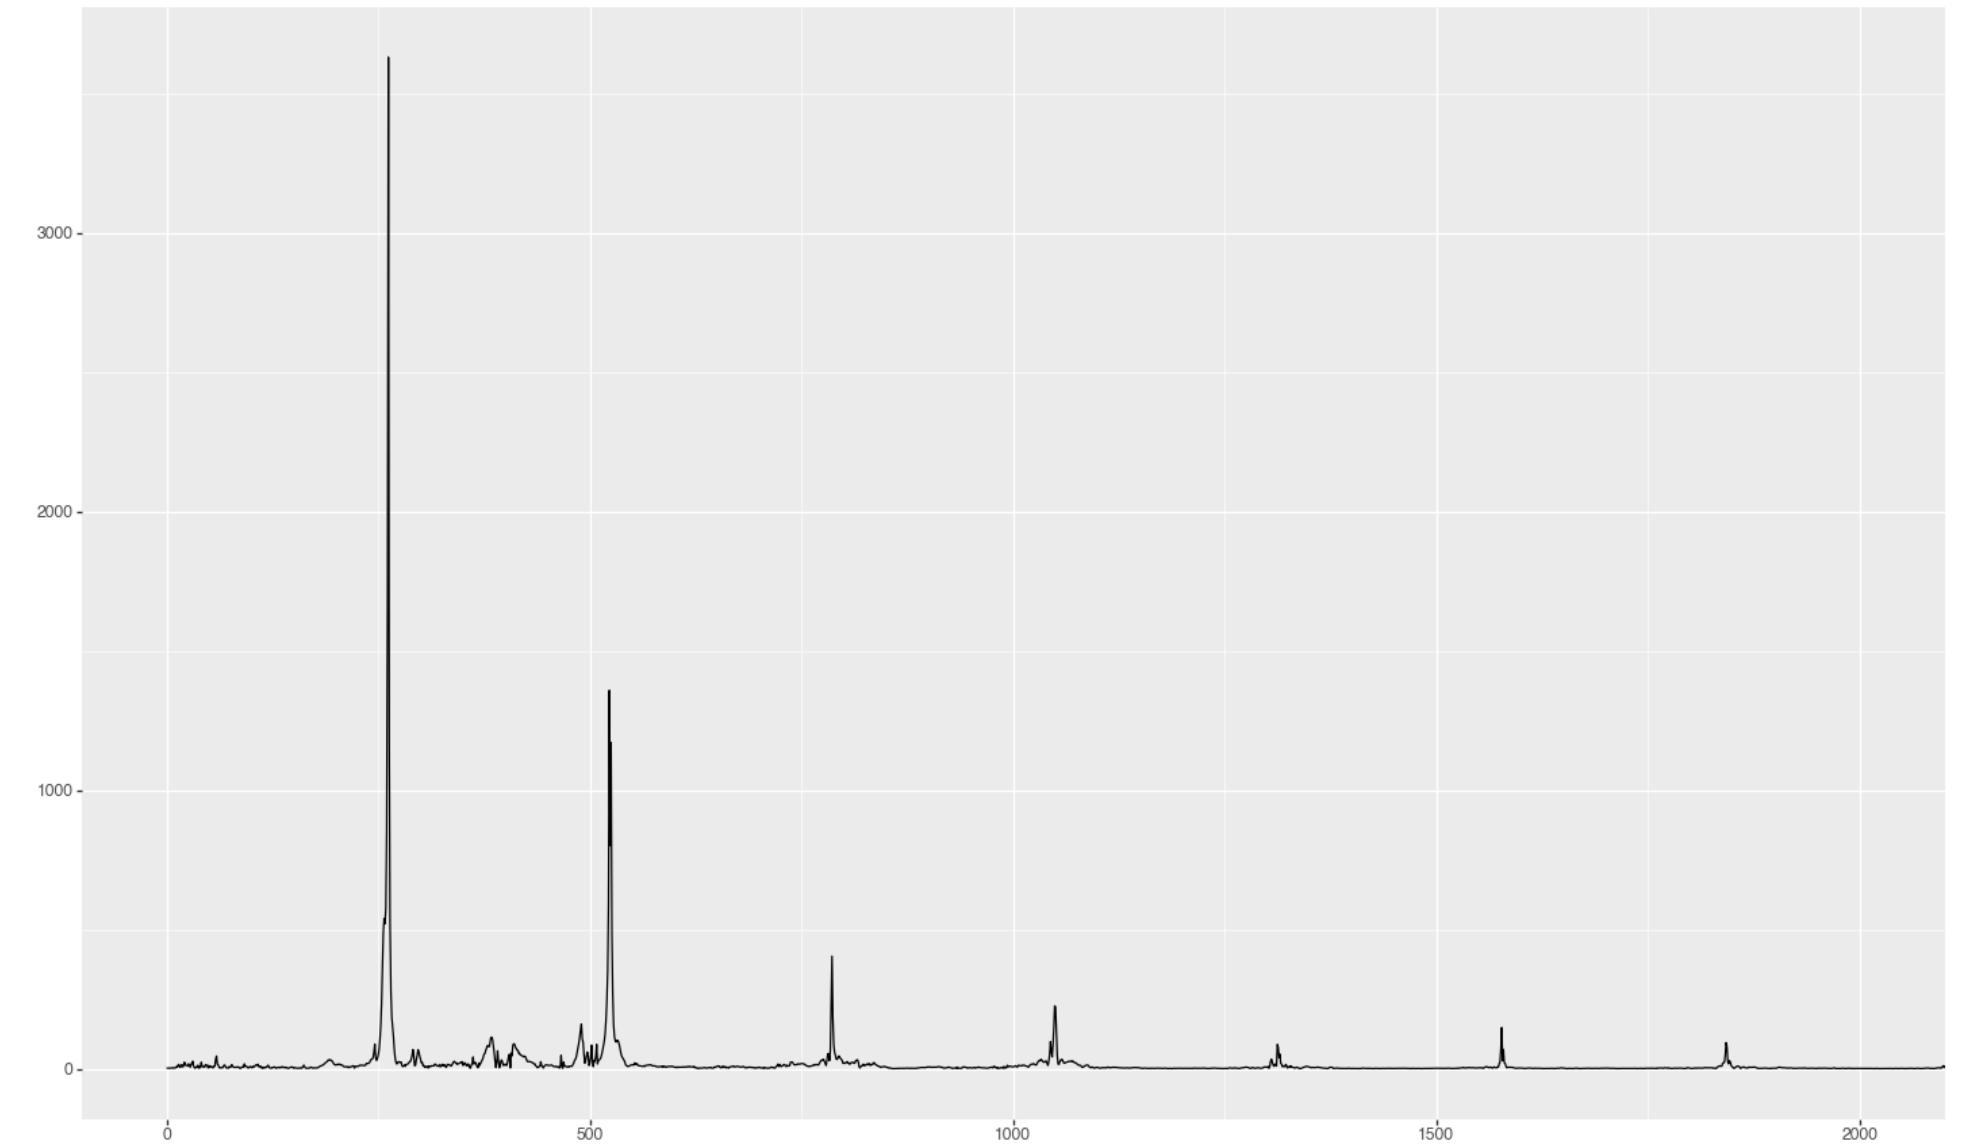
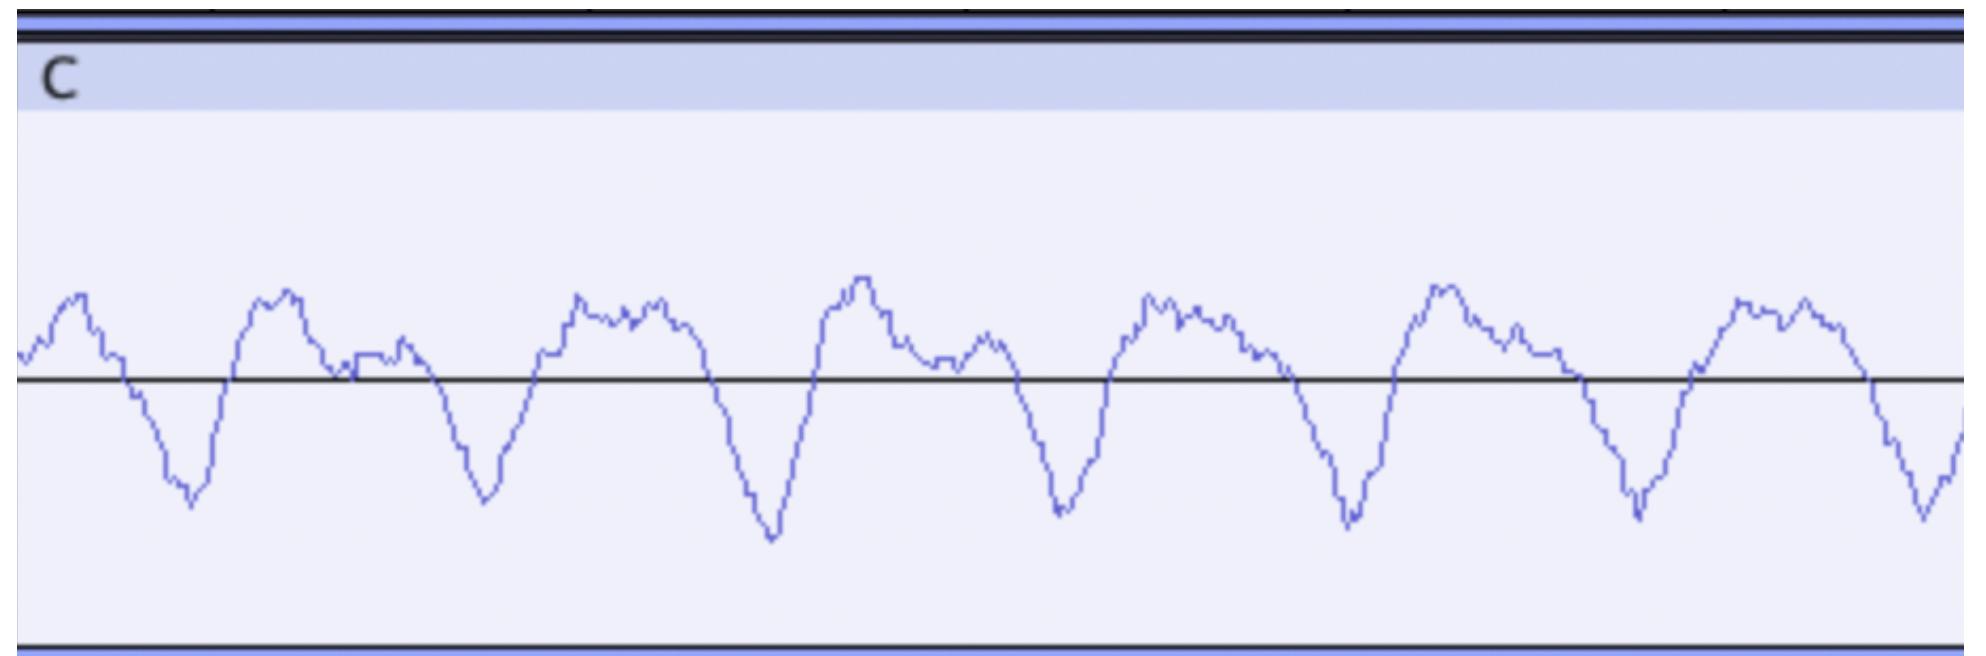
Ling 282/482: Deep Learning for Computational Linguistics

C.M. Downey

Fall 2025

Last time

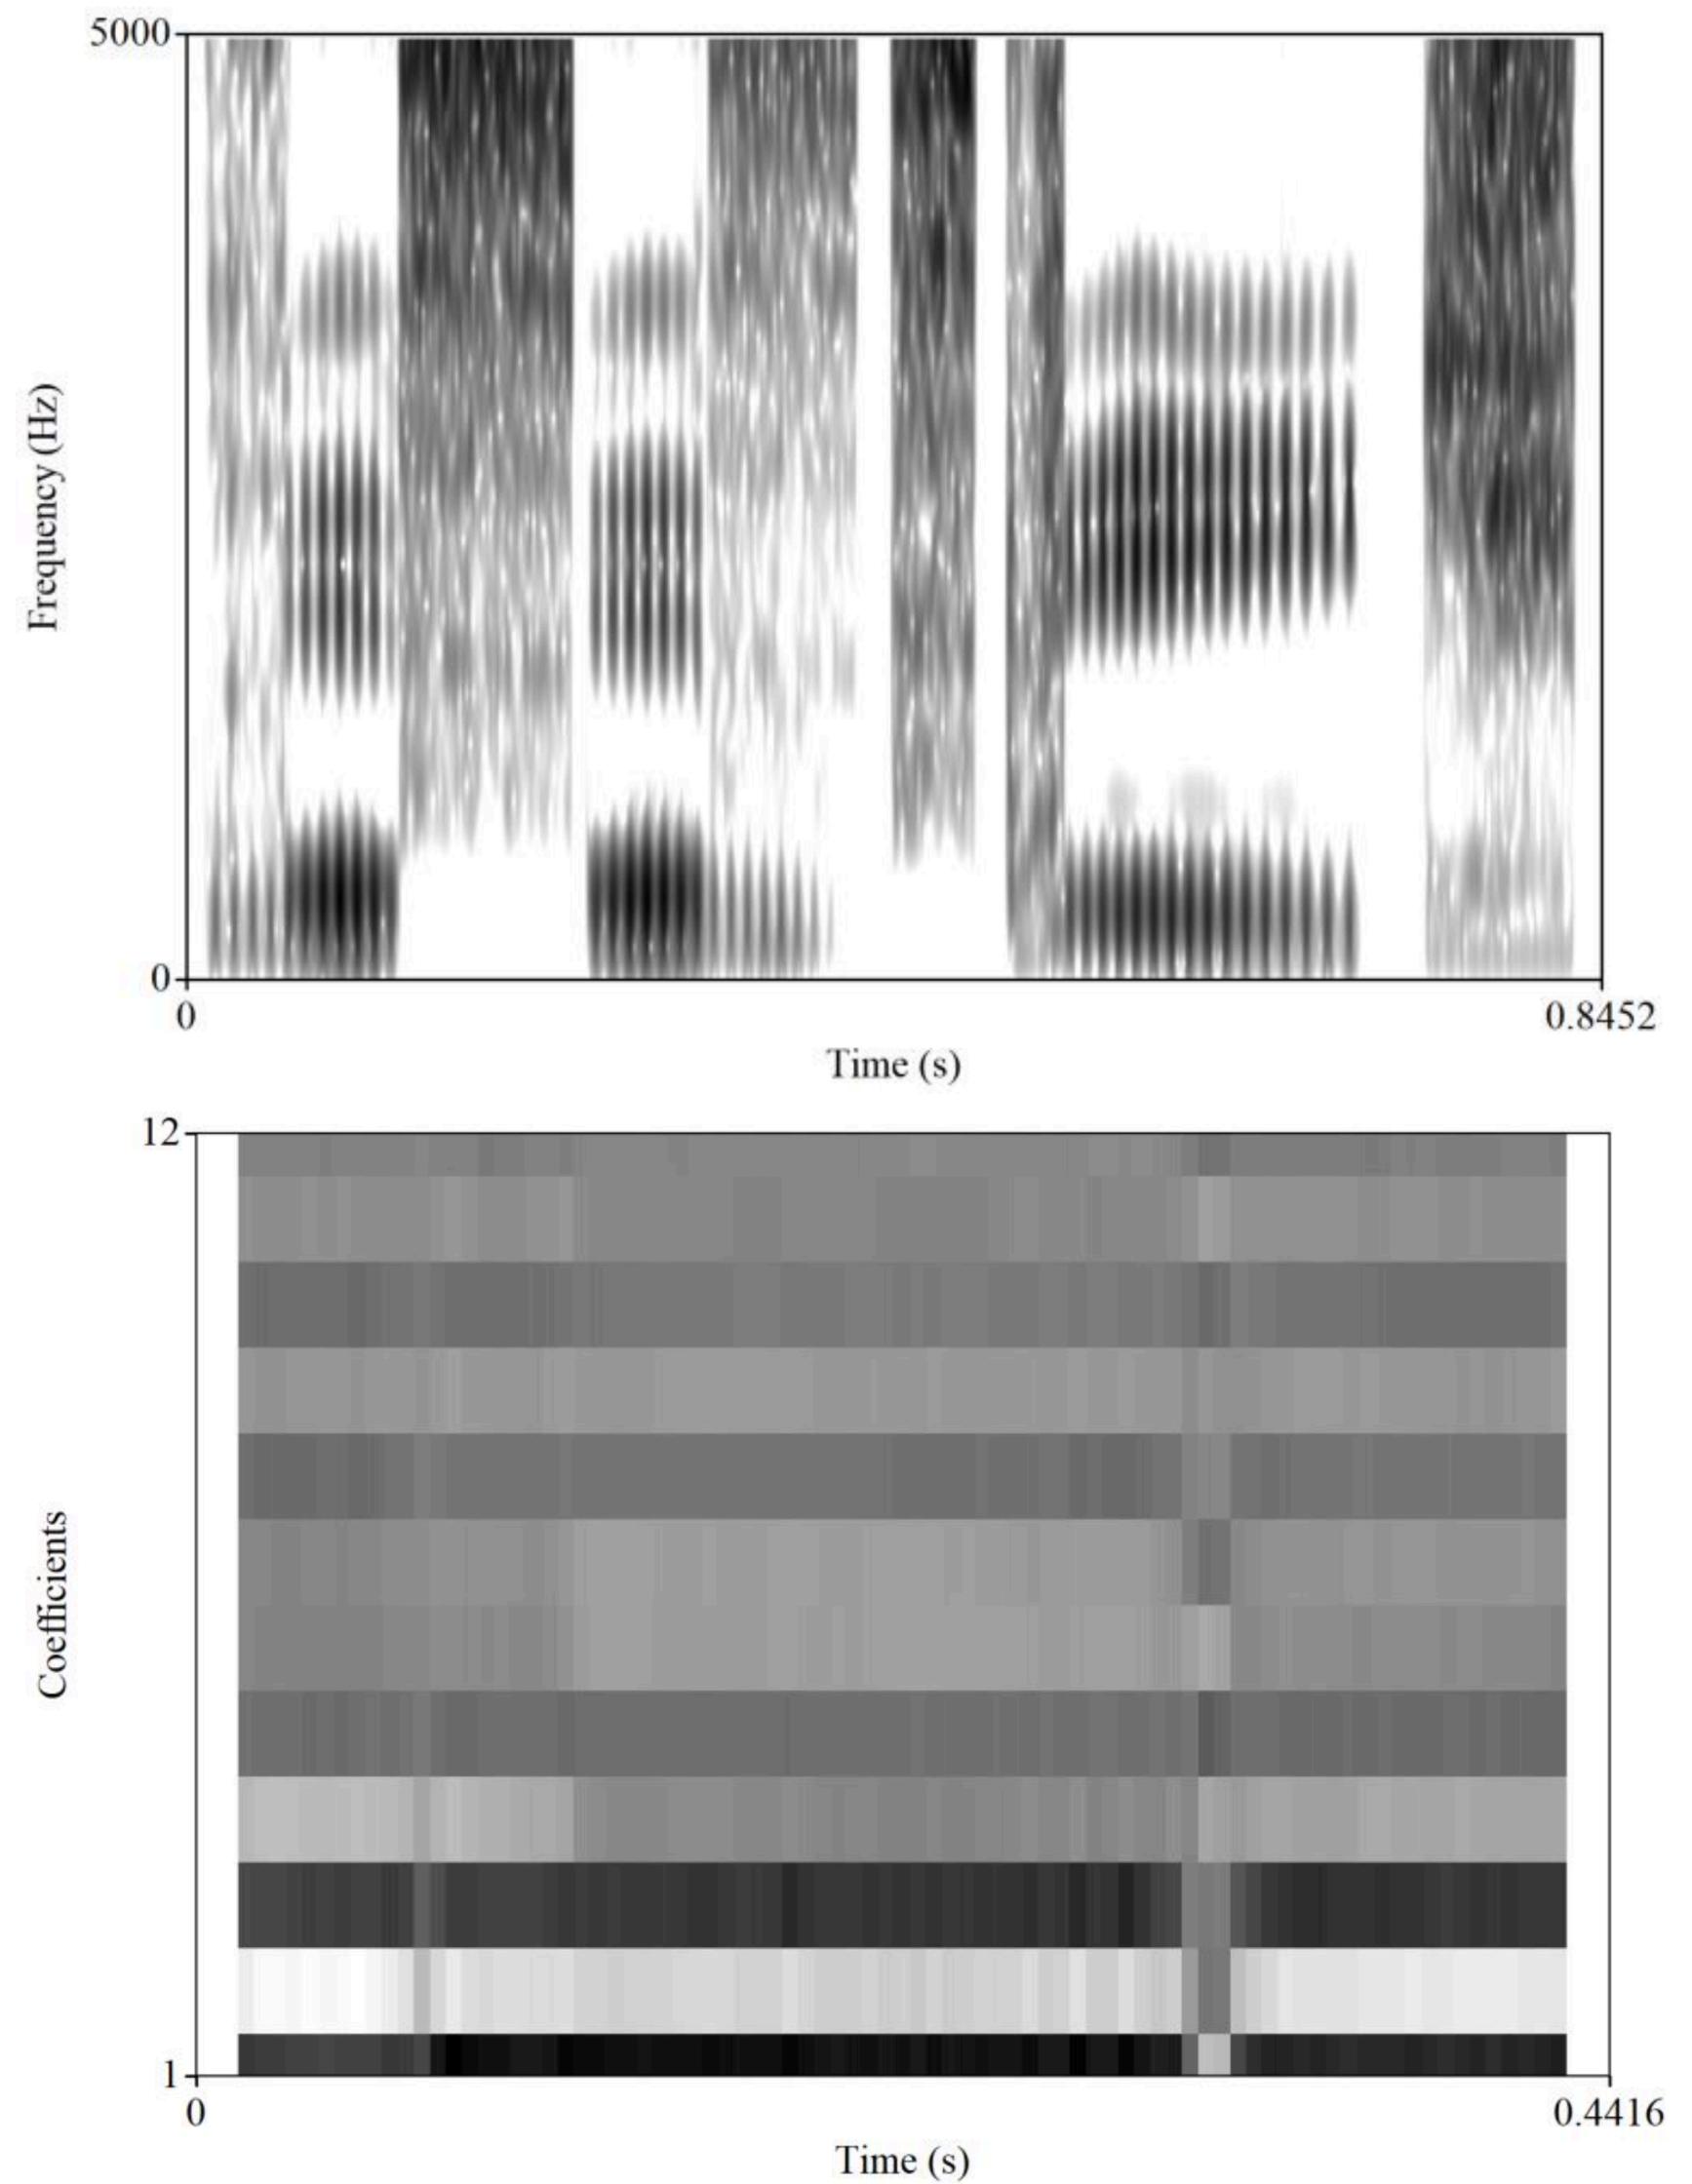
- Introduction to **acoustic data**
- Raw sound data known as the **waveform**
 - Just **amplitude across time** ("time domain" data)
- Apply the **Fourier Transform** to get the **spectrum** of component frequencies instead ("frequency domain")



Additional signal transformations

MFCCs

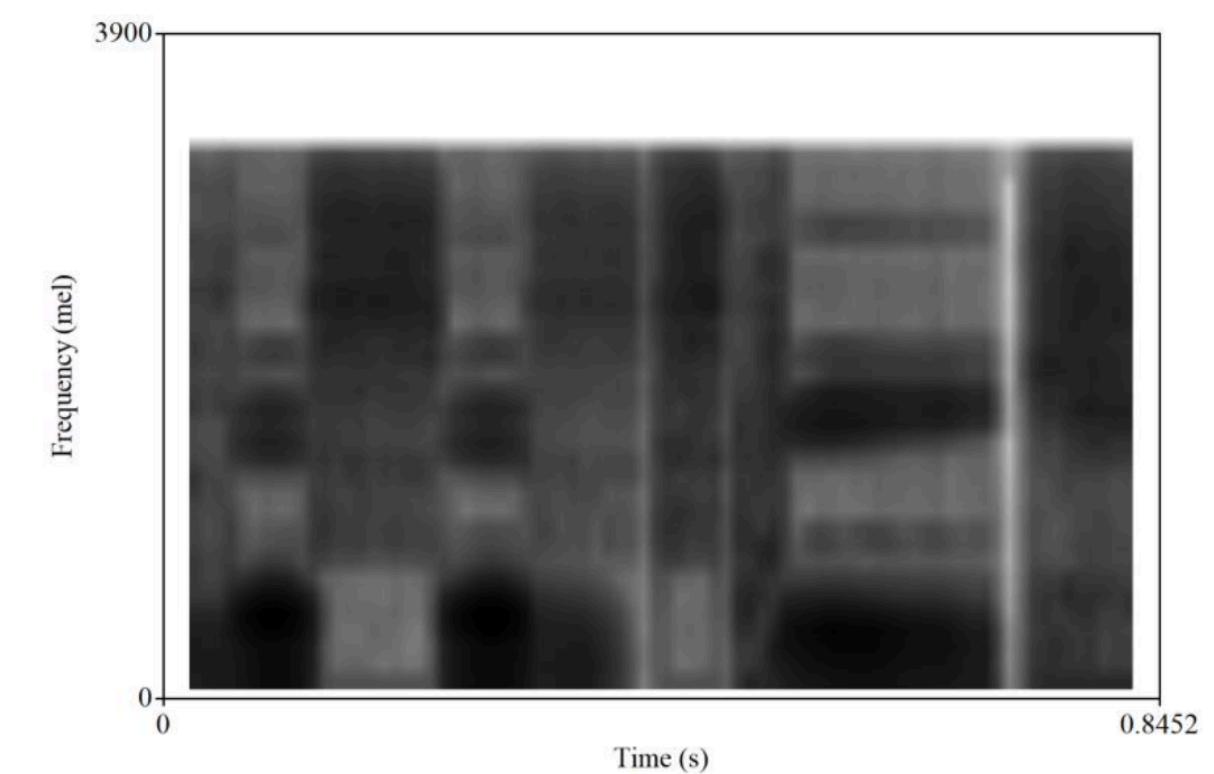
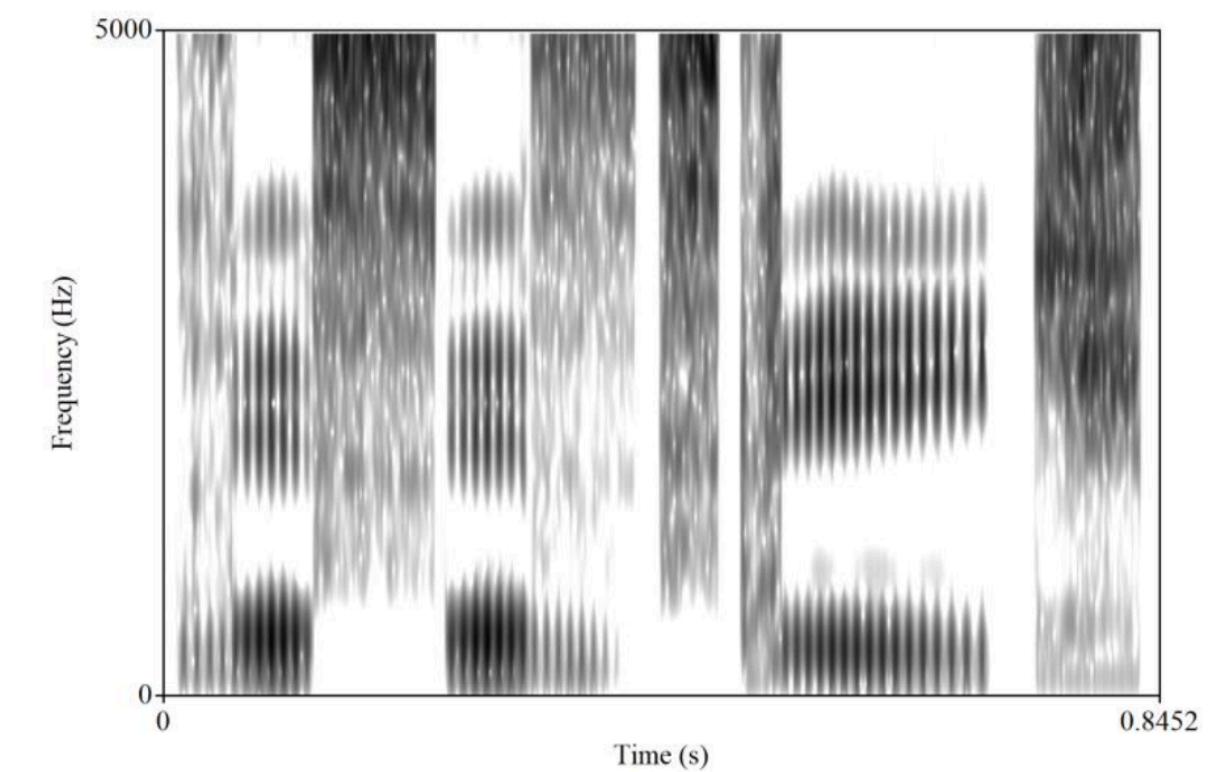
- The full FT spectrum is often re-sampled into **Mel Frequency Cepstral Coefficients (MFCCs)**
 - This is fairly advanced signal processing
 - The main point is to **reduce the dimensionality** of the spectrum
 - Can be thought of as a **compression** of the full spectrum
- Traditionally used as **input to machine learning models**



Log Mel Spectrogram

- Spectrogram may also be re-sampled according to the **(Log) Mel Scale**, but not reduced to MFCCs
 - The result called a **Log Mel Spectrogram**
 - Has **more information** than MFCCs
 - Popular for **Neural Networks**, which don't need so much feature extraction in the input

Standard spectrogram



Mel spectrogram

Neural Modeling Overview

Speech Encoder-Decoder

- Standard task for speech processing is called **Automatic Speech Recognition (ASR)**
 - Essentially: **speech-to-text**
- Neural approaches have typically used an **Encoder-Decoder** model
 - Encoder learns how to **usefully represent speech signal**
 - Decoder outputs the **corresponding text sequence**

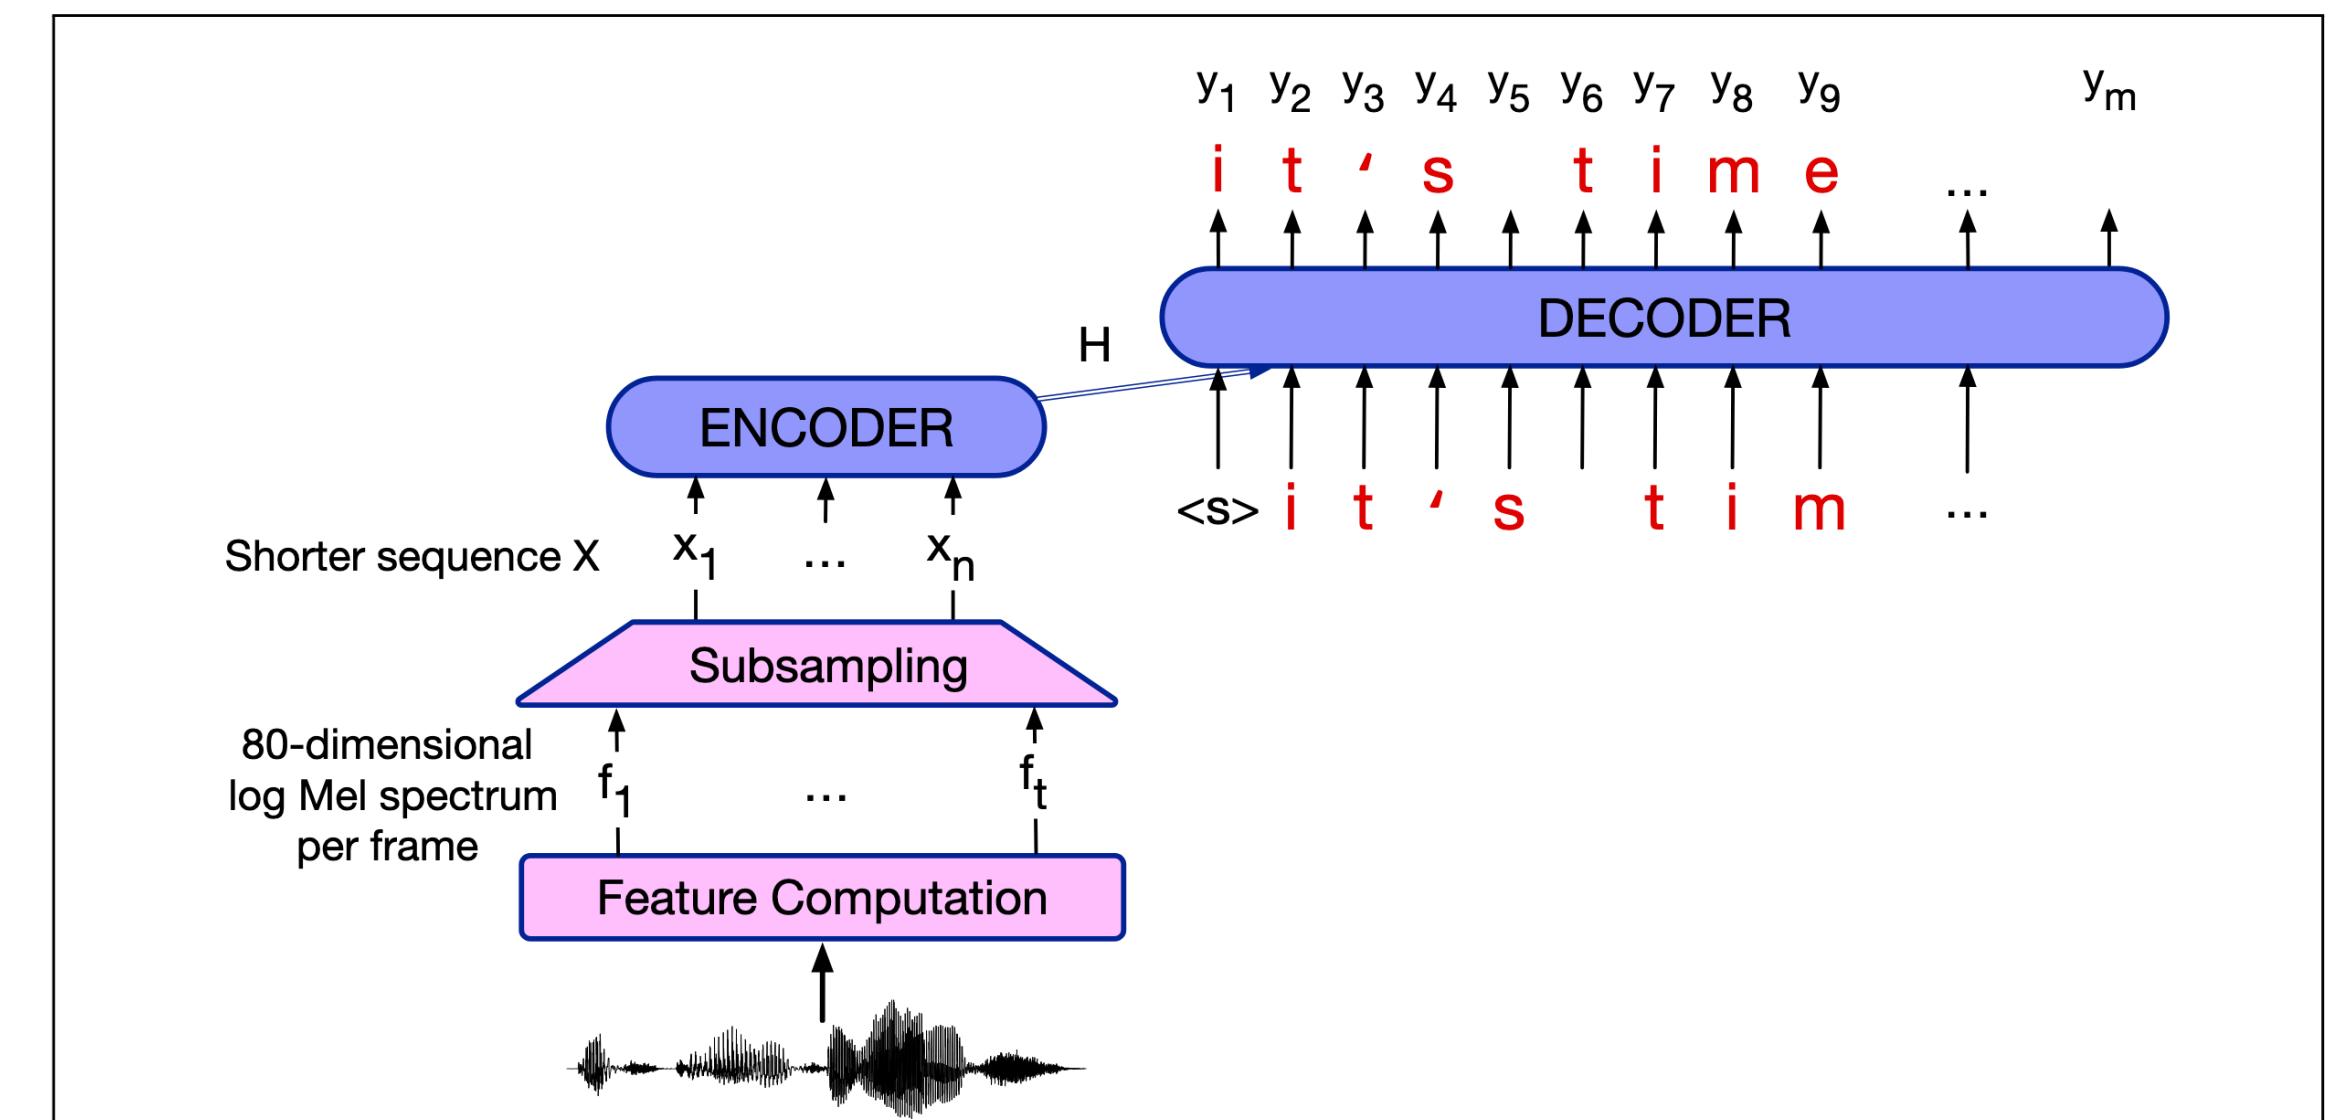


Figure 15.5 Schematic architecture for an encoder-decoder speech recognizer.

Speech Encoder-Decoder

- The main bodies of these models are often **familiar architectures** like the Transformer
- However, each has **additional components** to handle the more complex ASR task
- Almost all encoders use **Convolutional Neural Networks (CNNs)**

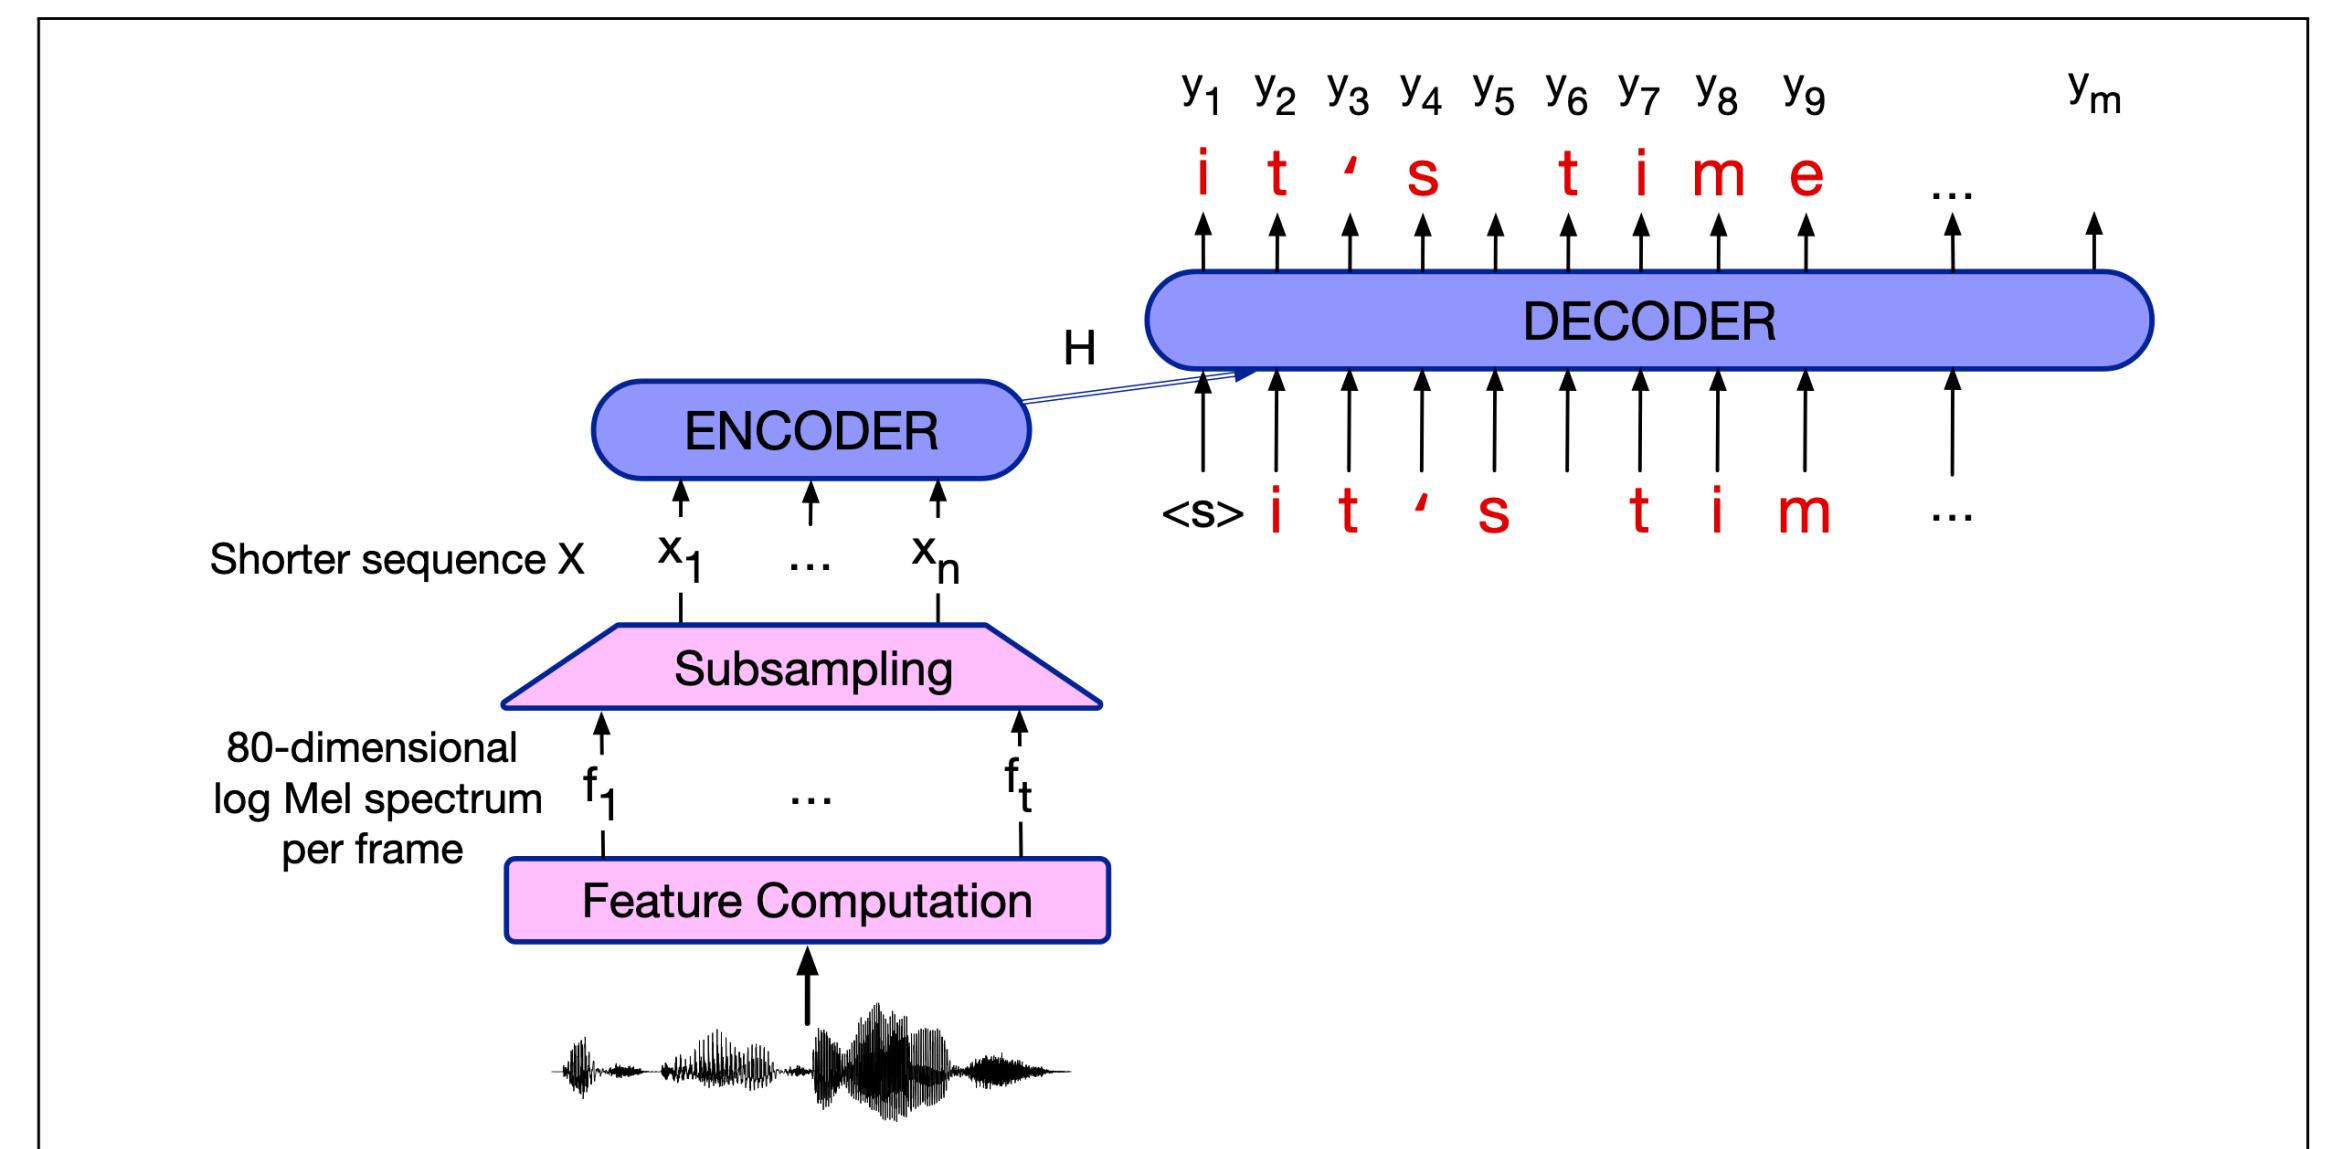


Figure 15.5 Schematic architecture for an encoder-decoder speech recognizer.

Convolutional Neural Networks

- Key intuition: a "sliding" filter called a **kernel**
 - Extracts information from **contiguous regions** of the input
 - In audio, filter is over a **span of time**
 - In images, an **area of the image**
- Essentially, take the **dot product** of the kernel with the input window
- Name comes from math concept of Convolution (which is slightly different)

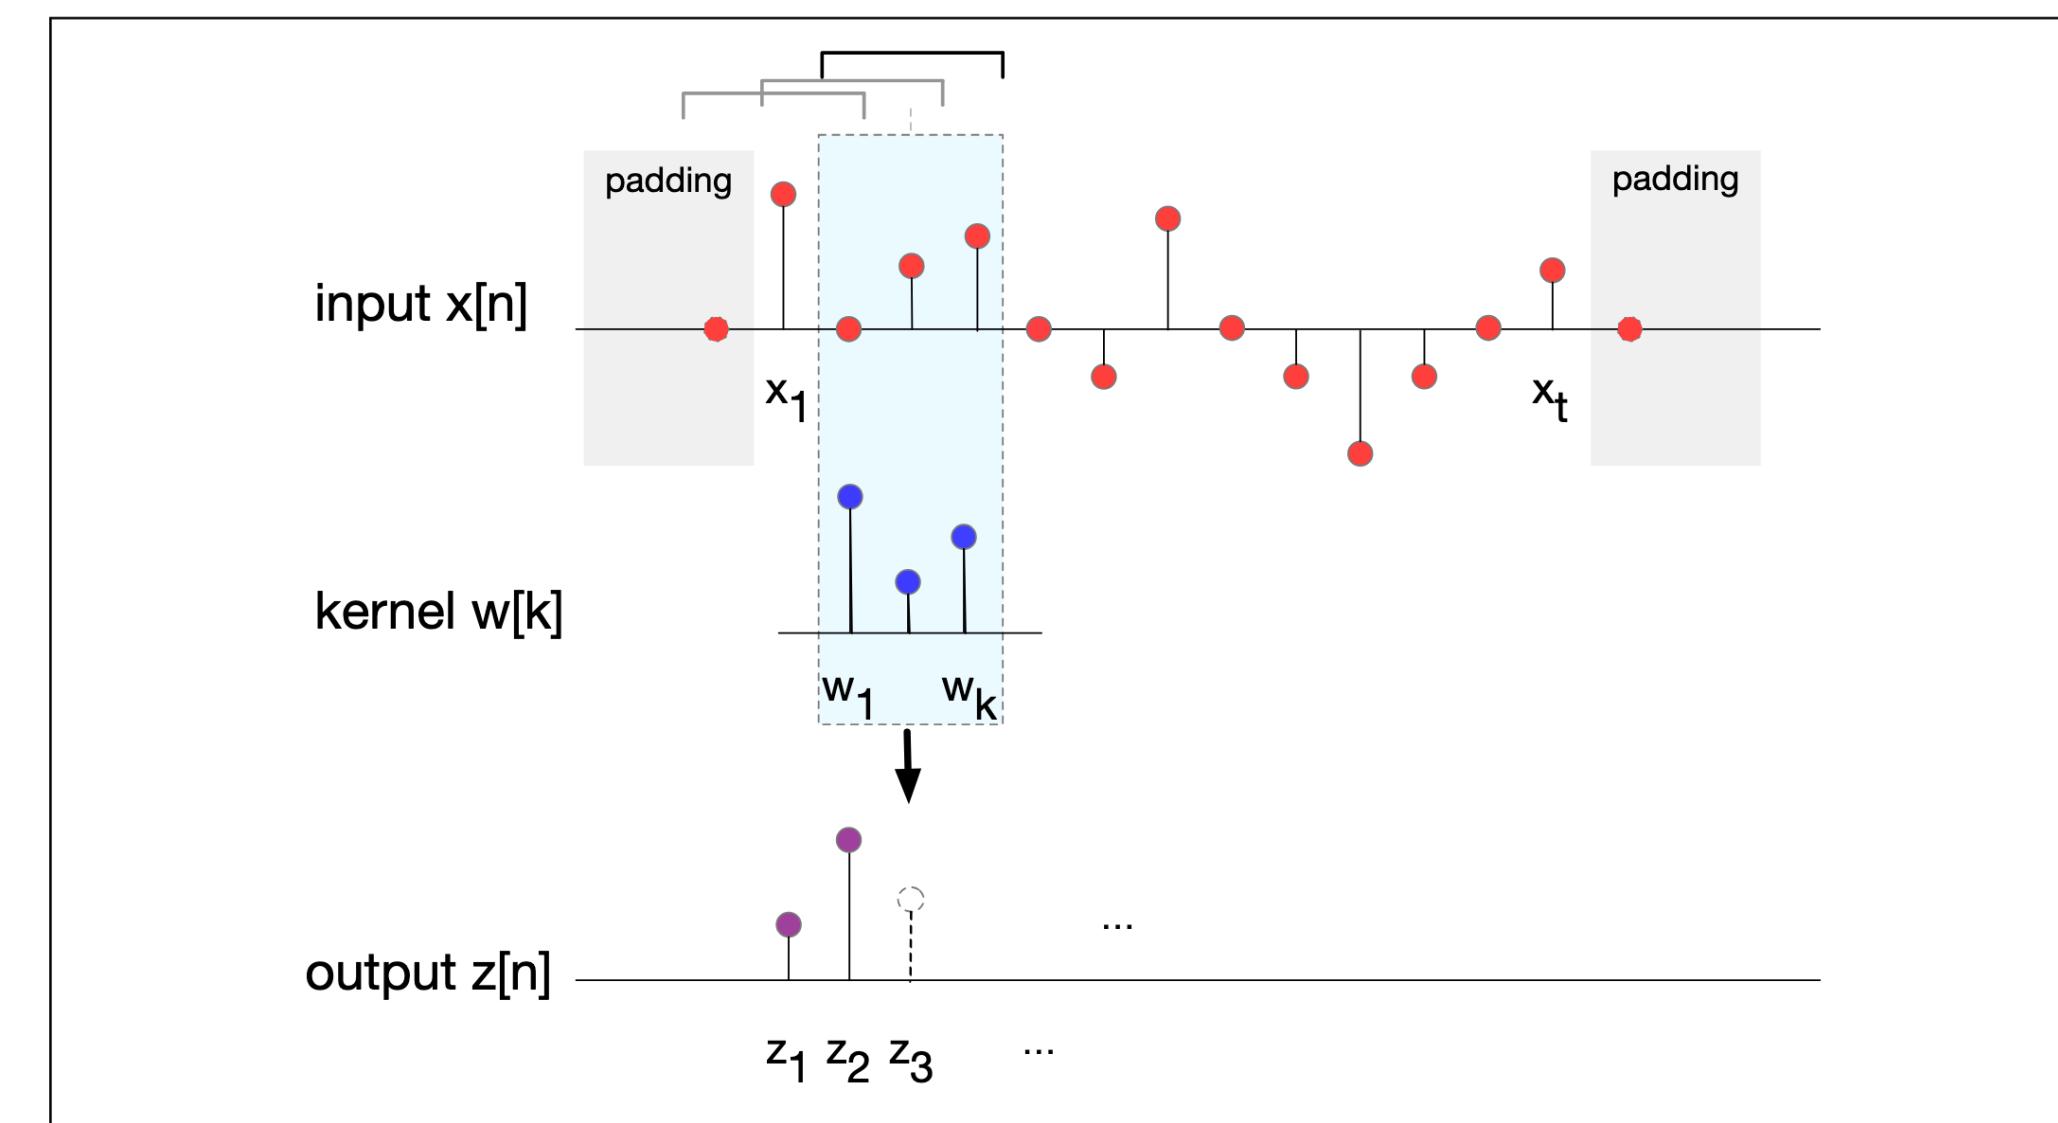


Figure 15.3 A schematic view of convolution with a kernel (filter) width of 3, and with a padding of 1, showing a zero value added at the start and end of the signal. The (already flipped) kernel is walked across the input, and the output at each frame z_i is the dot product of the kernel with the input in the window. The figure shows the computation of z_3 .

Speech CNN Layer

- The kernel contains the **learned weights**
- Kernel is "**slid**" across time
 - Has an option called **stride**, which controls how far to move in each step
 - Higher stride → output sequence is **shorter** than input
- Each embedding dimension usually called a **channel**

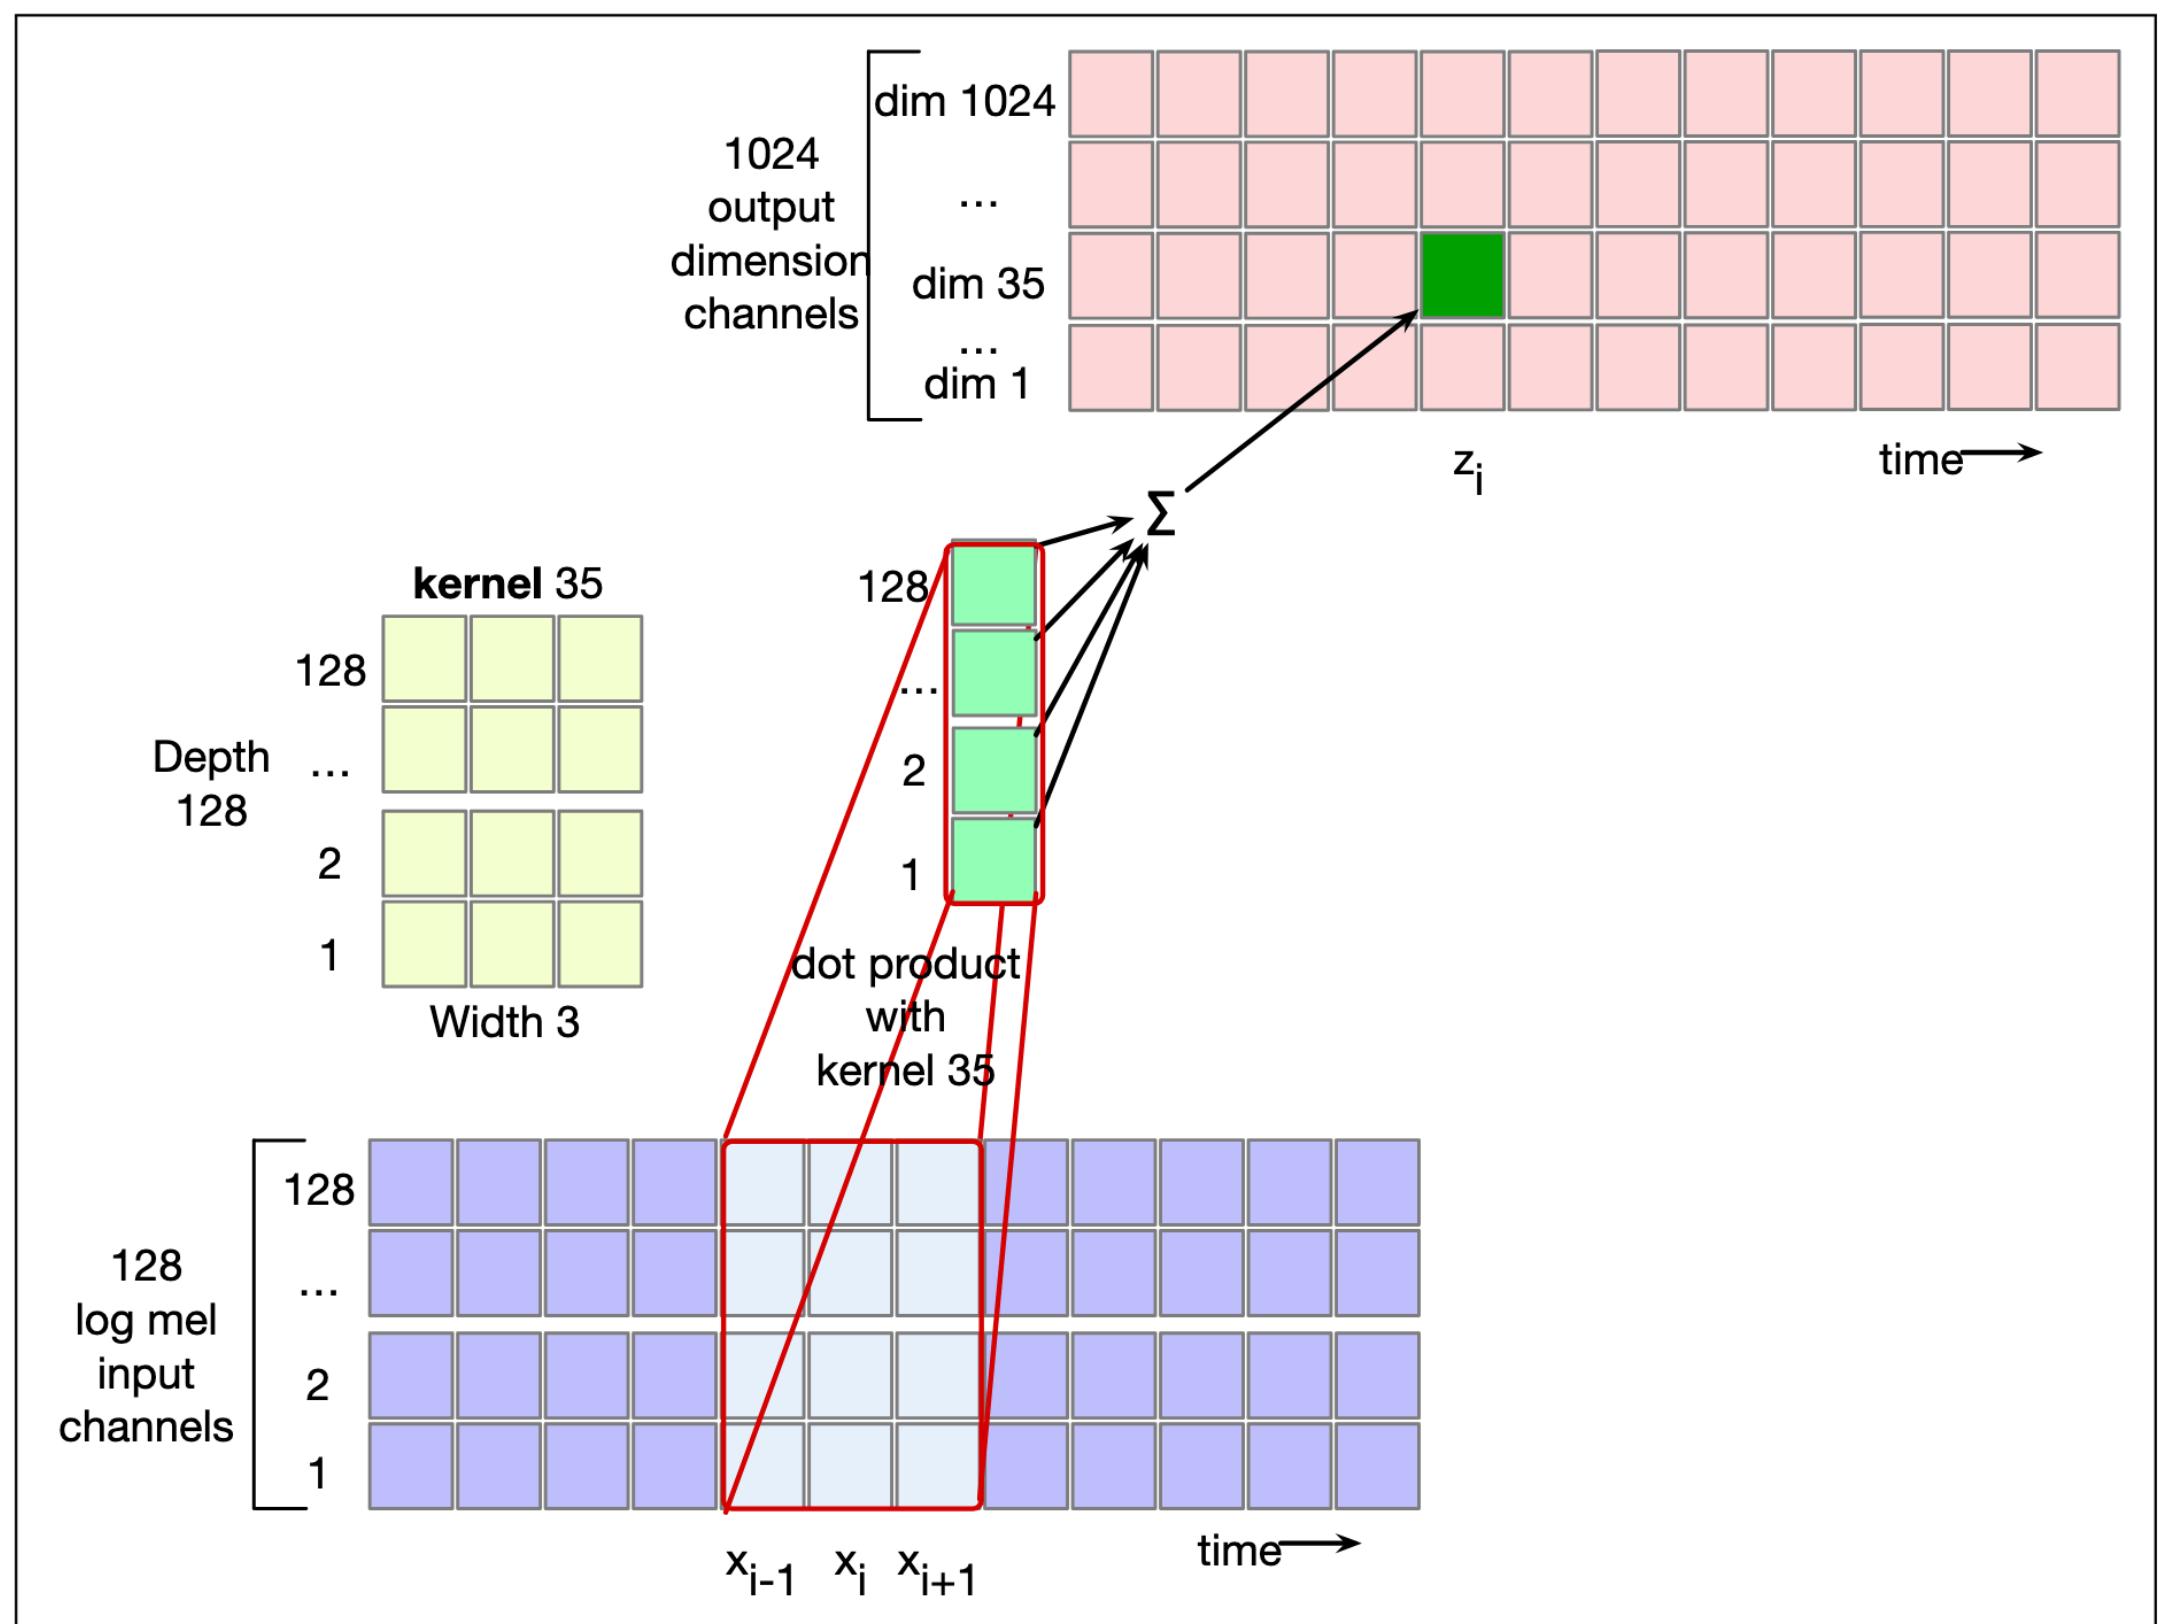


Figure 15.4 A schematic view of a convolutional net with 128 input channels and 1024 output channels. We see how at time point i one of the 1024 kernels ("kernel 35", each of depth 128 and width 3) is dot-product-ed with (each of) the 128 log mel spectrum input vectors, and then summed to produce a single value for one dimension of the output embedding at time i .

Speech CNN Layer

- For kernel length 3, 128 channels:
 - Kernel becomes size **128x3**
 - However, we still **treat this like a dot product**: sum is taken over whole 128x3 area
 - Another way to think about it: treat the input area and kernel as **single vectors**, and take the dot product
 - Gives the value for a **single output channel!**
 - **Each output channel gets a different kernel**

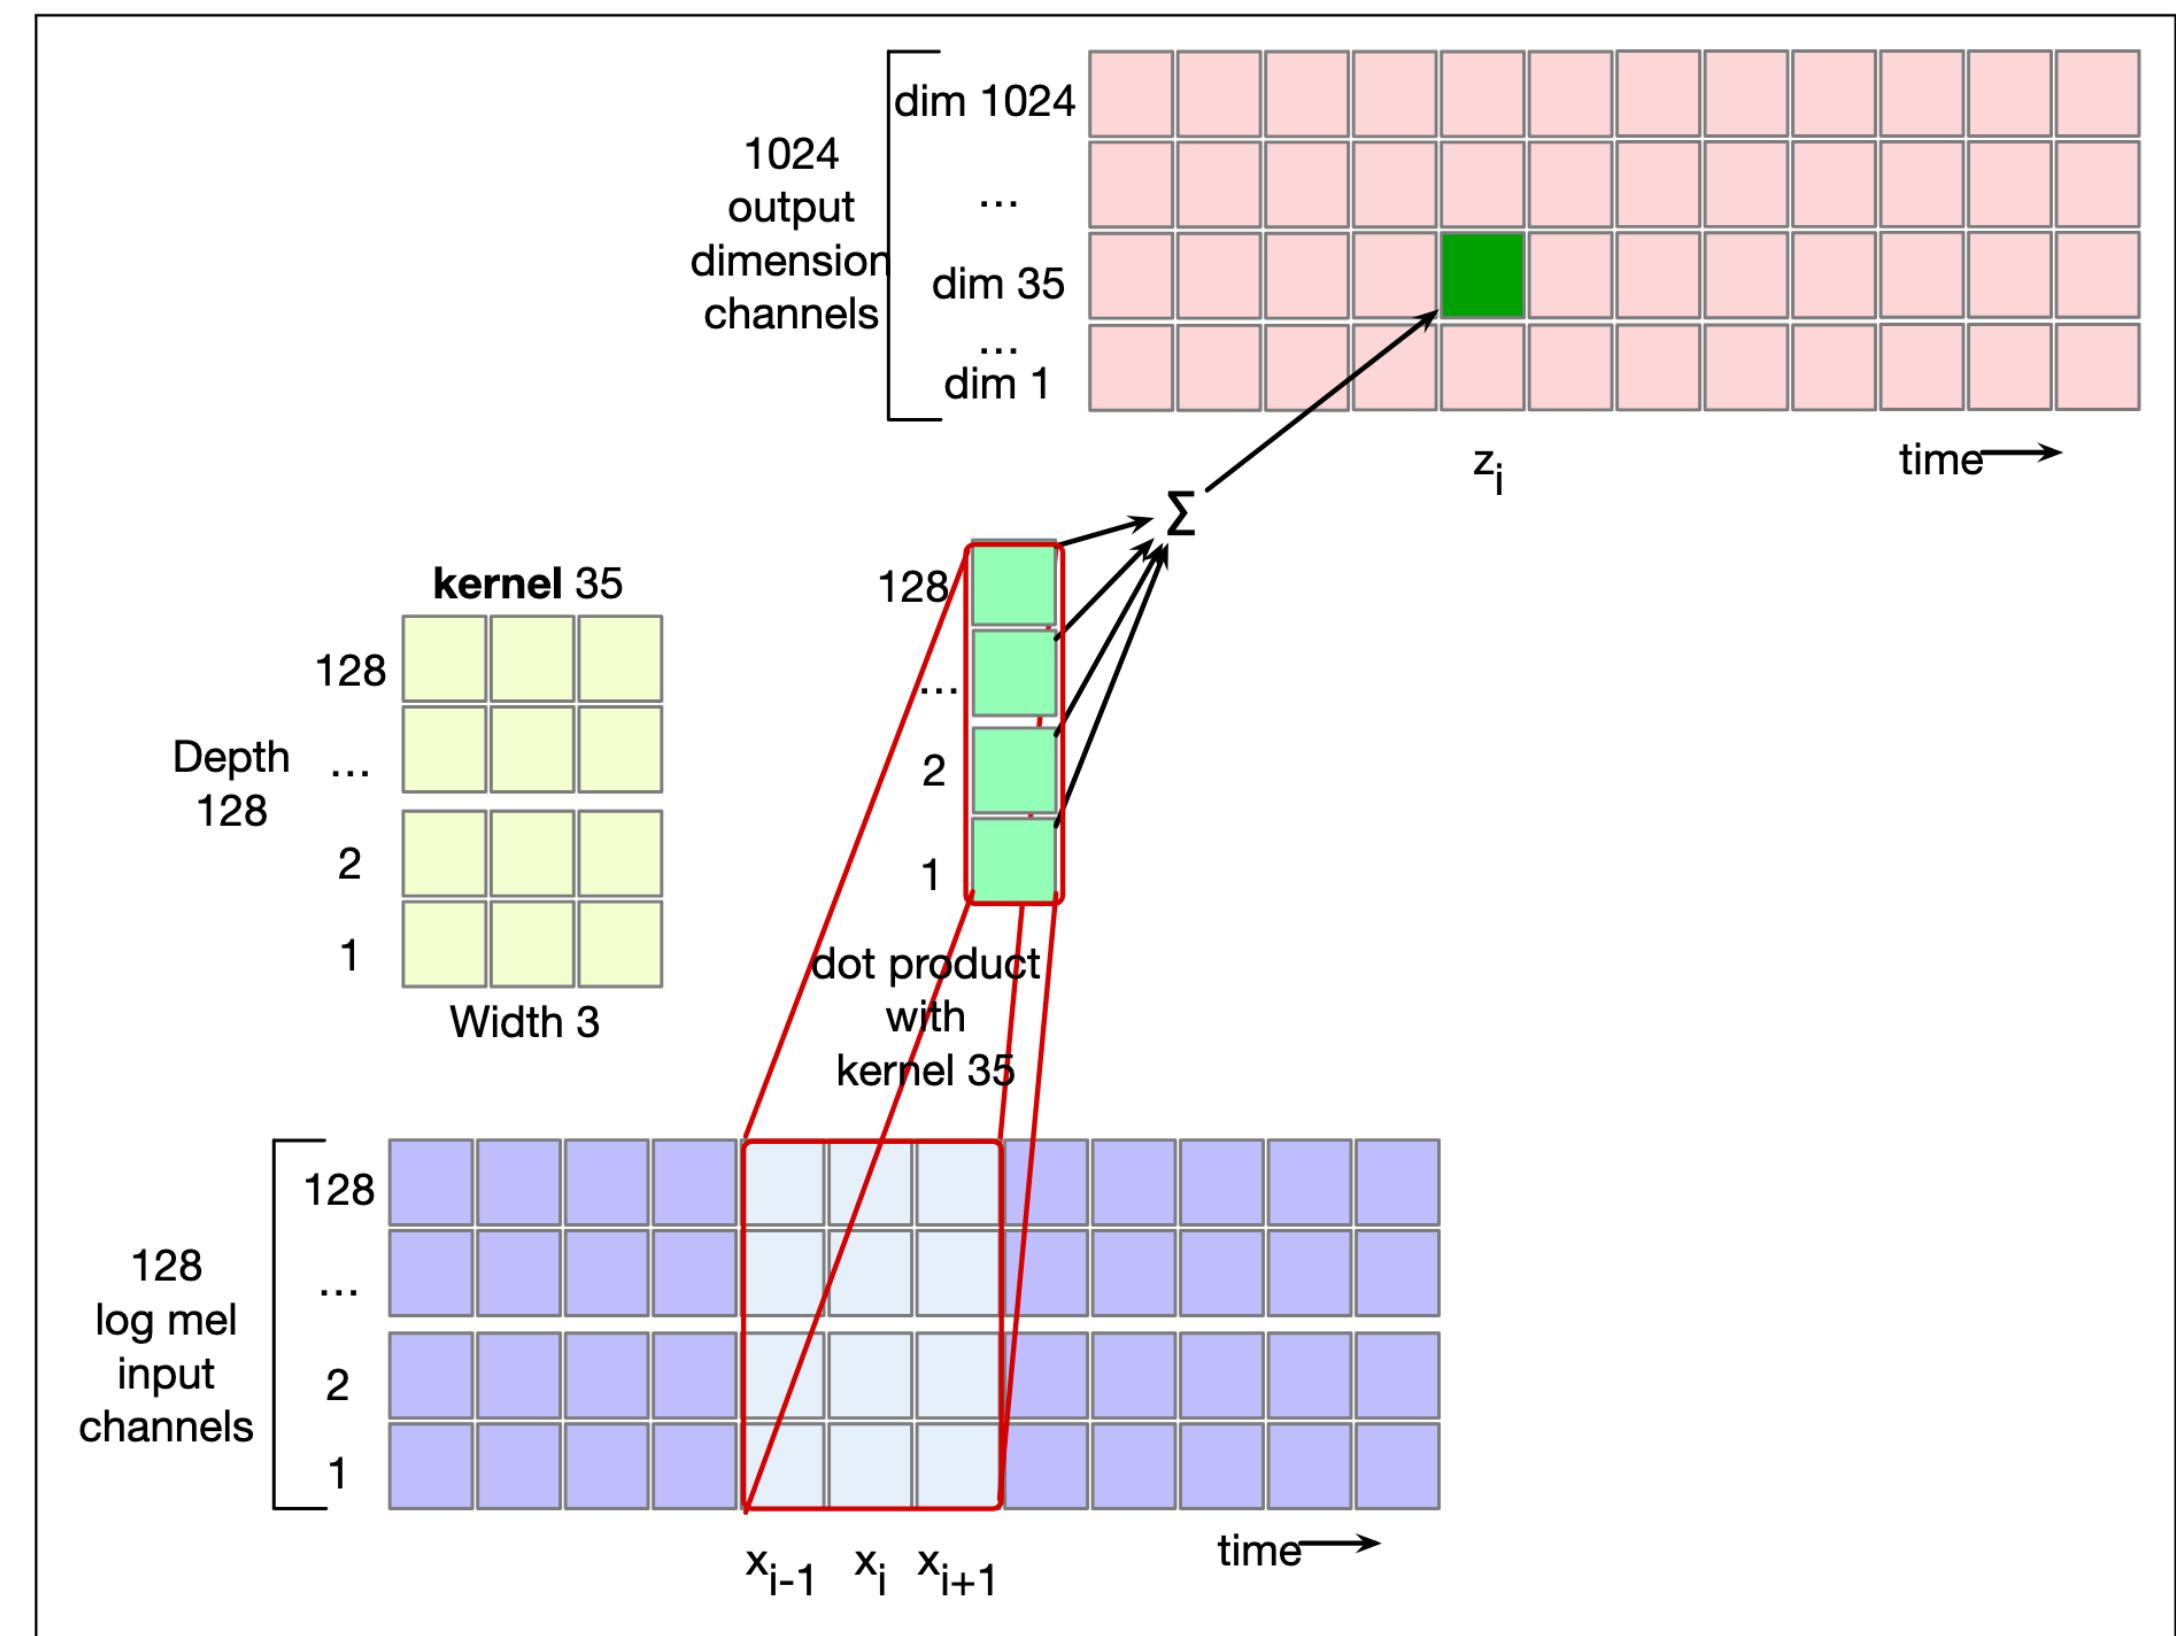


Figure 15.4 A schematic view of a convolutional net with 128 input channels and 1024 output channels. We see how at time point i one of the 1024 kernels (“kernel 35”, each of depth 128 and width 3) is dot-product-ed with (each of) the 128 log mel spectrum input vectors, and then summed to produce a single value for one dimension of the output embedding at time i .

CTC



UNIVERSITY OF ROCHESTER

Why we need CTC

- In audio, each input timestep corresponds to e.g. a **25ms window**
- Speech sounds might stretch across **multiple time-steps**
- Text transcription **doesn't tell us** how to align characters to time!
- How do we figure out this alignment?

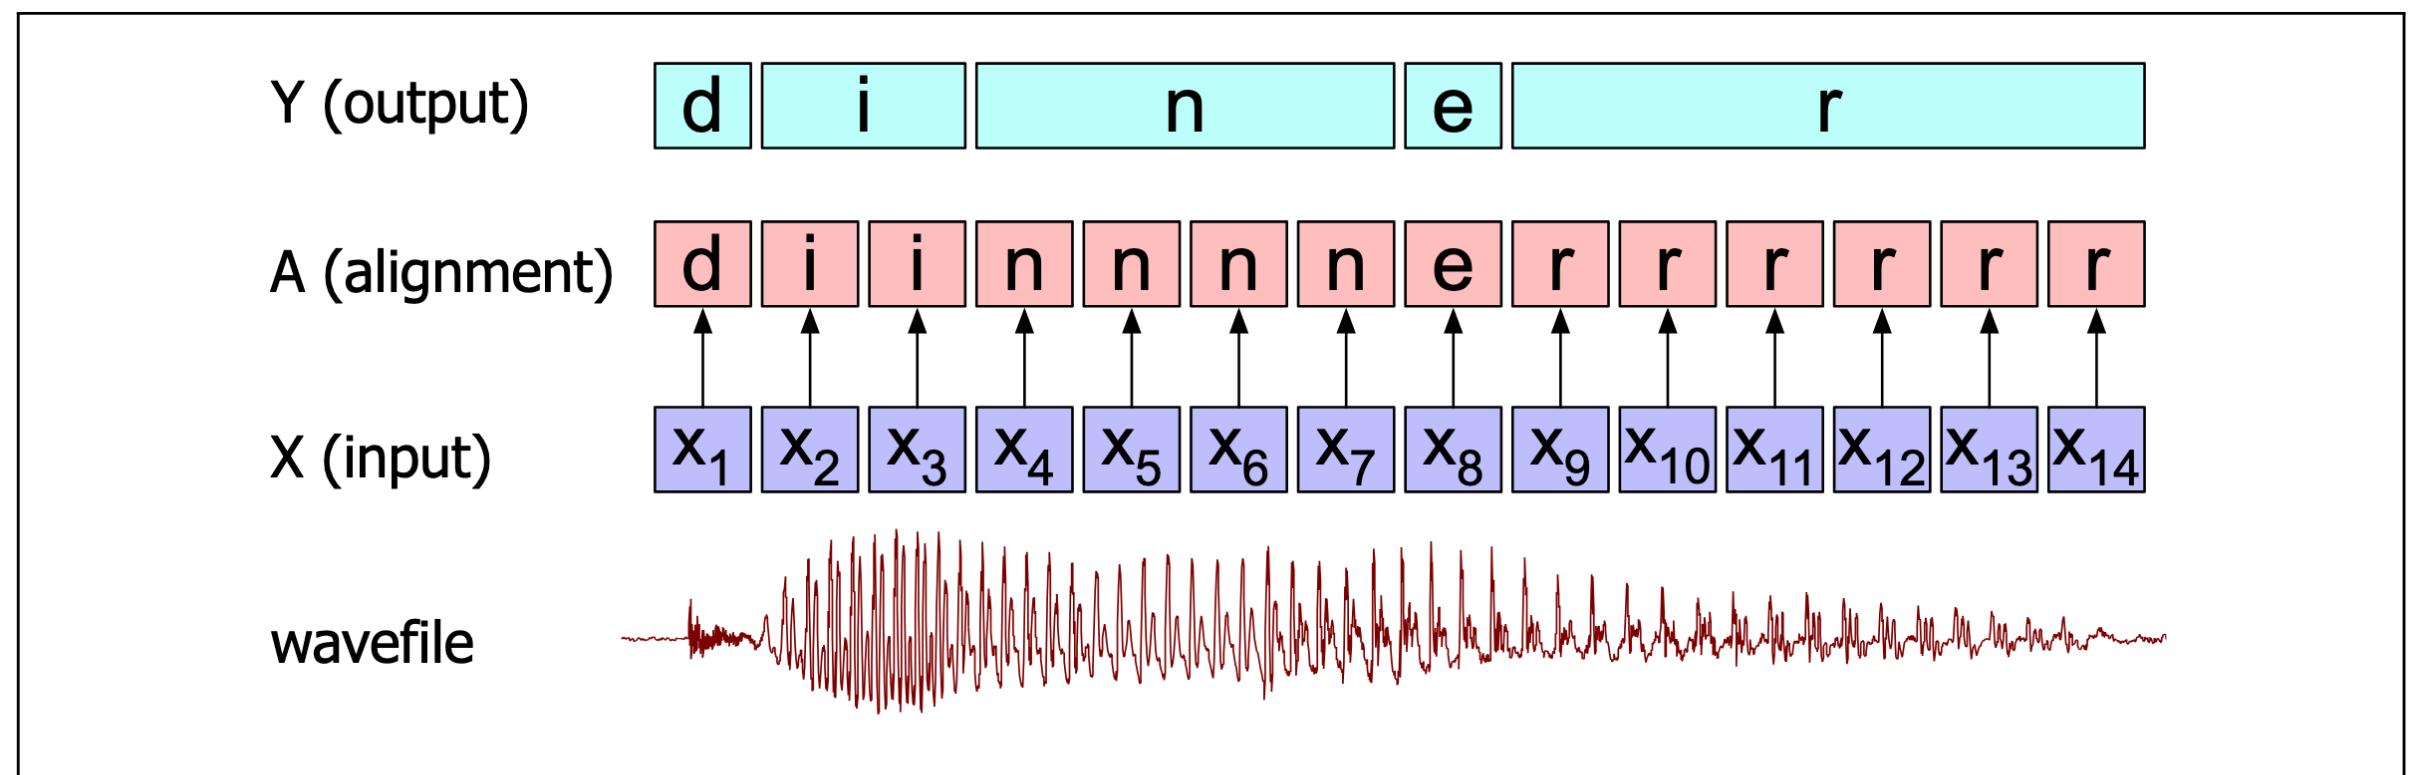


Figure 15.12 A naive algorithm for collapsing an alignment between input and letters.

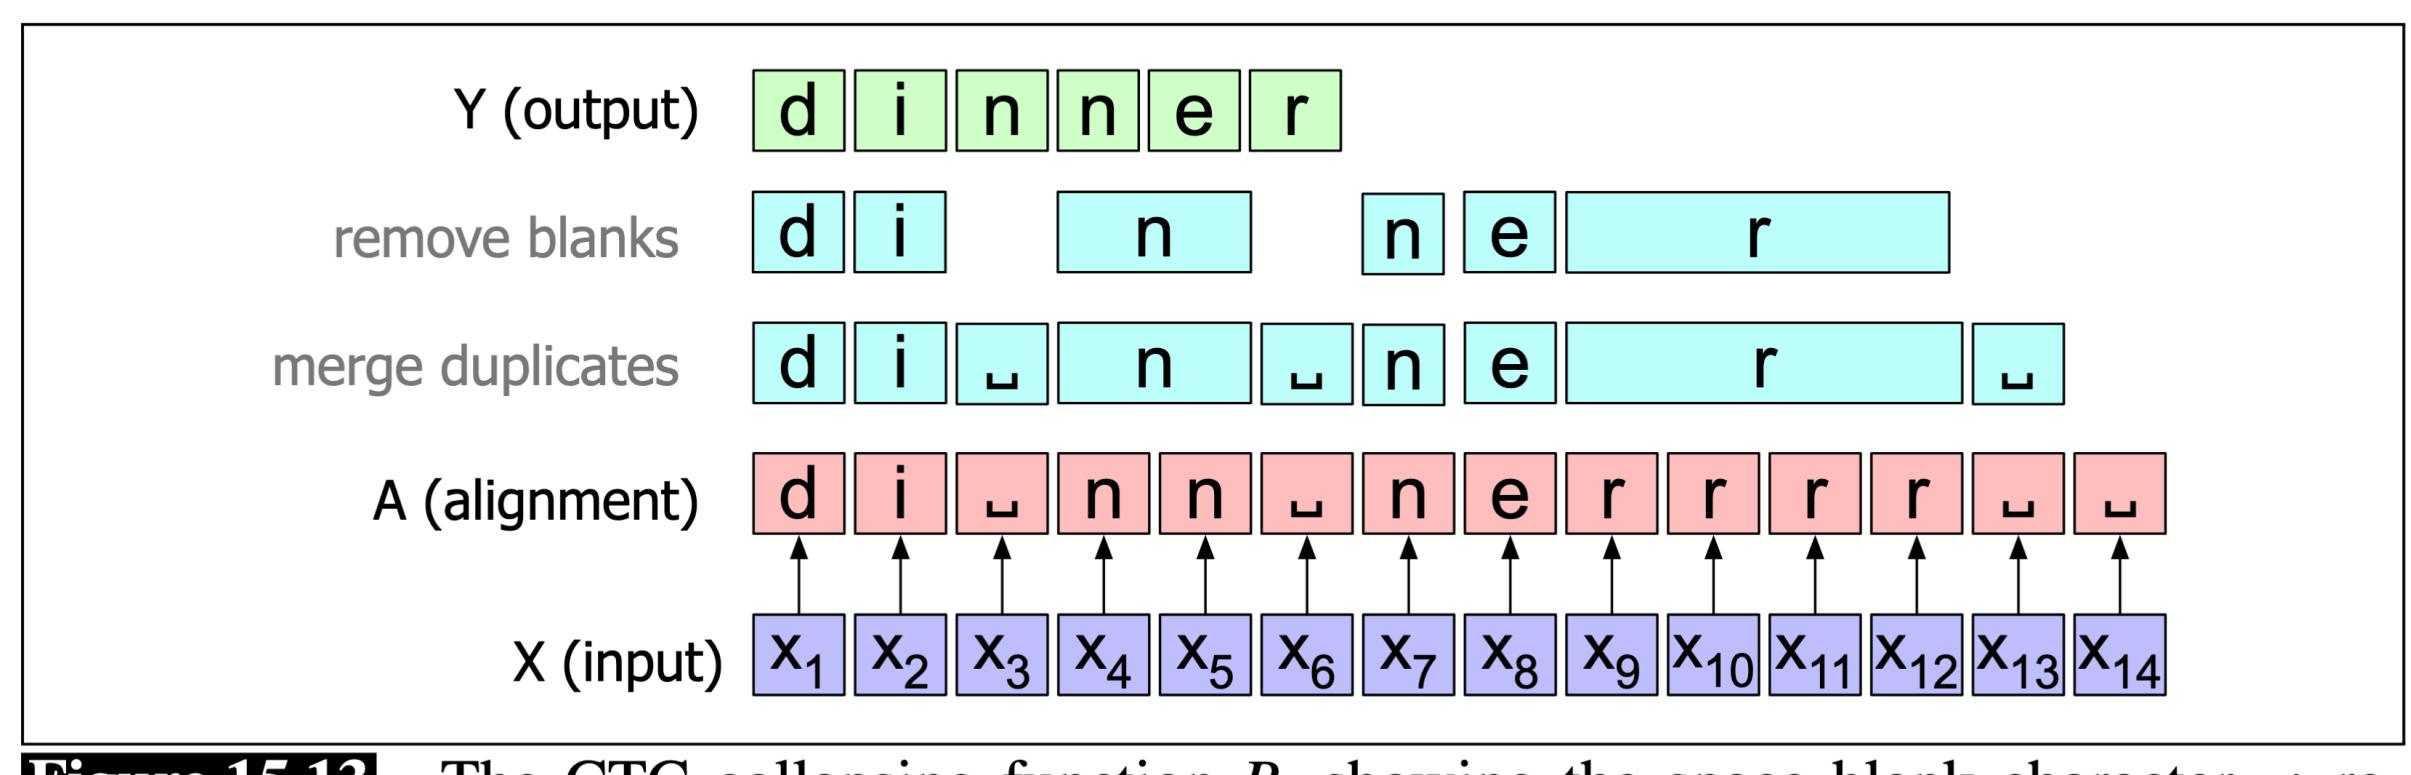


Figure 15.13 The CTC collapsing function B , showing the space blank character u ; repeated (consecutive) characters in an alignment A are removed to form the output Y .

CTC

- CTC: Connectionist Temporal Classification
 - Essentially, a way to **align symbols (characters)** to output time-steps
- Uses **blank symbol** for two purposes:
 - To represent **silence**
 - To **break up** duplicate characters
- Otherwise **merges duplicates**

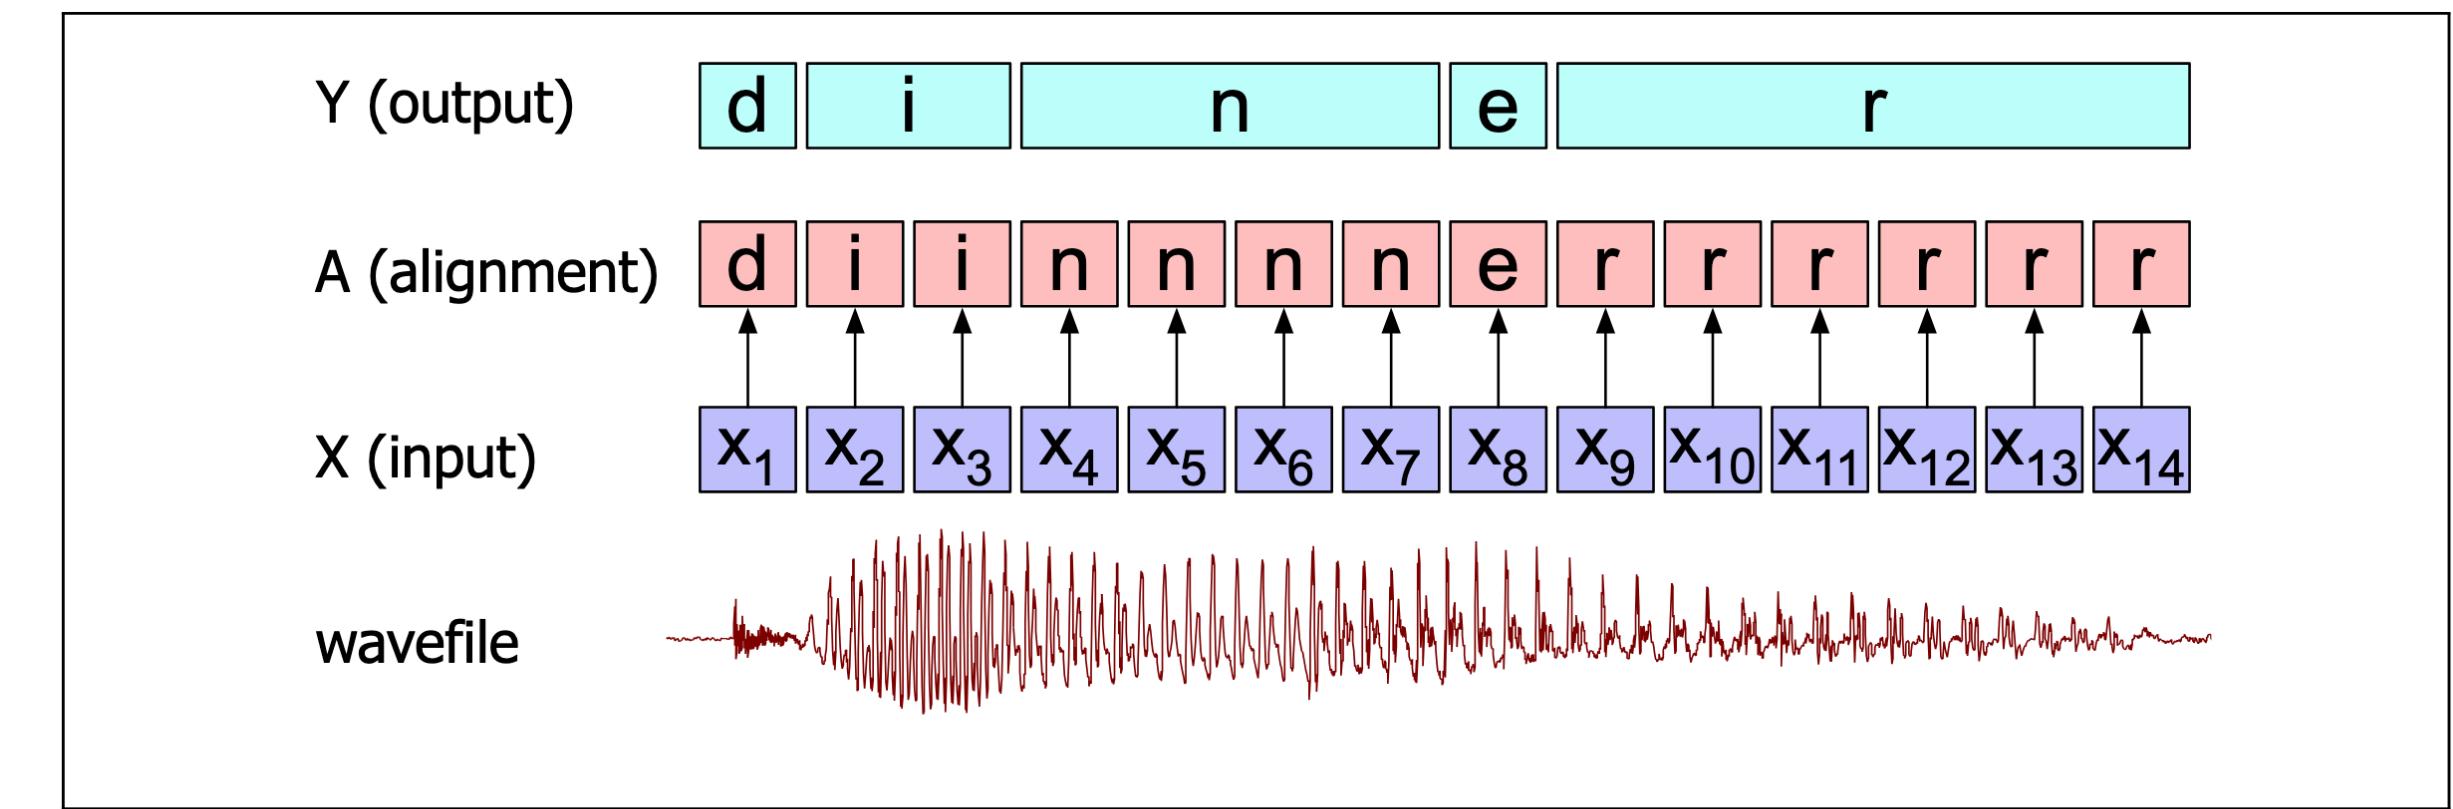


Figure 15.12 A naive algorithm for collapsing an alignment between input and letters.

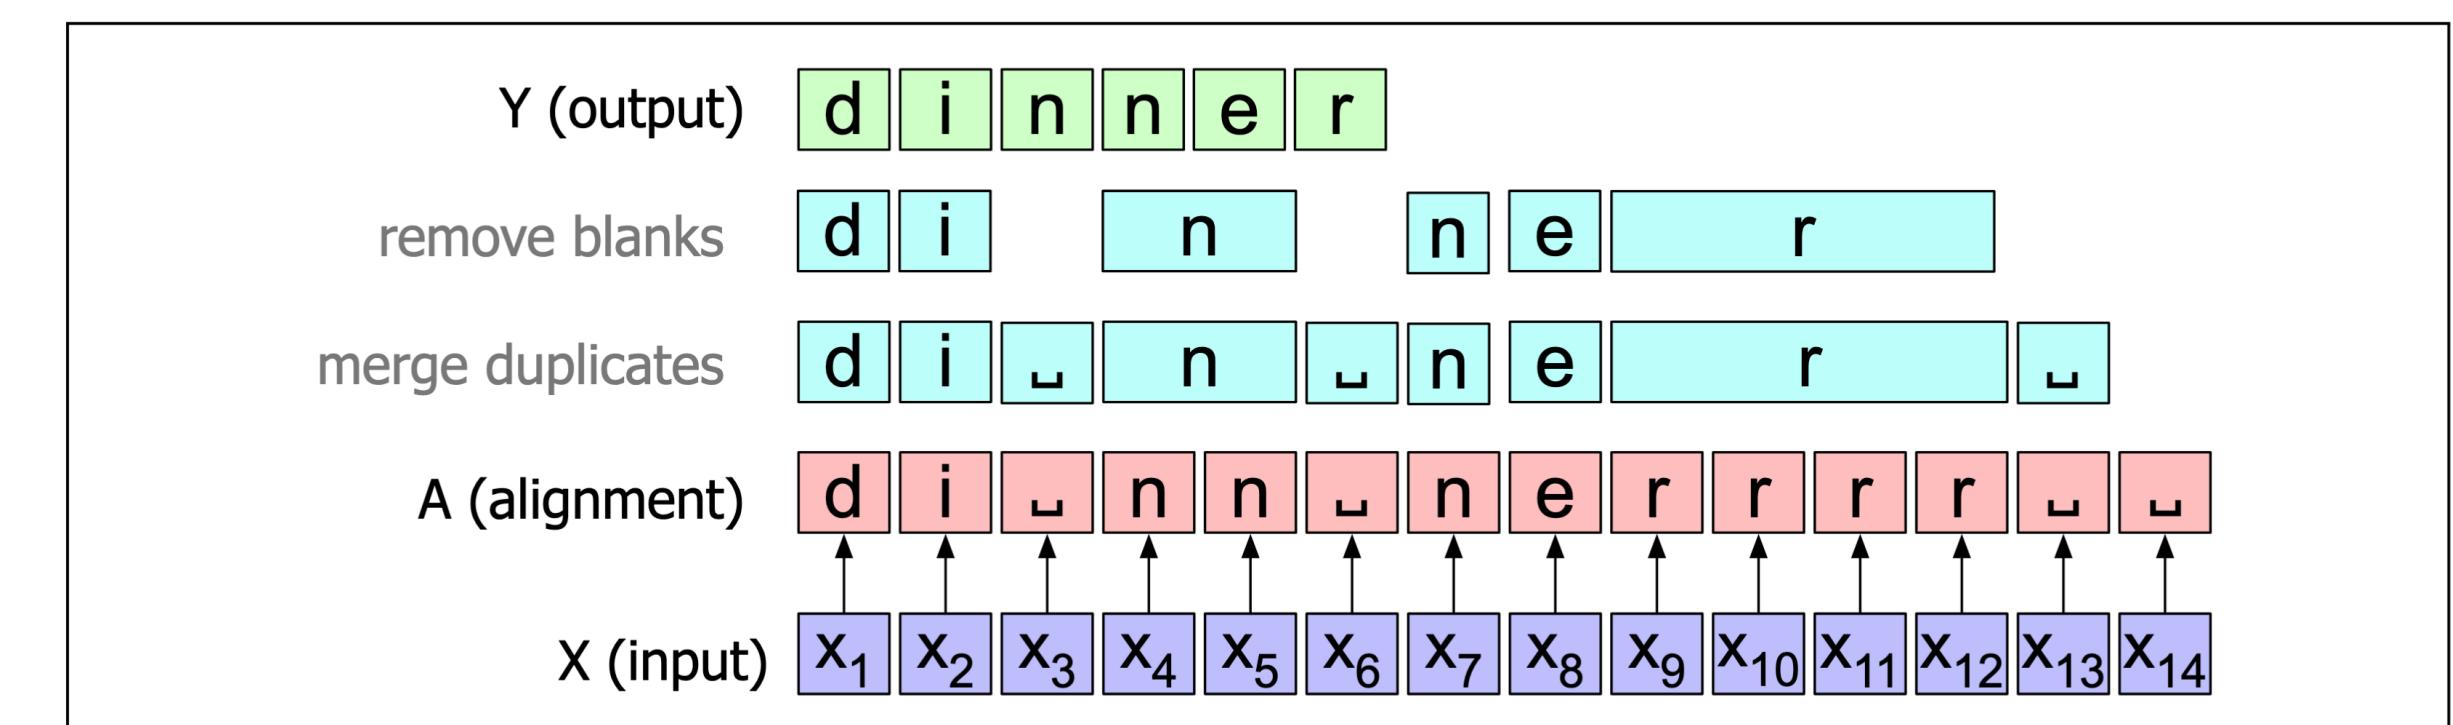


Figure 15.13 The CTC collapsing function **B**, showing the space blank character **u**; repeated (consecutive) characters in an alignment **A** are removed to form the output **Y**.

CTC

- Model still learns to output **one character per timestep**, including the special symbols
- **Multiple alignments** may be compatible with the same output
 - Because of this, calculating **loss** during training requires an approximation of summing over all alignments

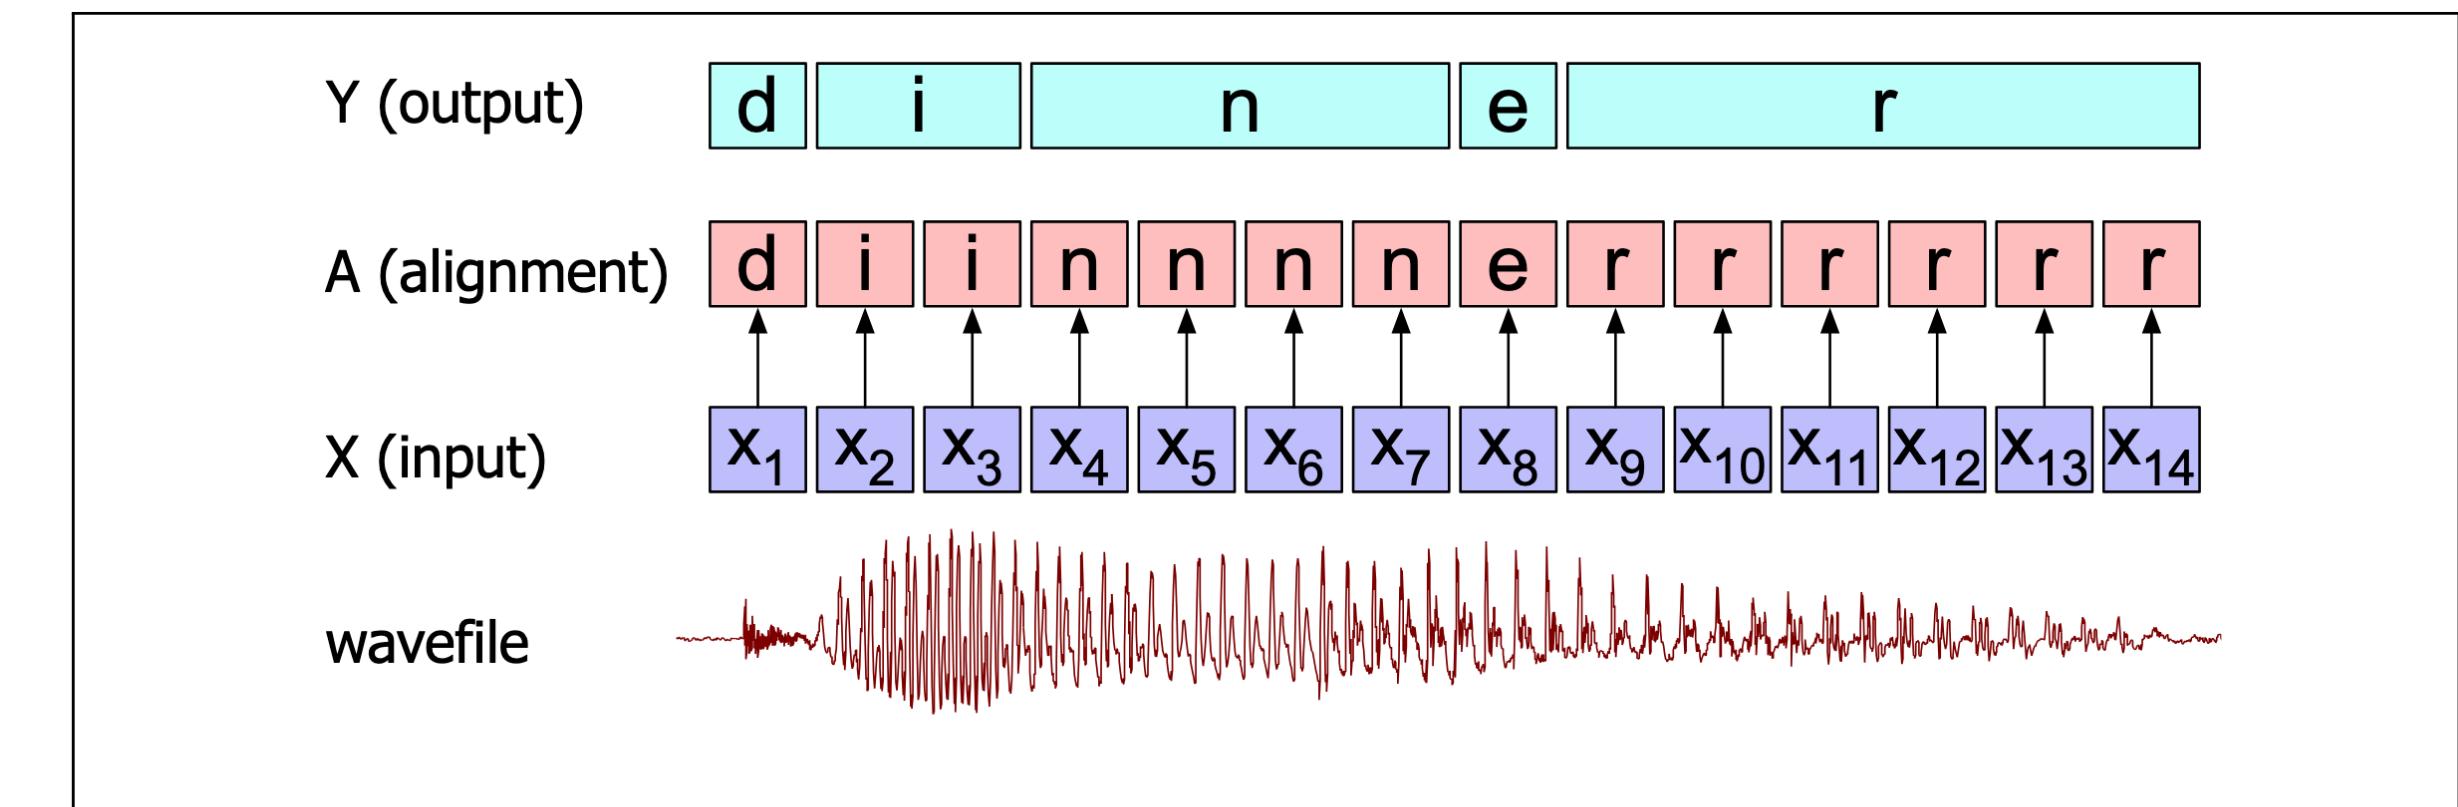


Figure 15.12 A naive algorithm for collapsing an alignment between input and letters.

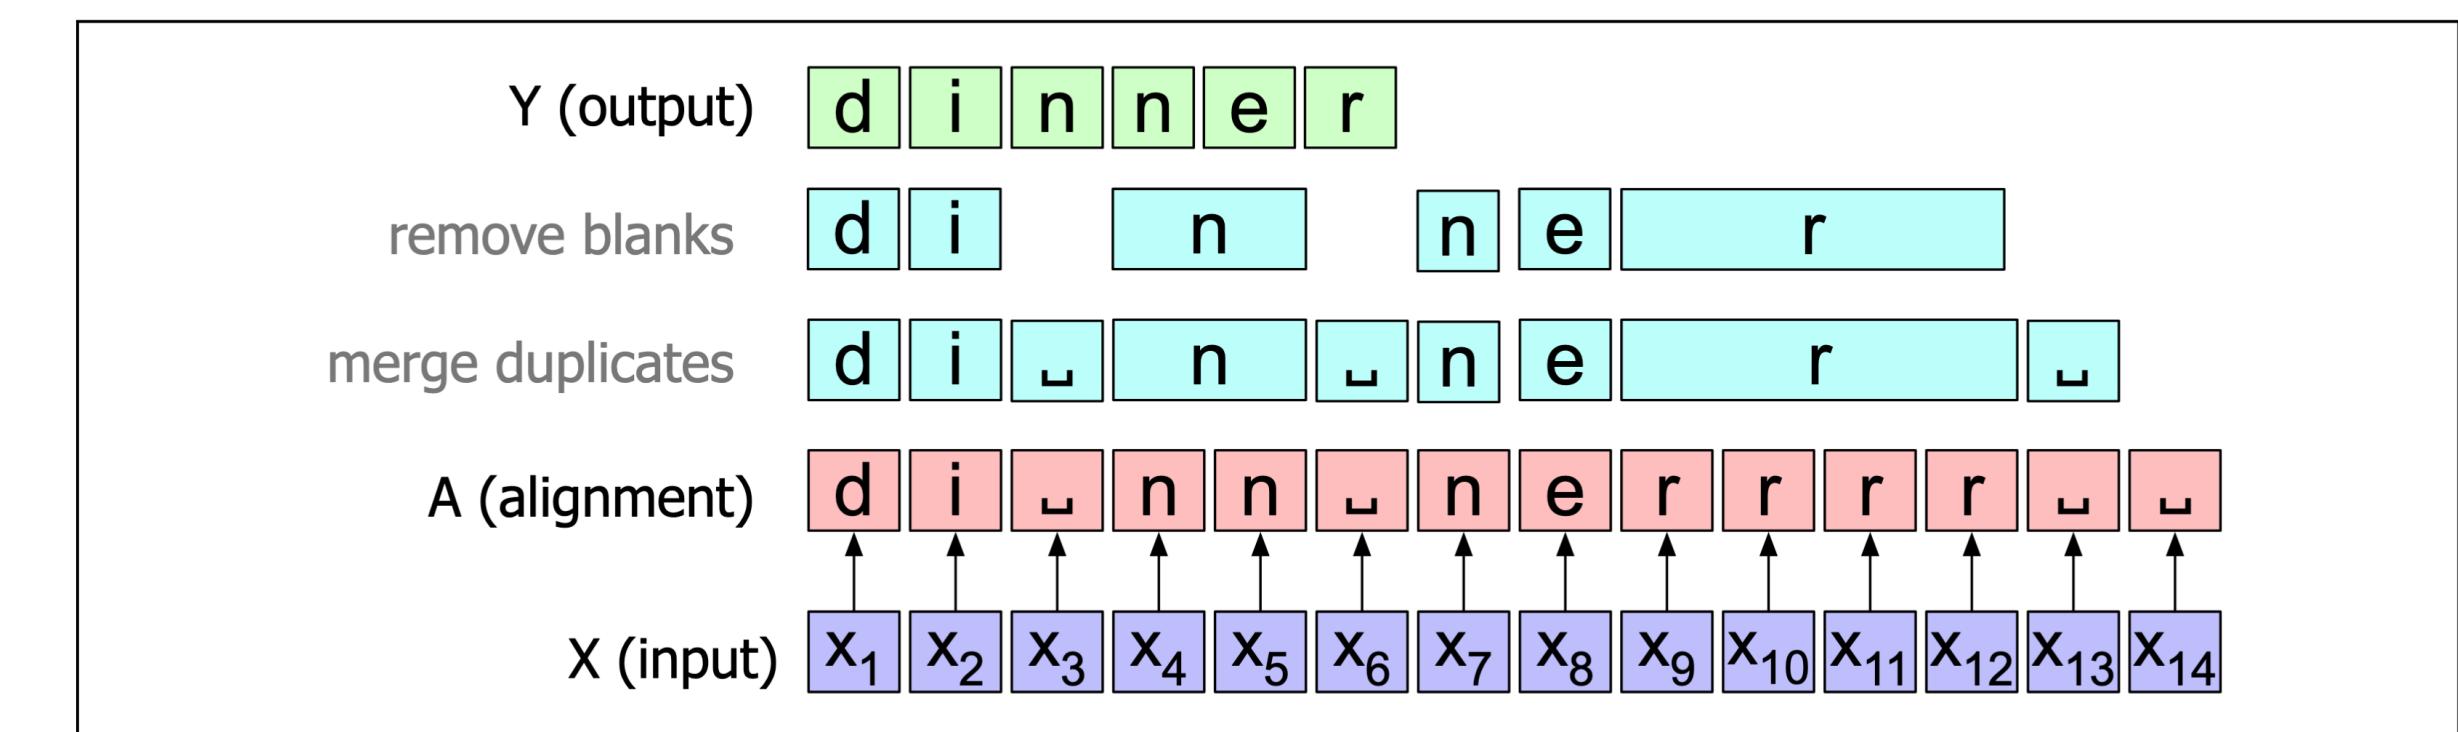


Figure 15.13 The CTC collapsing function B , showing the space blank character **u**; repeated (consecutive) characters in an alignment **A** are removed to form the output **Y**.

LM Integration

- ASR is a **difficult task**
 - Learns to output what seems **most likely given audio**, but might not be sensible language
- Solution: incorporate an **auxiliary language model**
 - LM trained **just on text**
 - Biases system towards **probable sequences of words**

ASR Model Overview

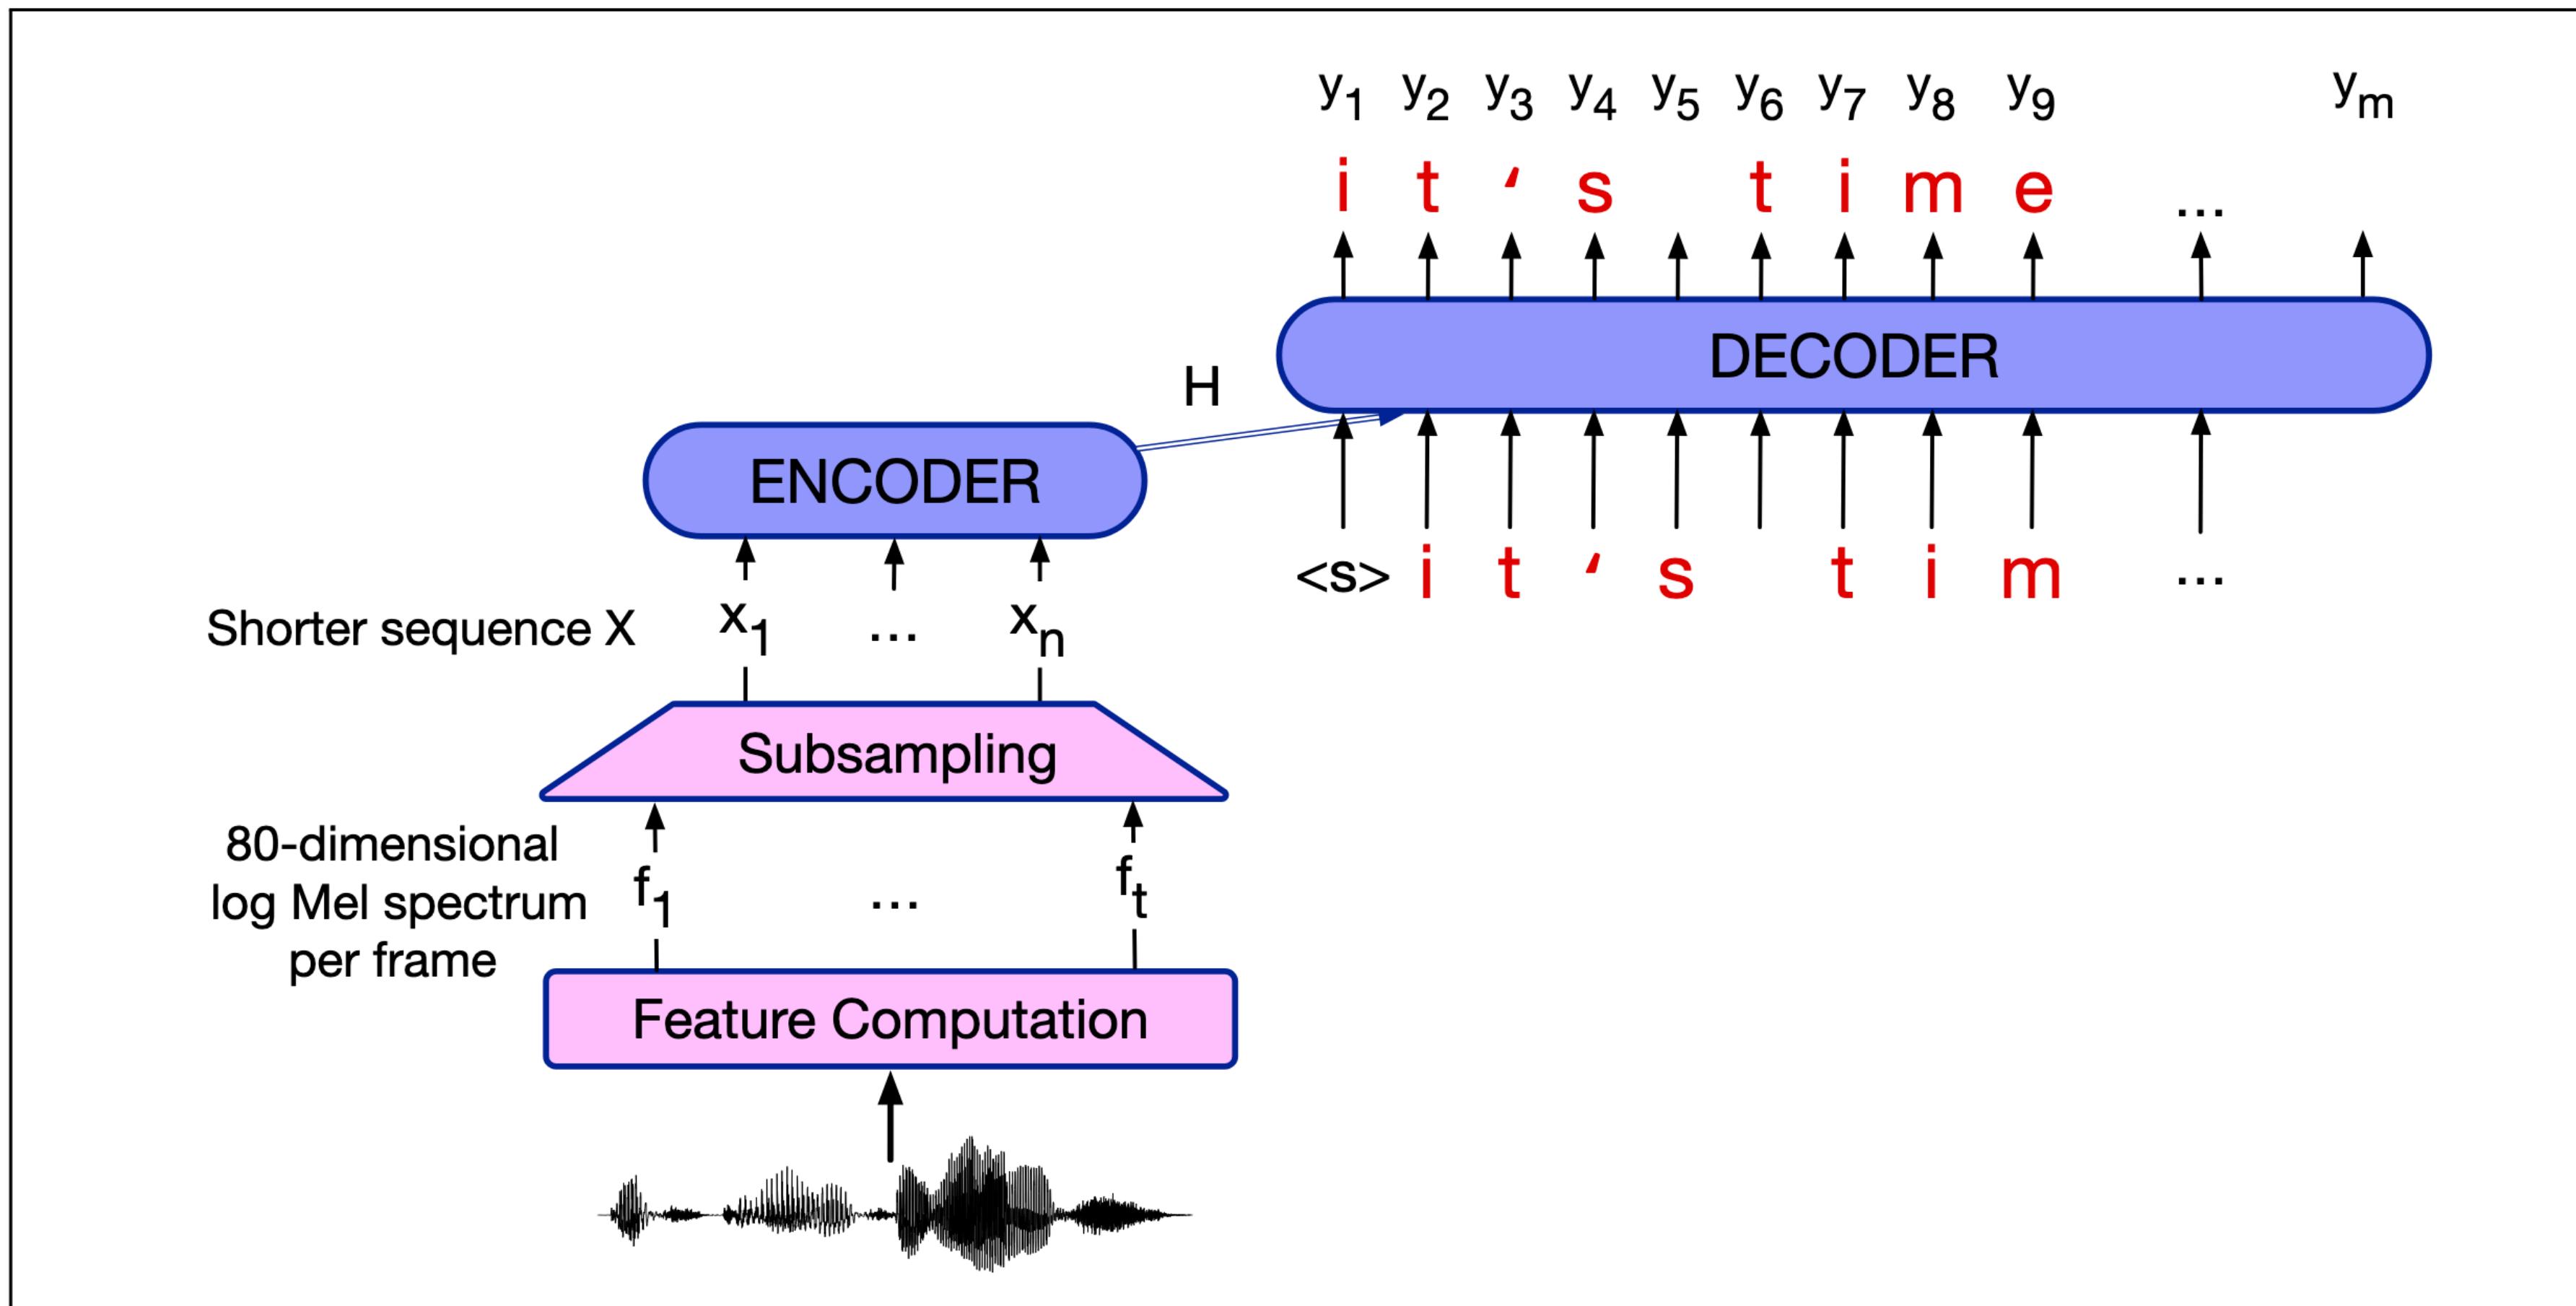
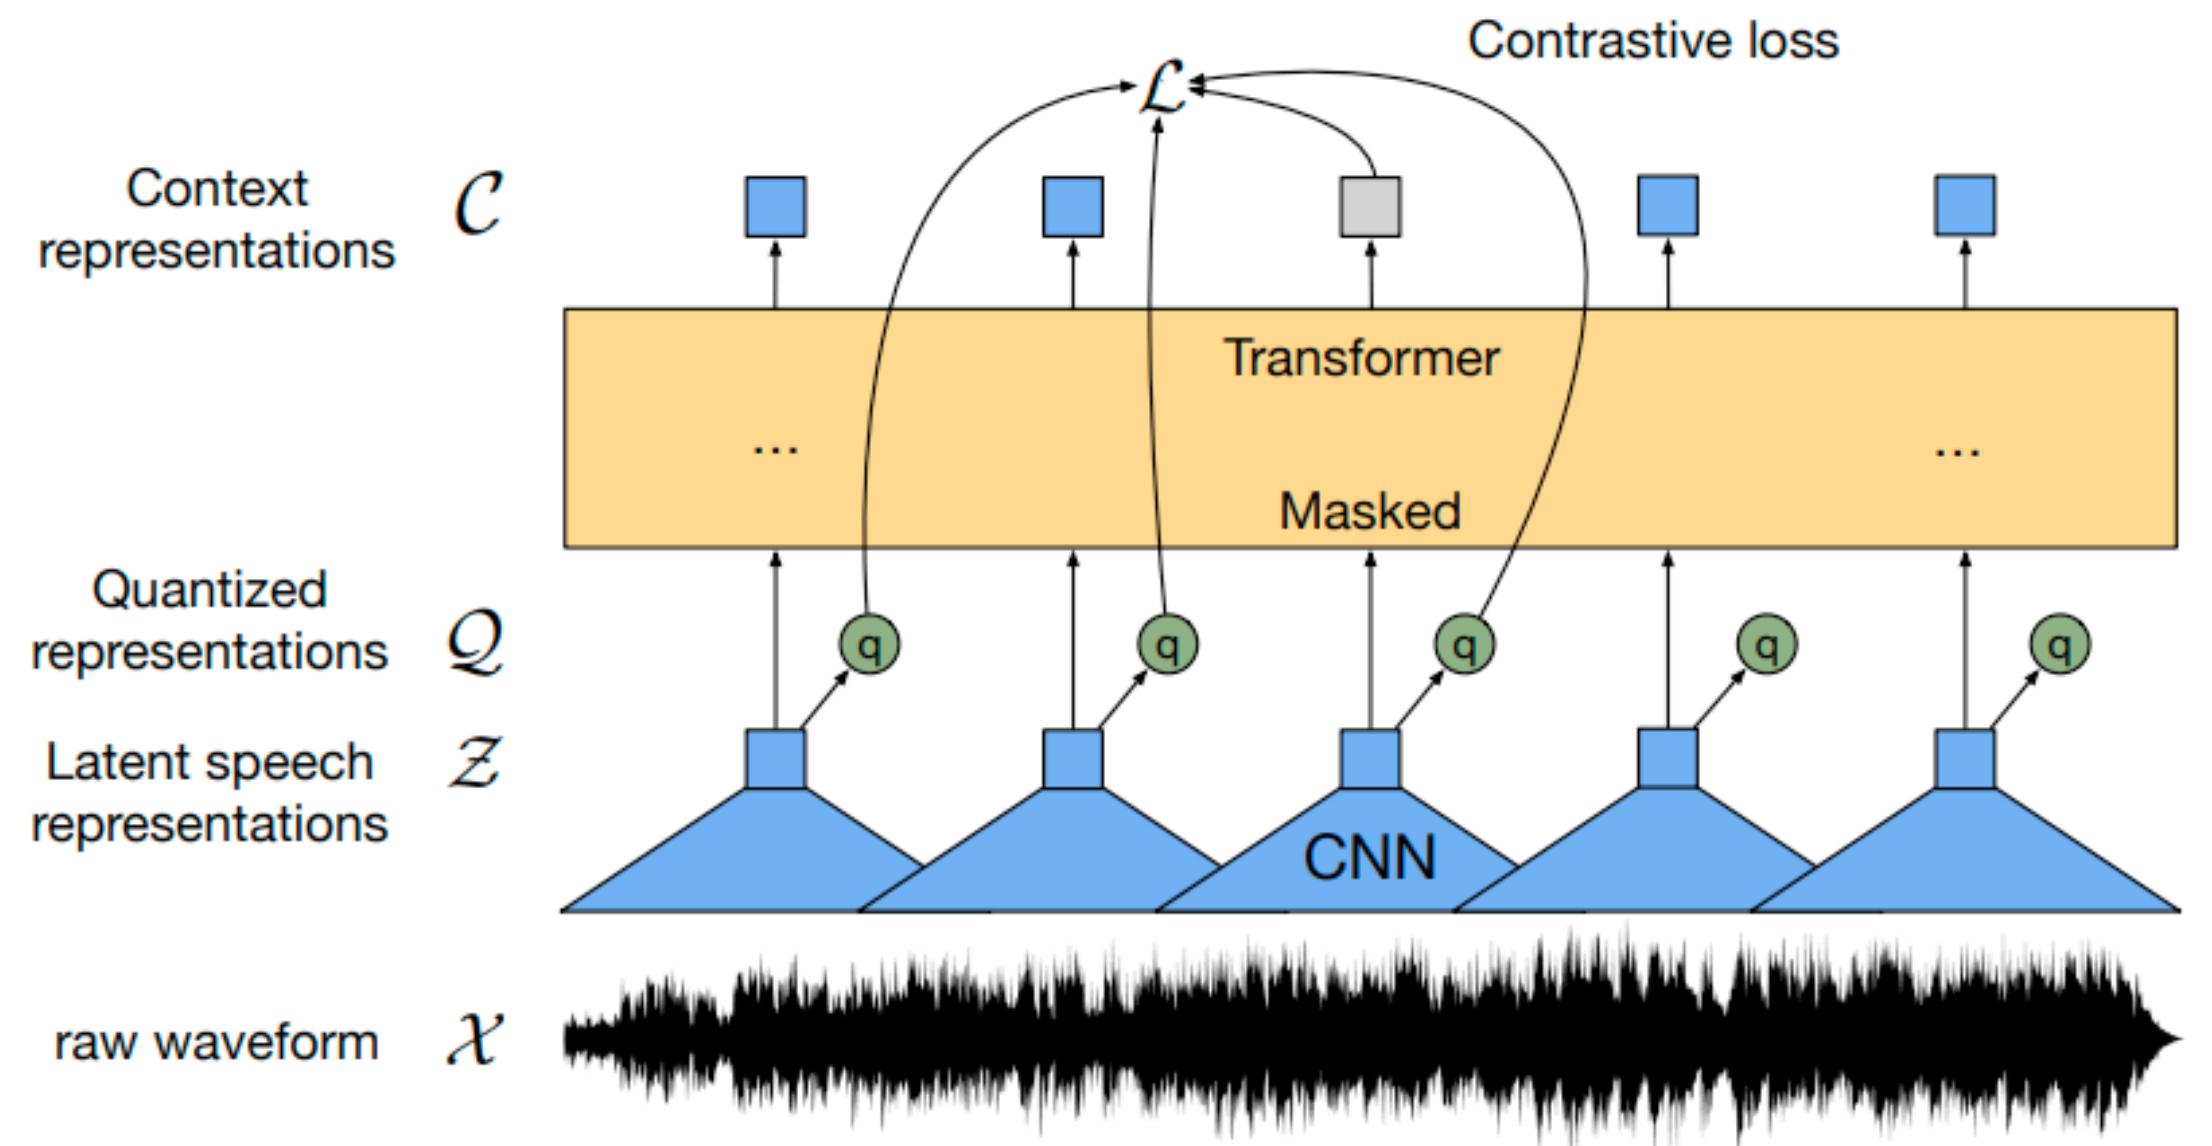


Figure 15.5 Schematic architecture for an encoder-decoder speech recognizer.

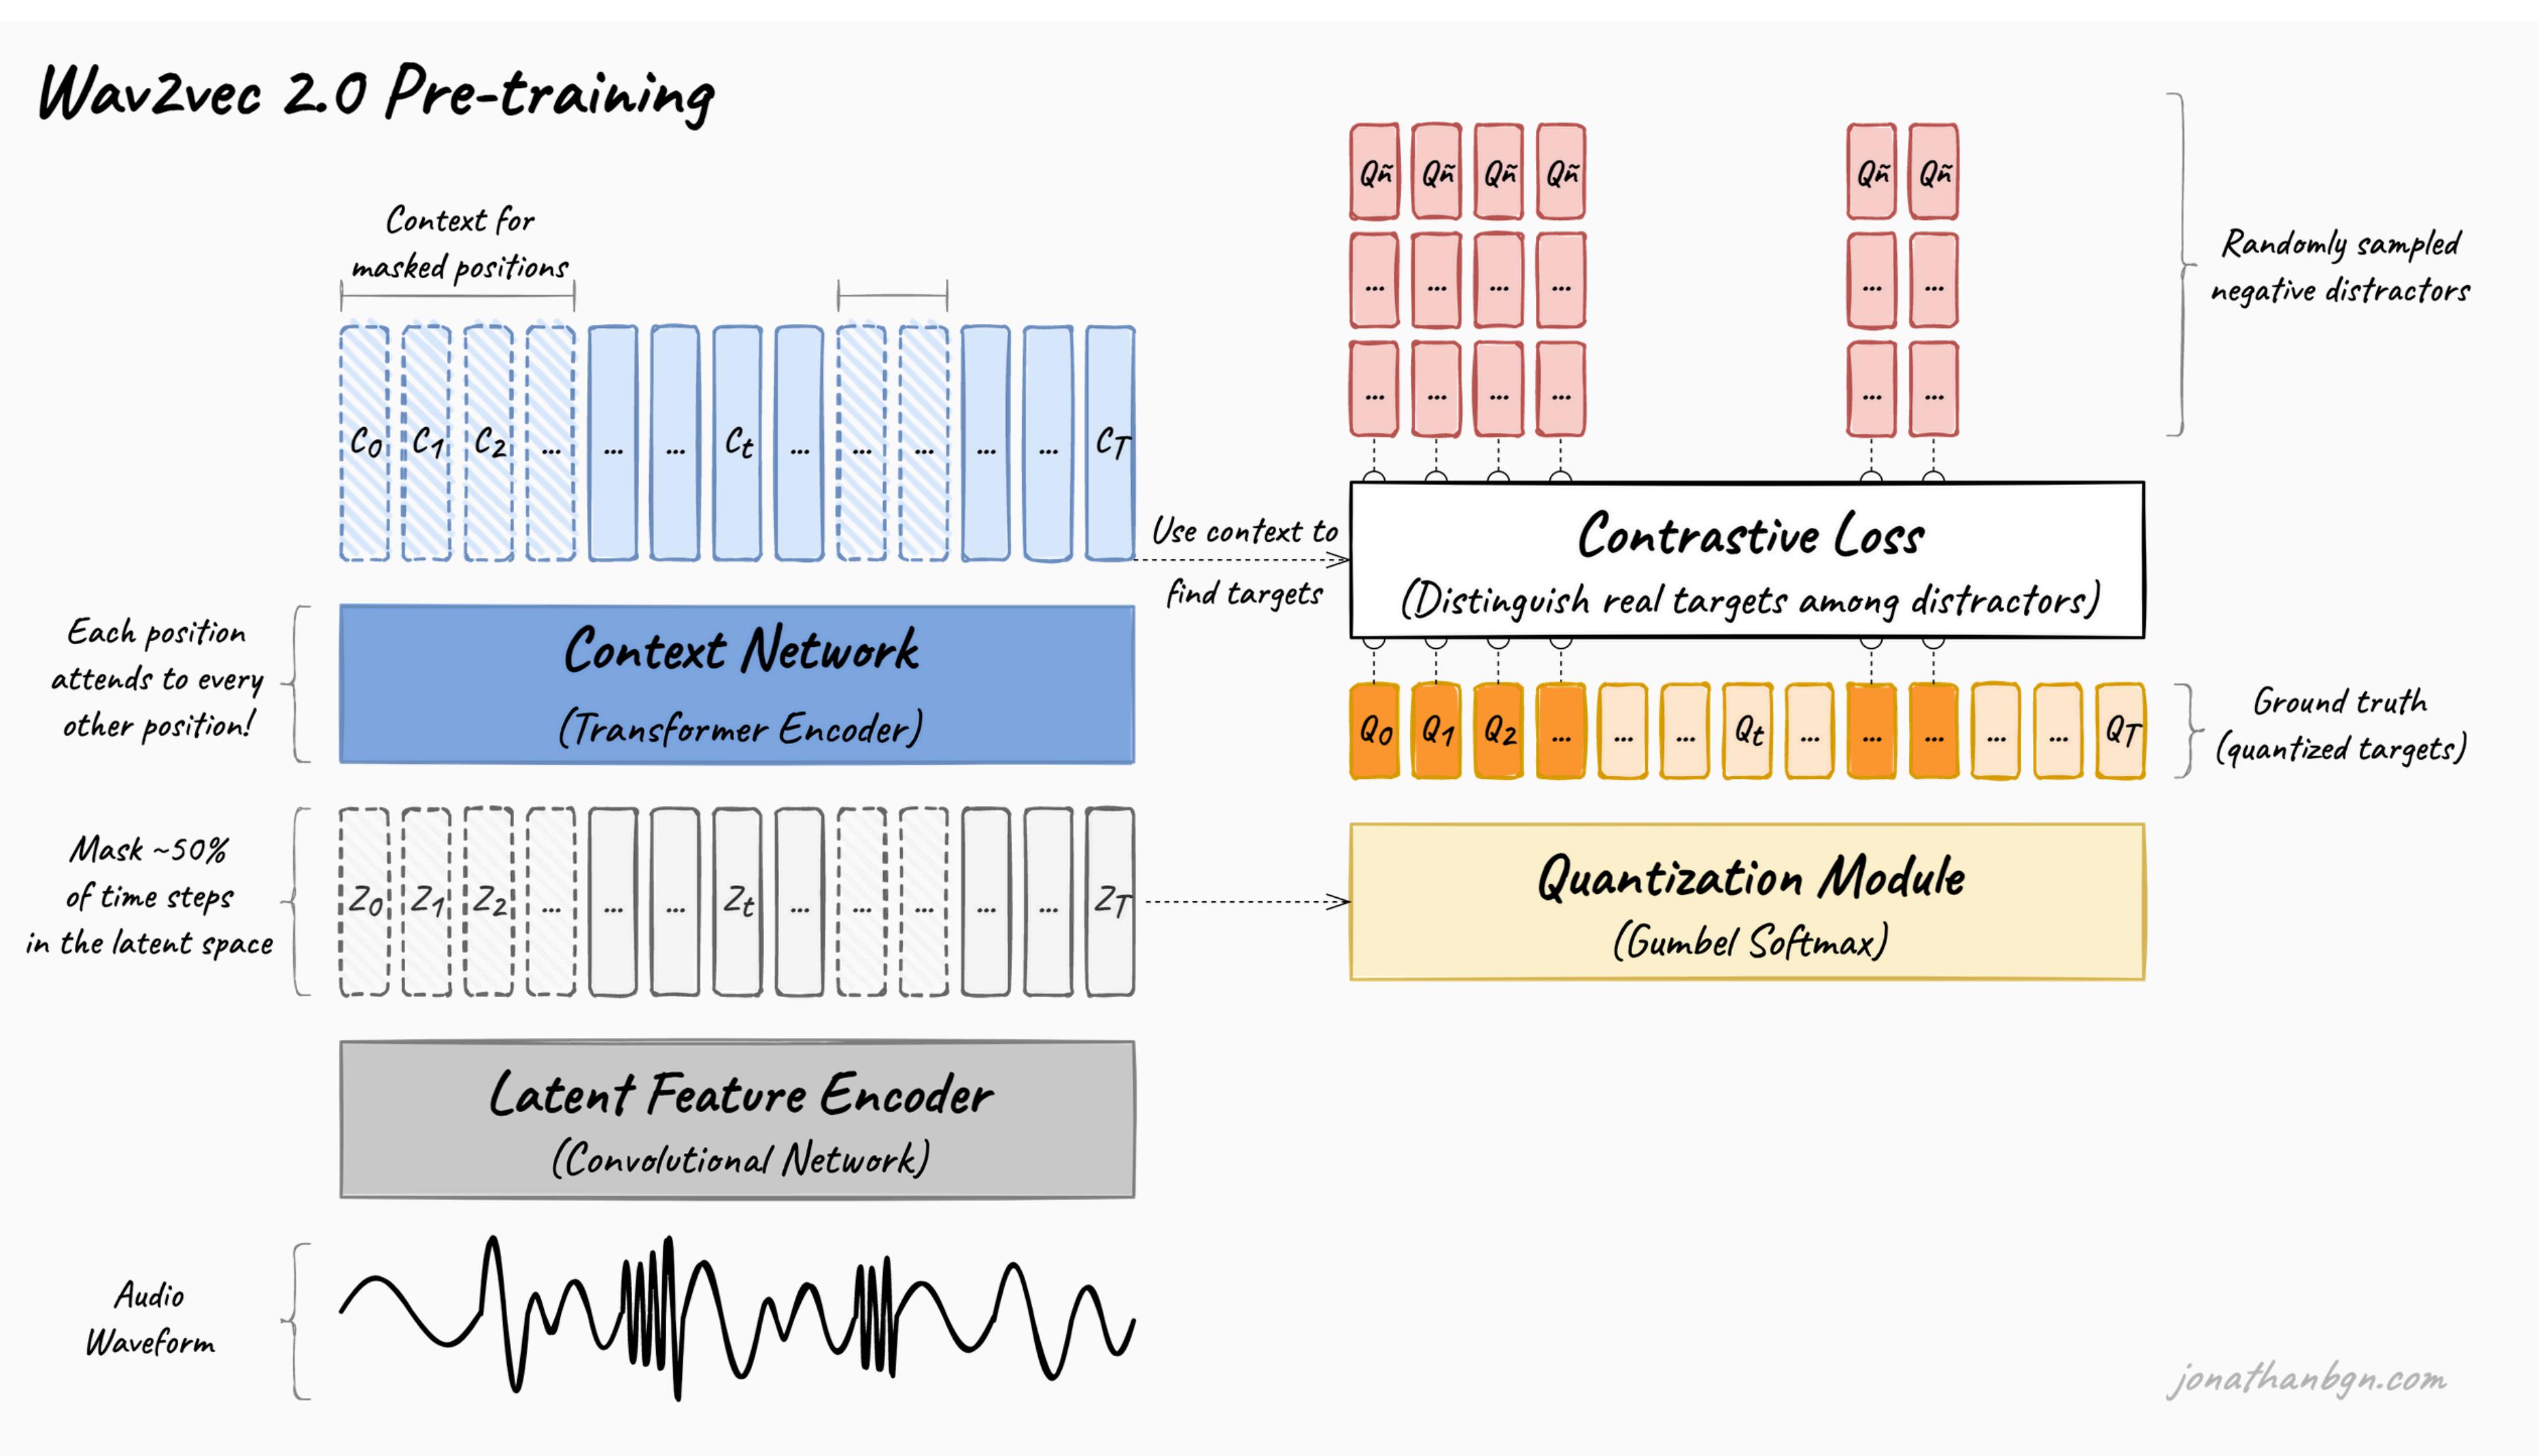
Self-supervised Models

Speech Self-Supervision

- How can we apply **pre-training** to speech models?
- Transcribed audio (needed for ASR training) is **rare?**
- How can we pre-train on **audio only?**
 - Models like **wav2vec** have attempted to do just this
- Visualizations on next slides from [this helpful blog post](#)

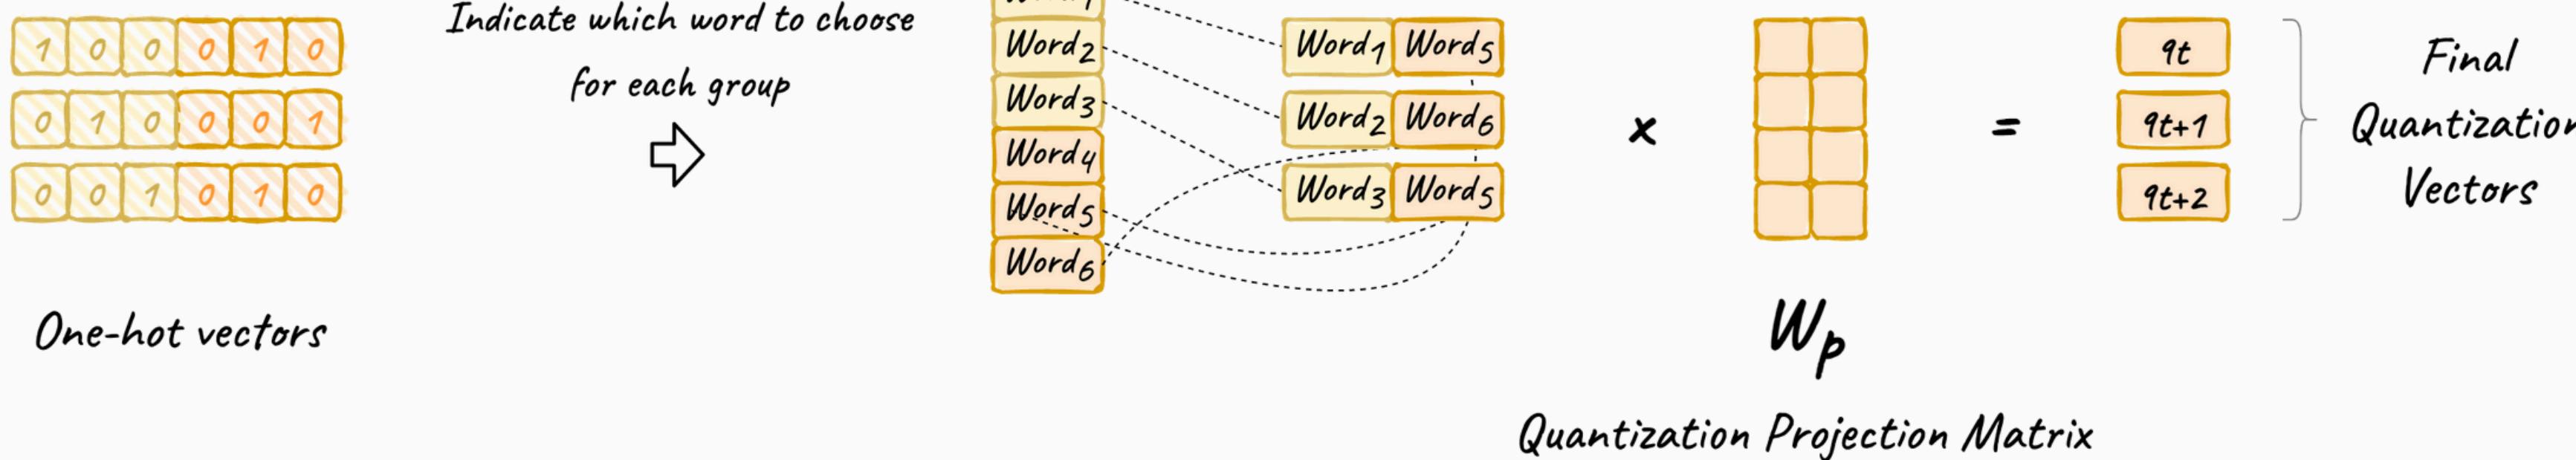
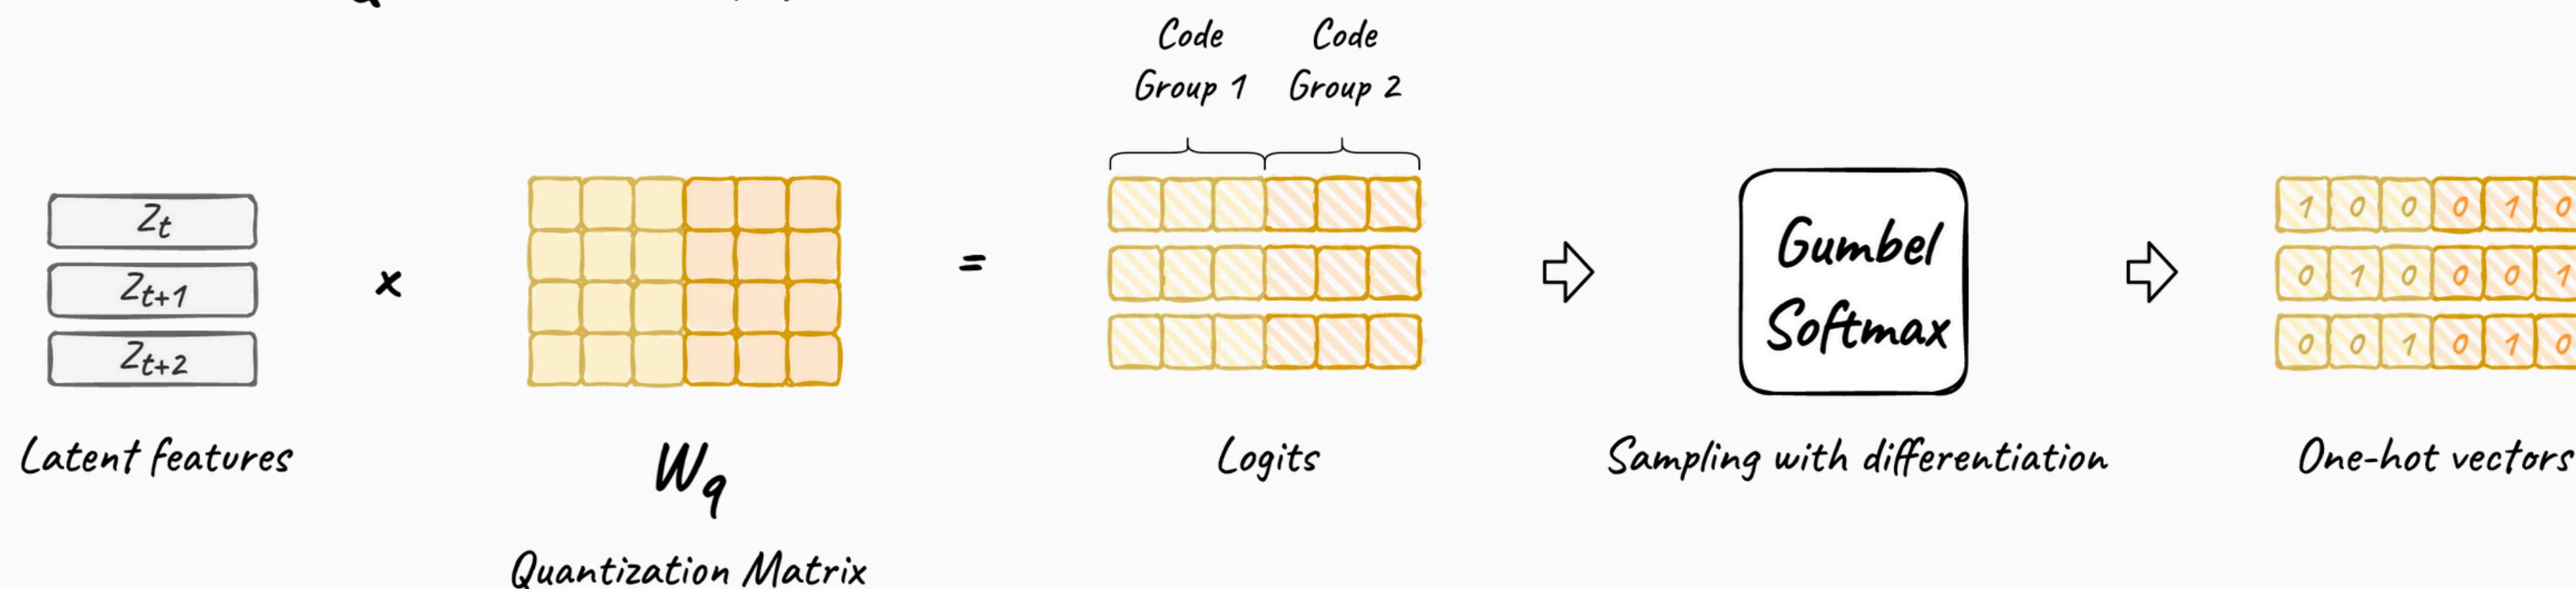


wav2vec Overview



wav2vec Quantization

Wav2vec 2.0 Quantization Module



wav2vec Contrastive Loss

Wav2vec 2.0 Contrastive Loss

